

Image Alignment

Dr. Randy C. Hoover

Department of Electrical and Computer Engineering
South Dakota School of Mines and Technology
Rapid City, SD 57701, USA

November 5, 2012

Review

- ▶ Last few lectures:
 - ▶ Finding distinctive features in multiple images
 - ▶ Describing said features using a feature descriptor
 - ▶ Corresponding each descriptor for alignment

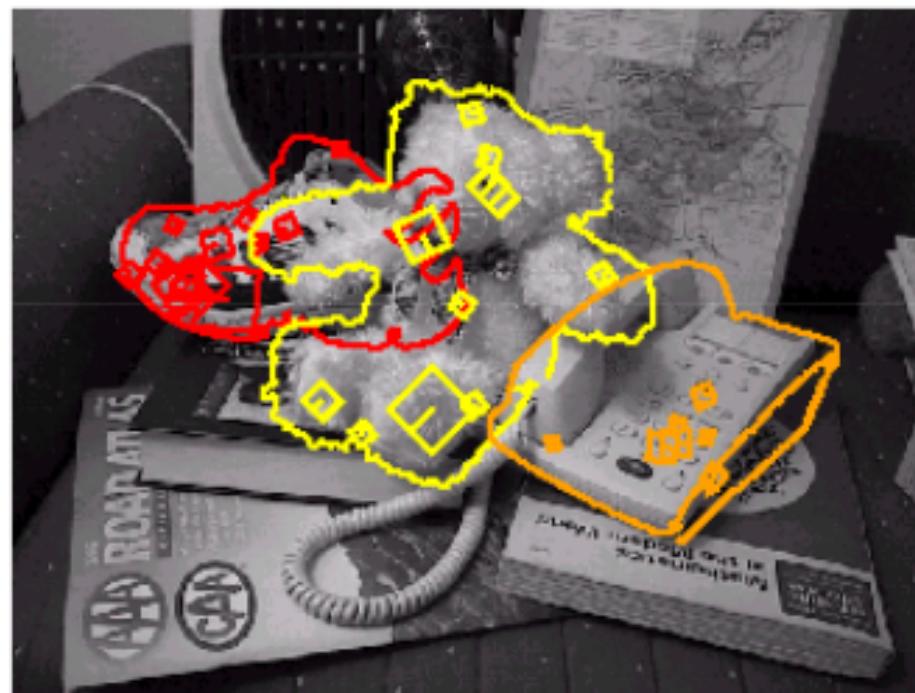
Today (and next few lectures)

- ▶ Image alignment

- ▶ Computing homographies
- ▶ Rotational panoramics
- ▶ RANSAC
- ▶ Global alignment
- ▶ Warping
- ▶ Blending

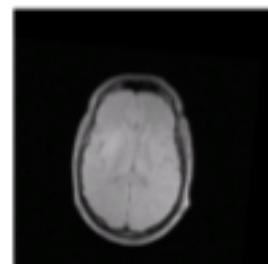
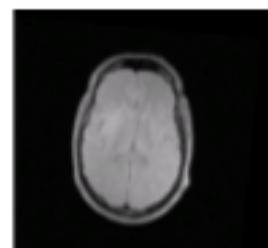
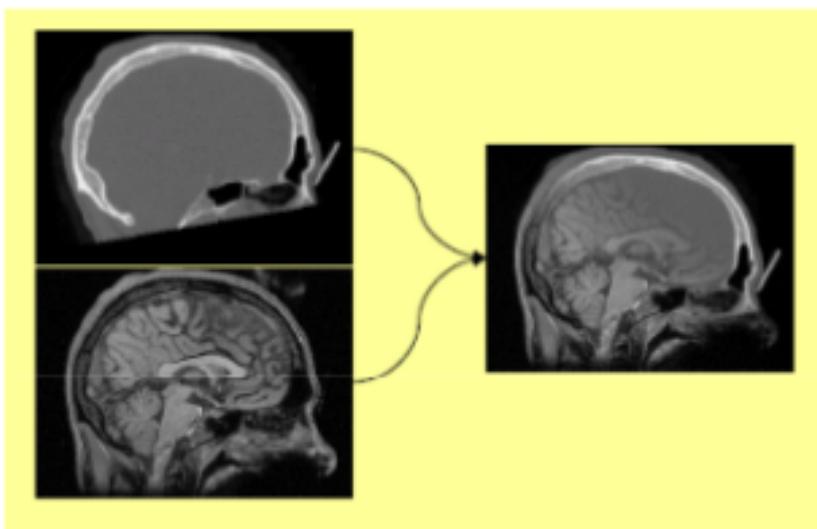


Motivation: Recognition in cluttered scenes



Figures from David Lowe

Motivation: Image registration



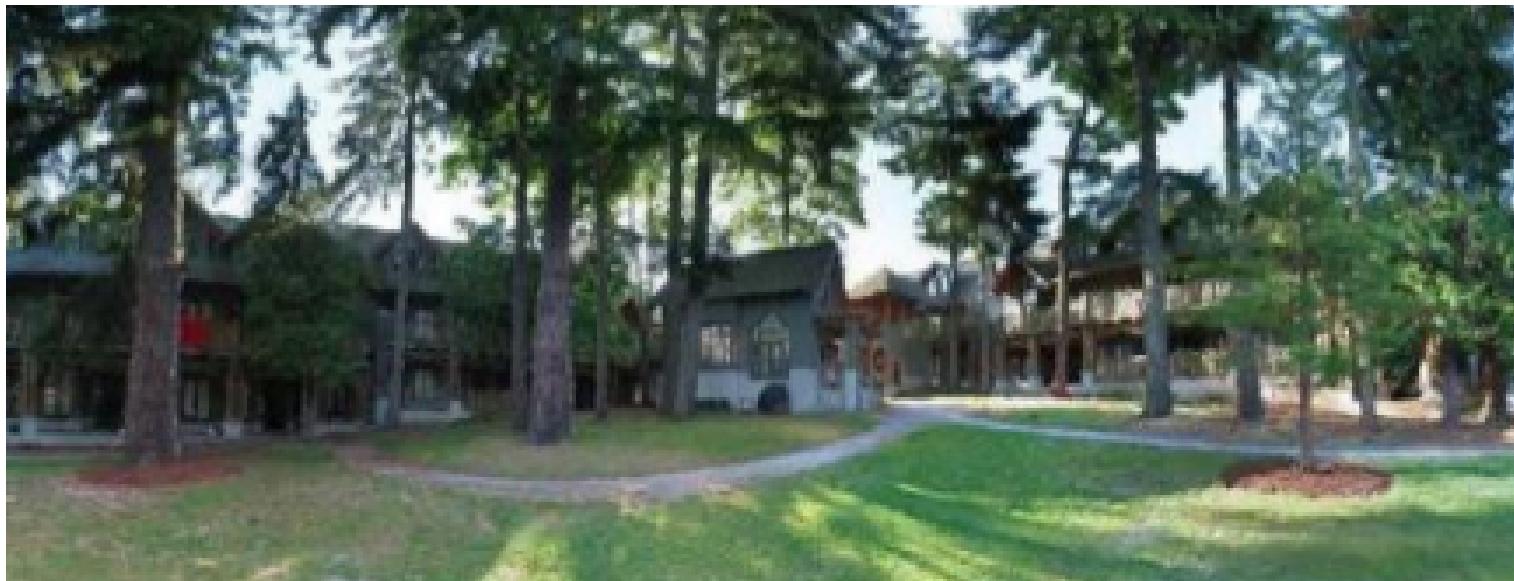
Motivation: Mosaics

- ▶ Getting the whole picture
 - ▶ Consumer camera: $50^\circ \times 35^\circ$



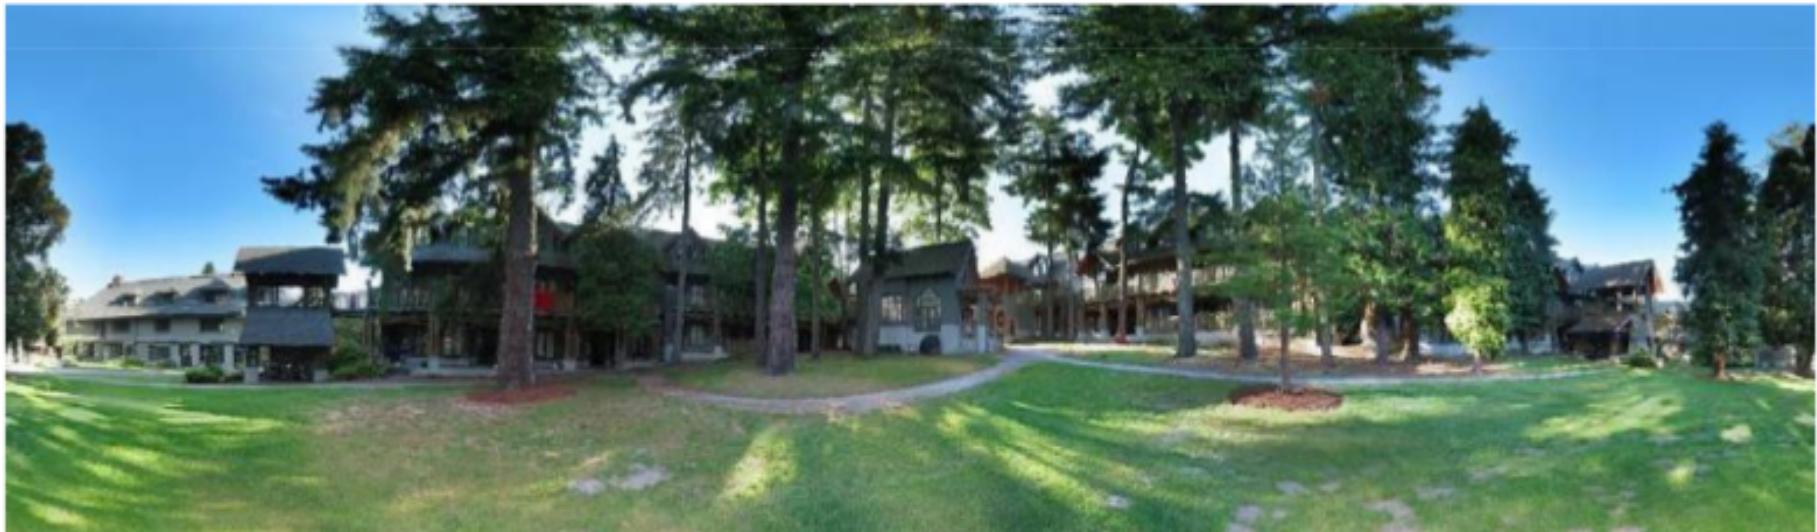
Motivation: Mosaics

- ▶ Getting the whole picture
 - ▶ Consumer camera: $50^\circ \times 35^\circ$
 - ▶ Human vision: $176^\circ \times 135^\circ$



Motivation: Mosaics

- ▶ Getting the whole picture
 - ▶ Consumer camera: $50^\circ \times 35^\circ$
 - ▶ Human vision: $176^\circ \times 135^\circ$



- Panoramic Mosaic = up to $360 \times 180^\circ$

Slide from Brown & Lowe 2003

General idea

- ▶ Correspondence
 - ▶ Detect interest points
 - ▶ Describe interest points
 - ▶ Match two or more points

General idea

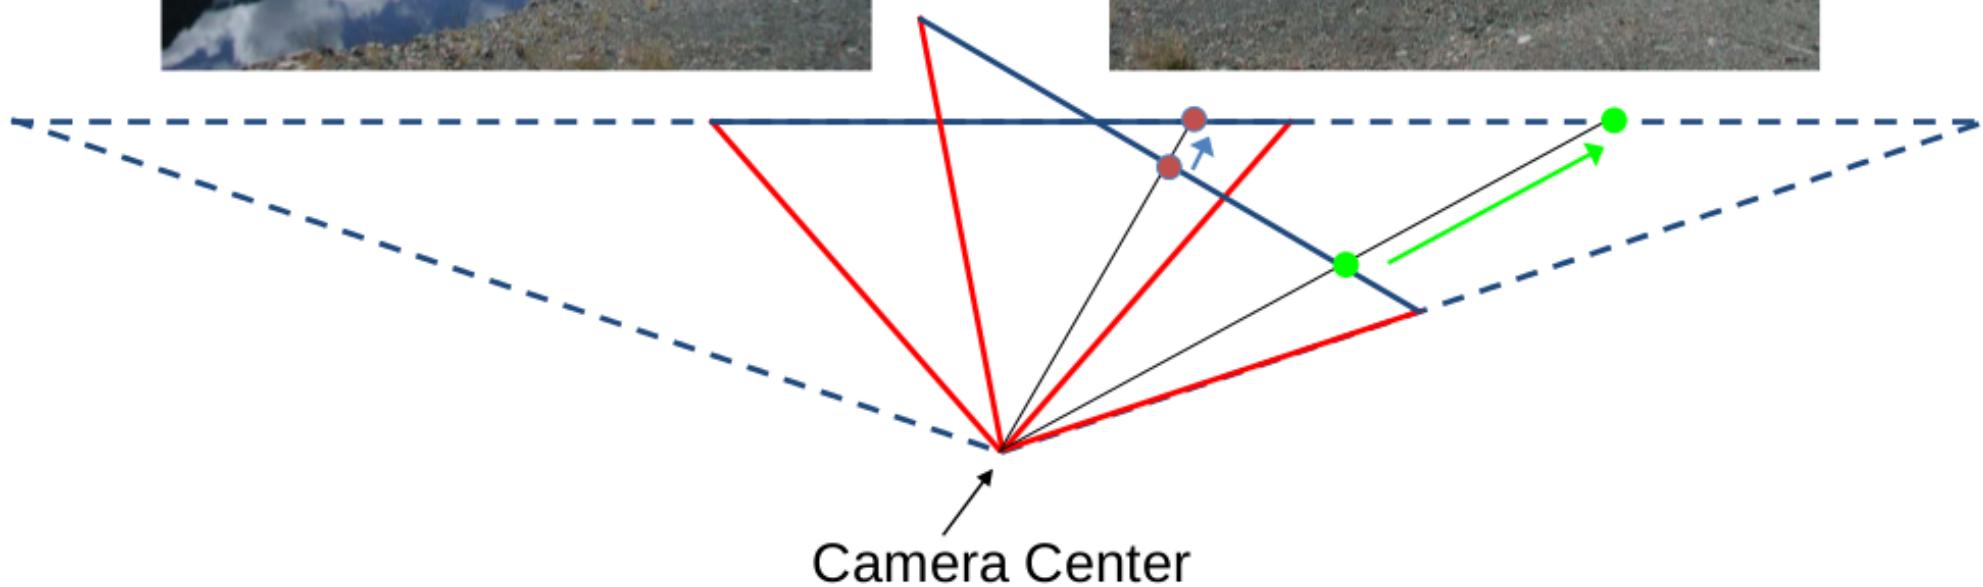
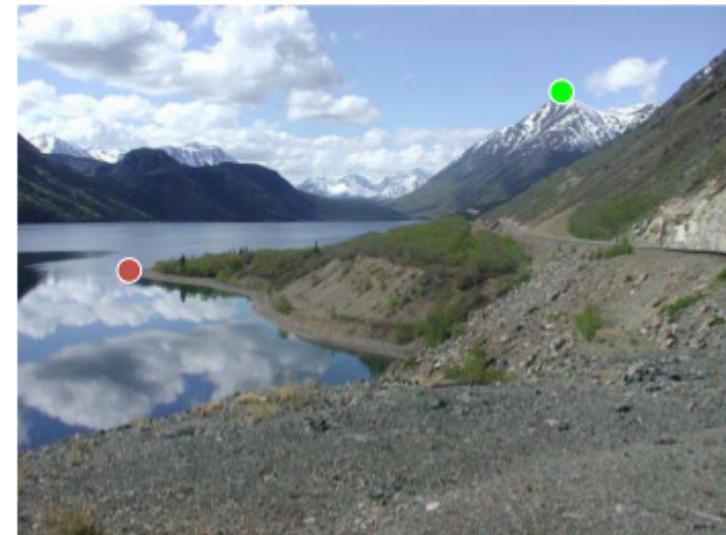
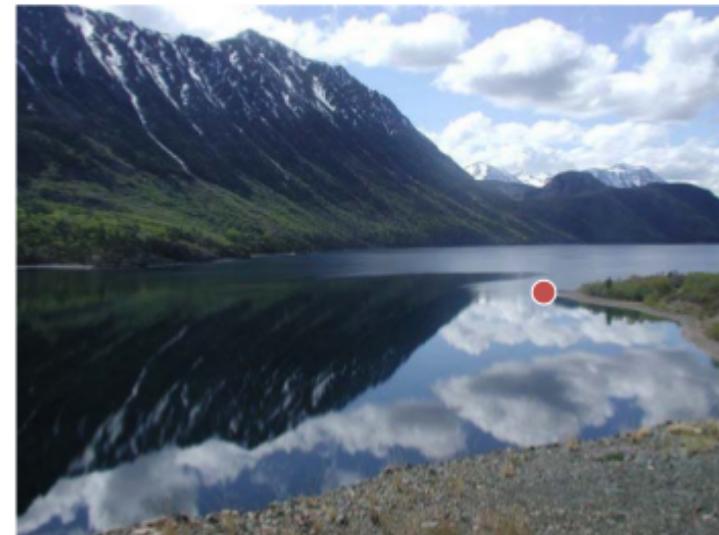
- ▶ Correspondence
 - ▶ Detect interest points
 - ▶ Describe interest points
 - ▶ Match two or more points
- ▶ Alignment and model fitting
 - ▶ Solve for motion or depth or camera movement from multiple matched points

General idea - image stitching

- ▶ Combine two or more overlapping images to make one larger one
- ▶ Cool HP vid. <http://www.youtube.com/watch?v=2RPl5vPEoQk>



Example



What types of motion can I get?

- ▶ What happens when we take two images from the same camera and try to align them?

- ▶ Translation?
- ▶ Rotation?
- ▶ Scale changes?
- ▶ Affine transformation?
- ▶ Perspective?
- ▶ Other....?



Recall: 2-D coordinate transformations

- ▶ Let $\mathbf{x} = (x, y)$
 - ▶ Translation: $\mathbf{x}' = \mathbf{x} + \mathbf{t}$
 - ▶ Rotation: $\mathbf{x}' = R\mathbf{x} + \mathbf{t}$, $R \in SO(2)$
 - ▶ Similarity: $\mathbf{x}' = \alpha R\mathbf{x} + \mathbf{t}$
 - ▶ Affine: $\mathbf{x}' = A\mathbf{x} + \mathbf{t}$
 - ▶ Perspective: $\bar{\mathbf{x}}' = H\bar{\mathbf{x}}$ where $\bar{\mathbf{x}}$ represents homogeneous coords:
 $\bar{\mathbf{x}} = (x, y, 1)$ and $H \in \mathbb{R}^{3 \times 3}$

Recall: 2-D coordinate transformations

- We can write these as 3×3 matrices using homogeneous coords.

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & t_x \\ 0 & 1 & t_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Translate

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} s_x & 0 & 0 \\ 0 & s_y & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Scale

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} \cos \Theta & -\sin \Theta & 0 \\ \sin \Theta & \cos \Theta & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Rotate

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & sh_x & 0 \\ sh_y & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Shear

Recall: 2-D affine transformations

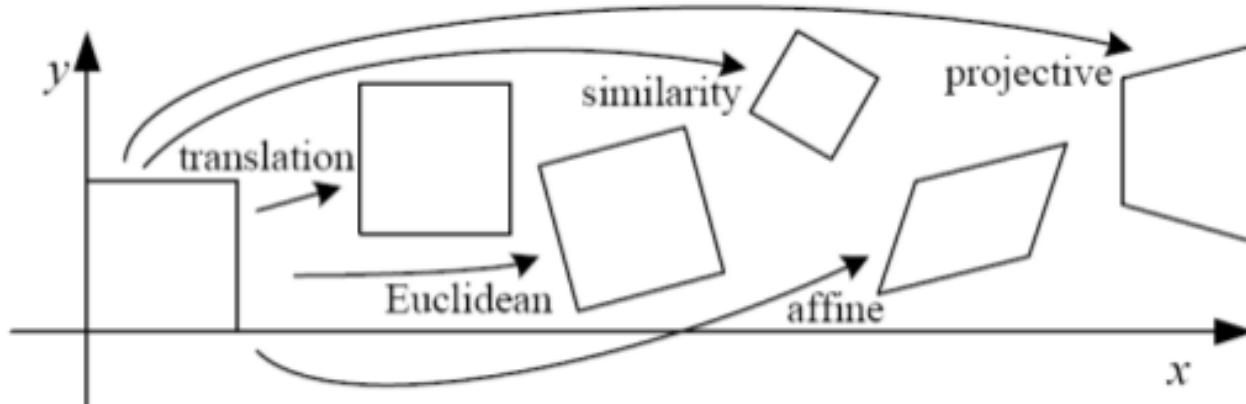
$$\begin{bmatrix} x' \\ y' \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

- ▶ Affine transformations are a combination of:
 - ▶ Linear transformations, and
 - ▶ Translations
 - ▶ \implies parallel lines remain parallel

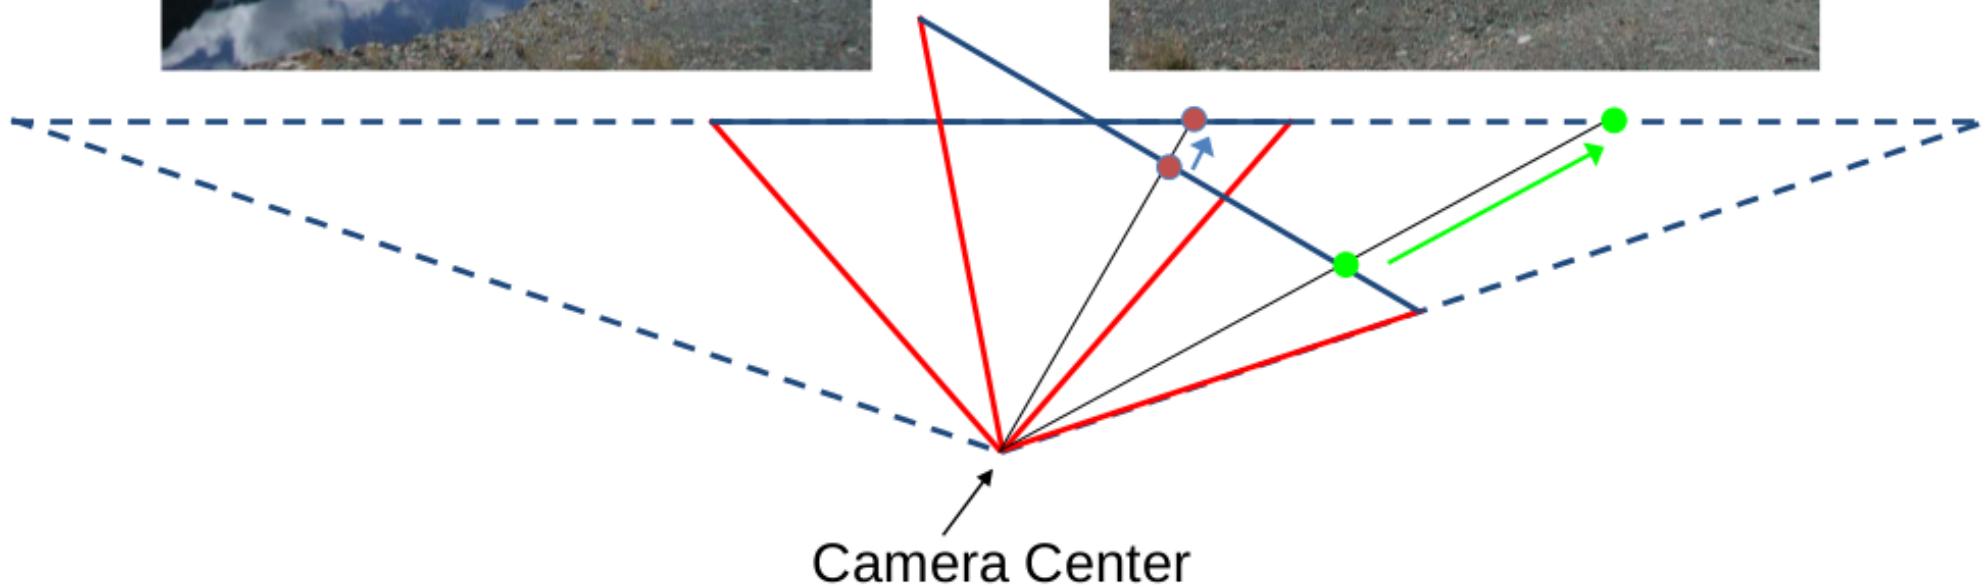
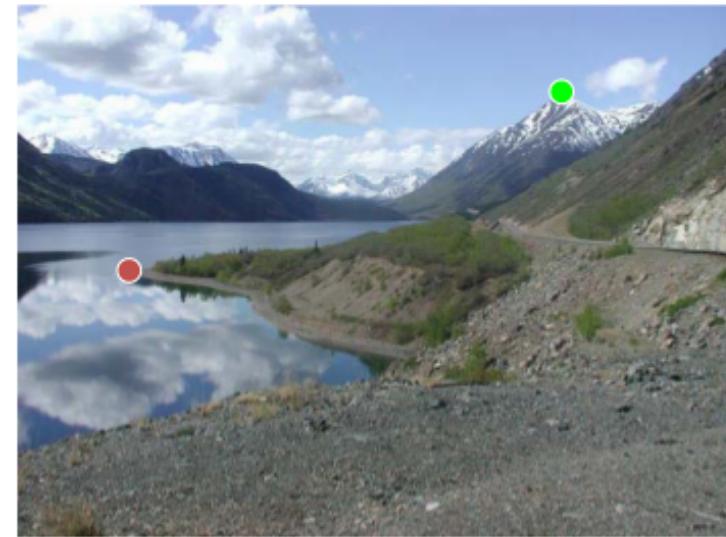
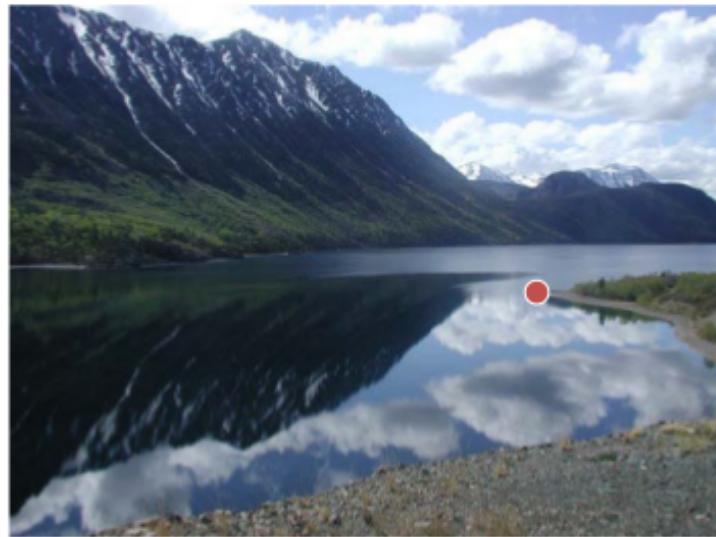
Recall: projective transformations

$$\begin{bmatrix} x' \\ y' \\ w \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix} \begin{bmatrix} x \\ y \\ w \end{bmatrix}$$

- ▶ Projective transformations are a combination of:
 - ▶ Affine transformations, and
 - ▶ Projective warping
 - ▶ \Rightarrow parallel lines **do not** remain parallel (in general)



Example - revisited



How do we analyze this?

- ▶ Review of projective geometry (in homogeneous coords)!
 - ▶ $\mathbf{x} = (\alpha R + \mathbf{t}) X$
 - ▶ $\mathbf{x}' = (\alpha' R' + \mathbf{t}') X$
 - ▶ Assume camera rotates about its optical axis $\implies \mathbf{t} = \mathbf{t}' = 0$
 - ▶ This $\implies \mathbf{x}' = H\mathbf{x}$ where,
 - ▶ $H = \alpha' R' R^{-1} \alpha^{-1}$
- ▶ Typically only R and f, f' change (4-parameters), however, in general, H has 8-parameters since $H(3, 3) = 1$

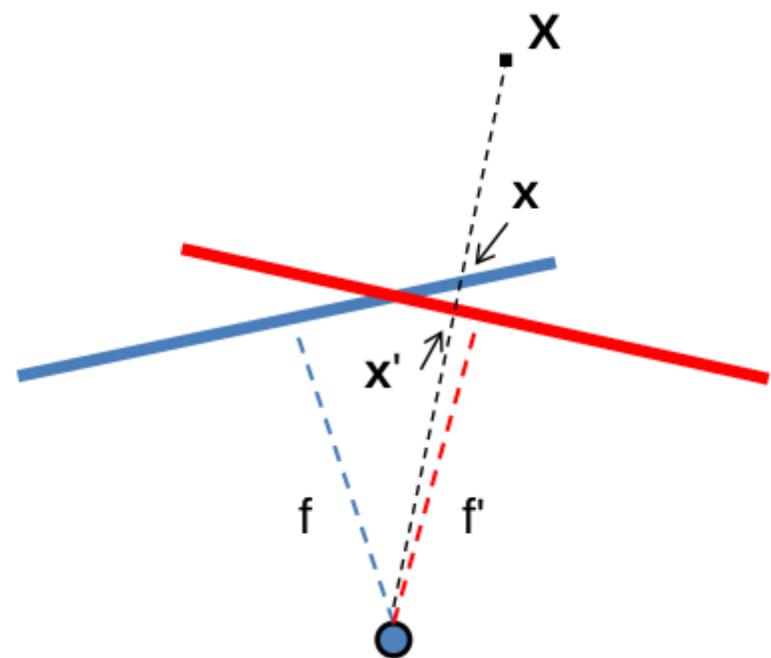
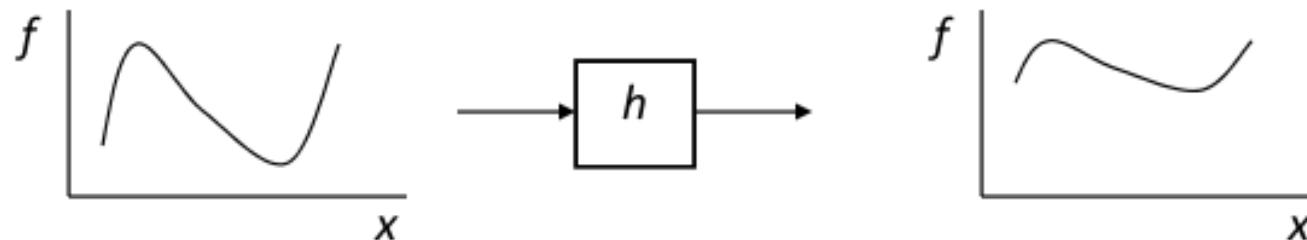


Image warping

- ▶ Image filtering: change the *range* of the image

- $g(x) = h(f(x))$



- ▶ Image warping: change the *domain* of the image

- $g(x) = f(h(x))$

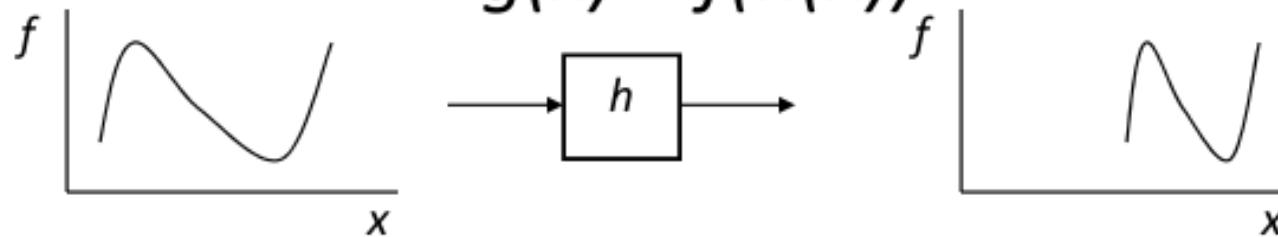
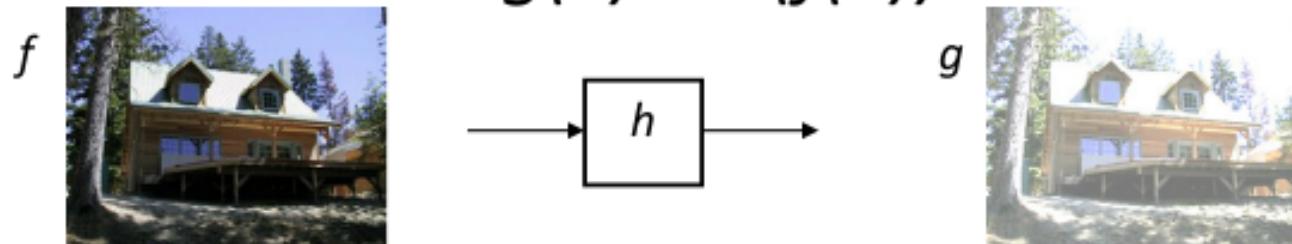


Image warping

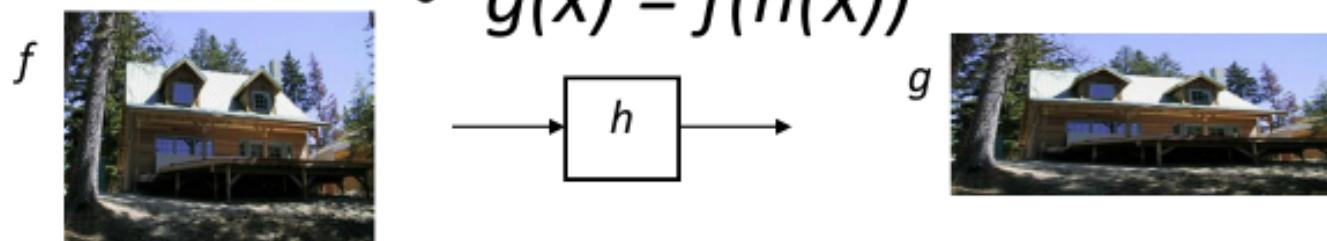
- ▶ Image filtering: change the *range* of the image

- $g(x) = h(f(x))$



- ▶ Image warping: change the *domain* of the image

- $g(x) = f(h(x))$



Parametric (global) warping

- ▶ Examples of parametric warping:



translation



rotation



aspect



affine



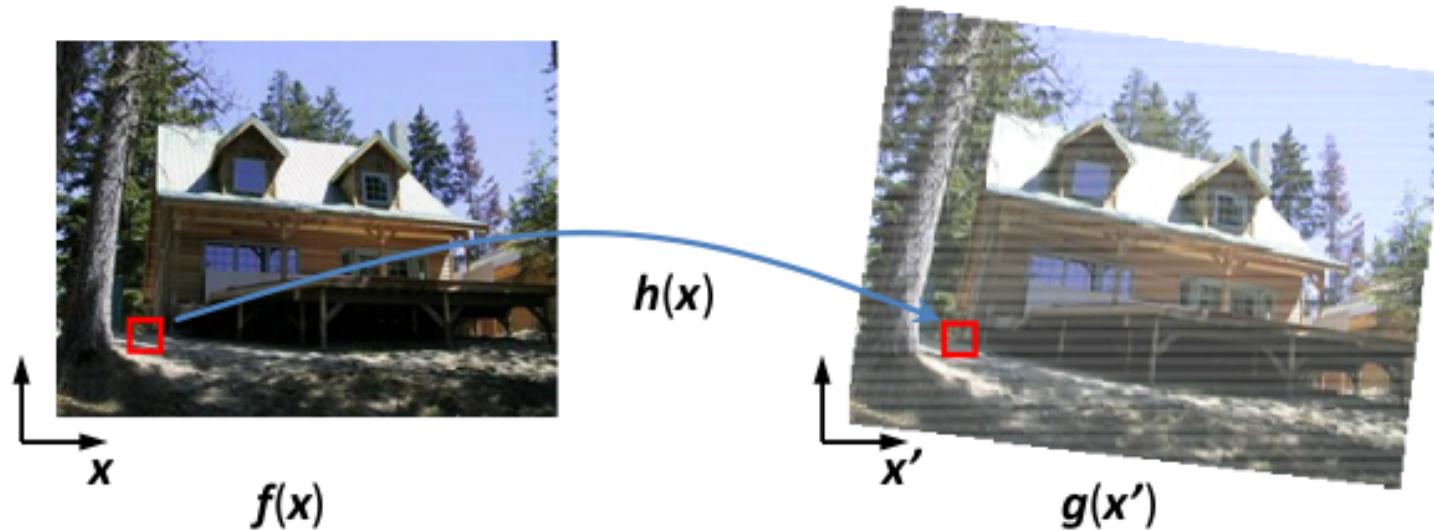
perspective



cylindrical

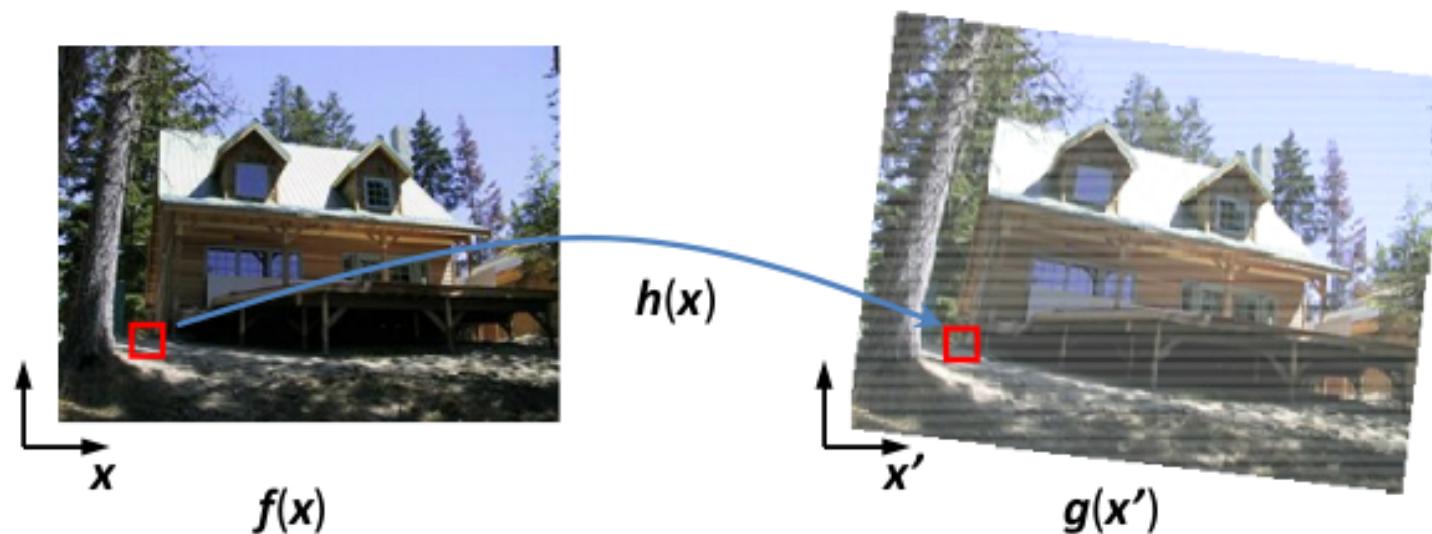
Image warping

- Given a coordinate transformation $x' = h(x)$ and a source image $f(x)$, how do we compute transformed image $g(x') = f(h(x))$?



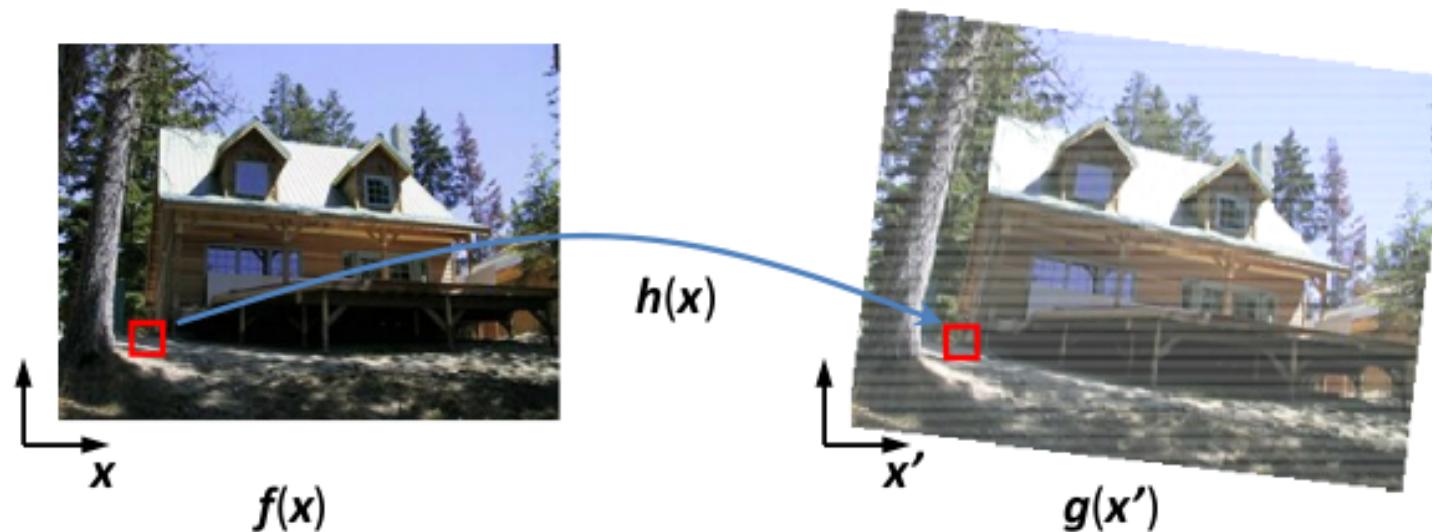
Forward warping

- Send each pixel in $f(x)$ to its corresponding location $x' = h(x)$ in $g(x')$



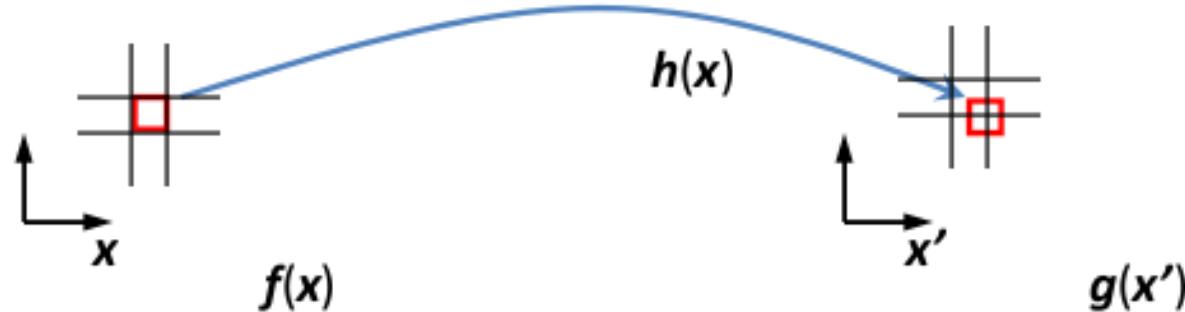
Forward warping

- ▶ Send each pixel in $f(x)$ to its corresponding location $x' = h(x)$ in $g(x')$
- ▶ What if a pixel “lands between” two pixels?



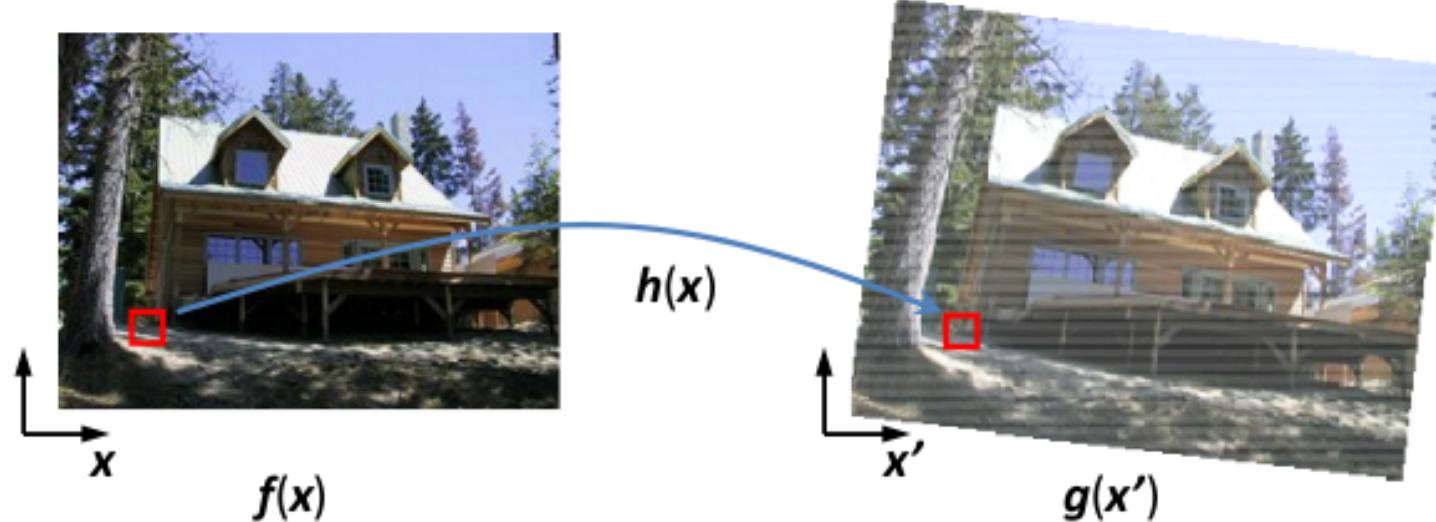
Forward warping

- ▶ Send each pixel in $f(x)$ to its corresponding location $x' = h(x)$ in $g(x')$
- ▶ What if a pixel “lands between” two pixels?
- ▶ Answer: add “contribution” to several pixels and normalize later (referred to as splatting)



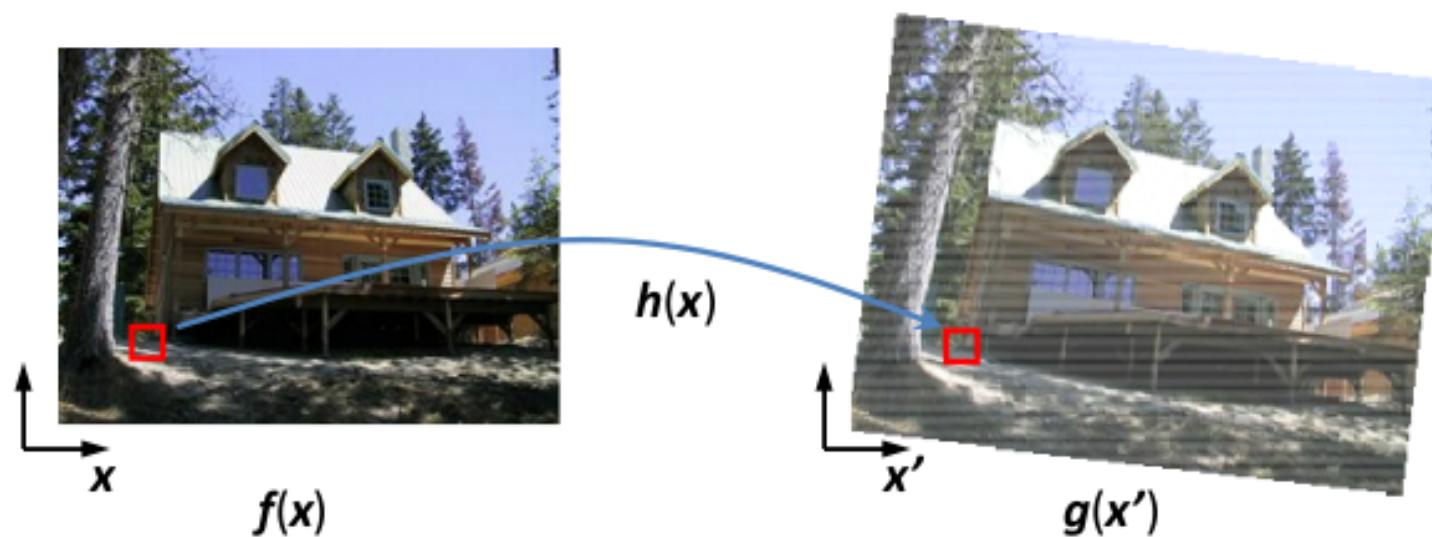
Inverse warping

- ▶ Get each pixel $g(x')$ from its corresponding location $x' = h(x)$ in $f(x)$



Inverse warping

- ▶ Get each pixel $g(x')$ from its corresponding location $x' = h(x)$ in $f(x)$
- ▶ What if a pixel “comes from between” two pixels?

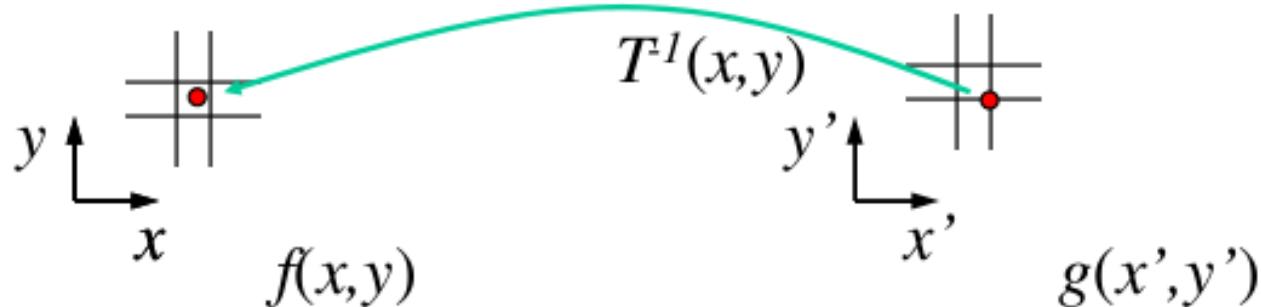


Inverse warping

- ▶ Get each pixel $g(x')$ from its corresponding location $x' = h(x)$ in $f(x)$
- ▶ What if a pixel “comes from between” two pixels?
- ▶ Answer: re-sample color value from *interpolated* (pre-filtered) source image

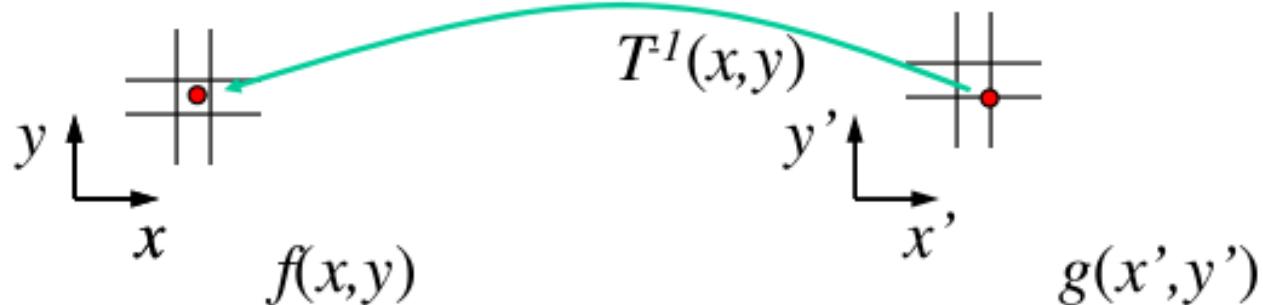


Inverse warping



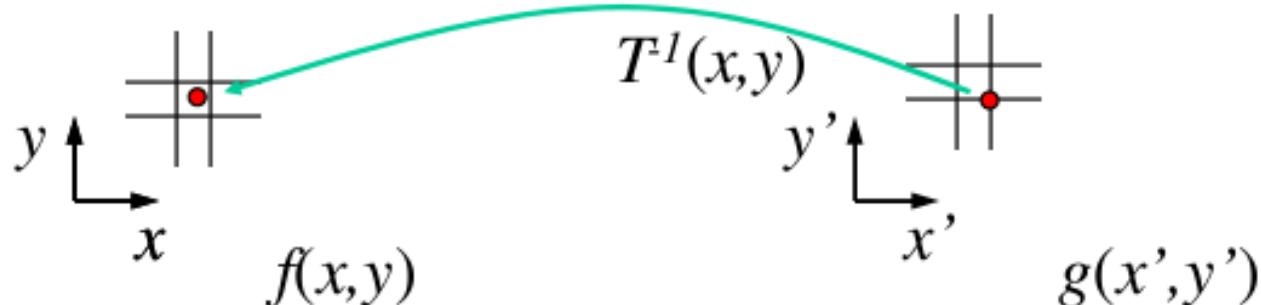
- ▶ Get each pixel $g(x', y')$ from its corresponding location
 - ▶ $(x, y) = T^{-1}(x', y')$ in the first image

Inverse warping



- ▶ Get each pixel $g(x', y')$ from its corresponding location
 - ▶ $(x, y) = T^{-1}(x', y')$ in the first image
- ▶ Q: What if a pixel “comes from between” two pixels?

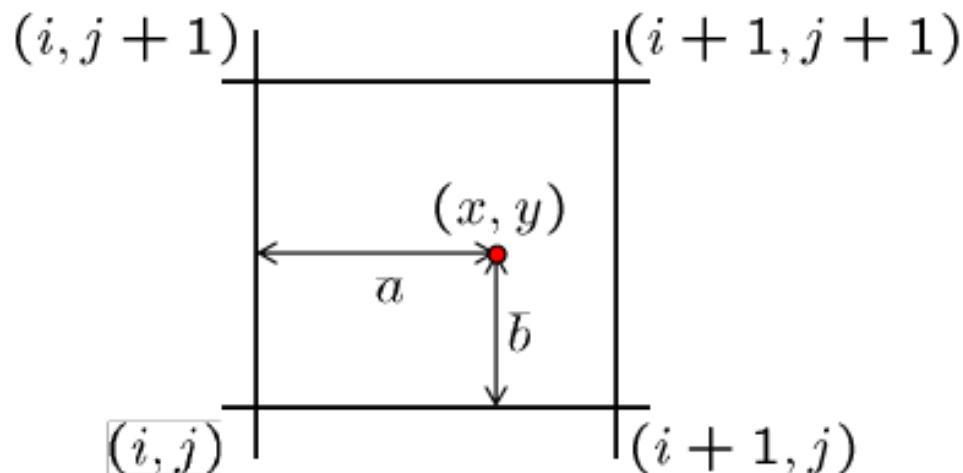
Inverse warping



- ▶ Get each pixel $g(x', y')$ from its corresponding location
 - ▶ $(x, y) = T^{-1}(x', y')$ in the first image
- ▶ Q: What if a pixel “comes from between” two pixels?
- ▶ A: Interpolate color values from neighbors (MATLAB: `help interp2`)
 - ▶ nearest neighbor, bilinear, bicubic, etc.

Bilinear interpolation

- ▶ Sampling at some $f(x, y)$



$$\begin{aligned} f(x, y) = & (1 - a)(1 - b) f[i, j] \\ & + a(1 - b) f[i + 1, j] \\ & + ab f[i + 1, j + 1] \\ & + (1 - a)b f[i, j + 1] \end{aligned}$$

Prefiltering

- ▶ Essential for *down-sampling* (decimation) to prevent *aliasing*
- ▶ MIP-mapping [Williams' 83] - similar to image pyramid but no down-sampling (transform to mipmap space)
 - ① Build a pyramid using a decimation filter:
 - ▶ block averaging
 - ▶ Burt & Adelson (5-tap bimodal)
 - ▶ 7-tap wavelet-based filter (best)
 - ② *trilinear* interpolation
 - ▶ bilinear within each 2 adjacent levels
 - ▶ linear blend *between* levels (based on pixel size)

Look at Forsyth & Ponce for details

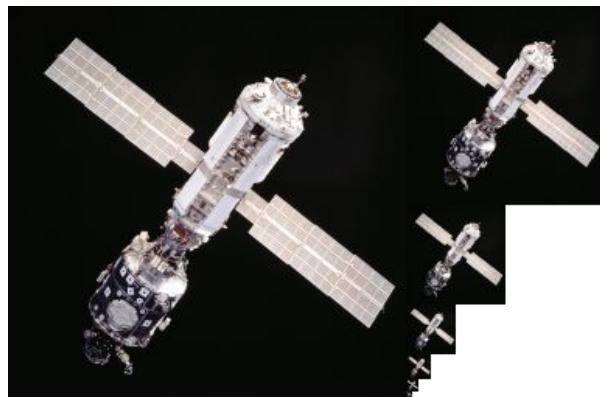
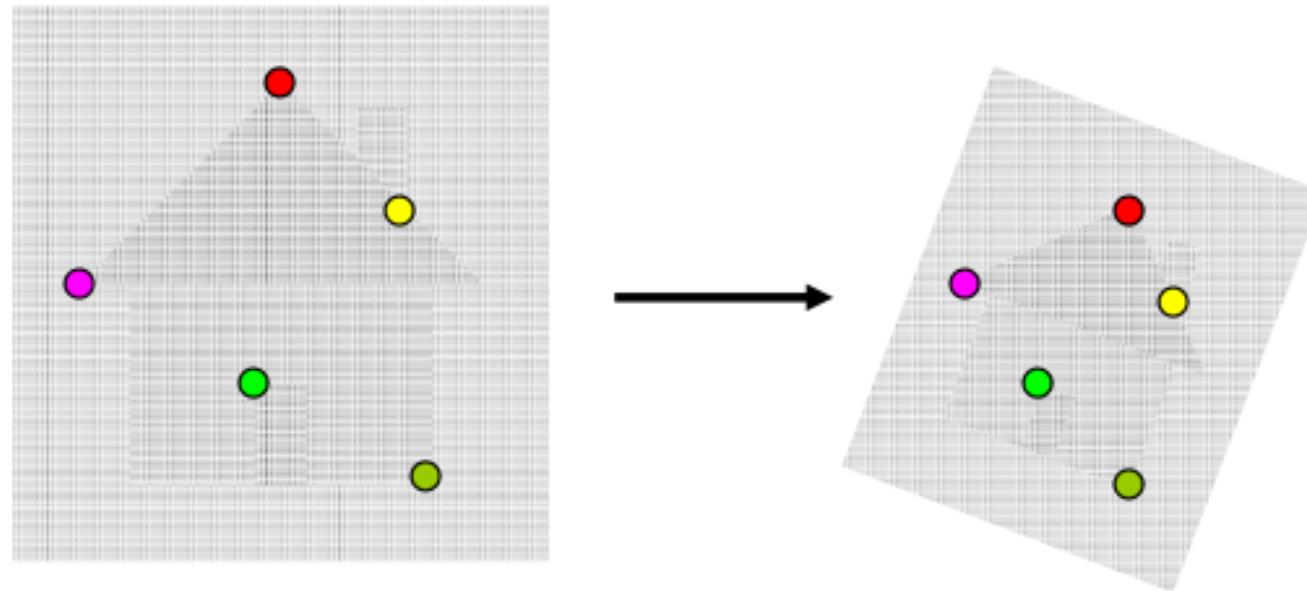
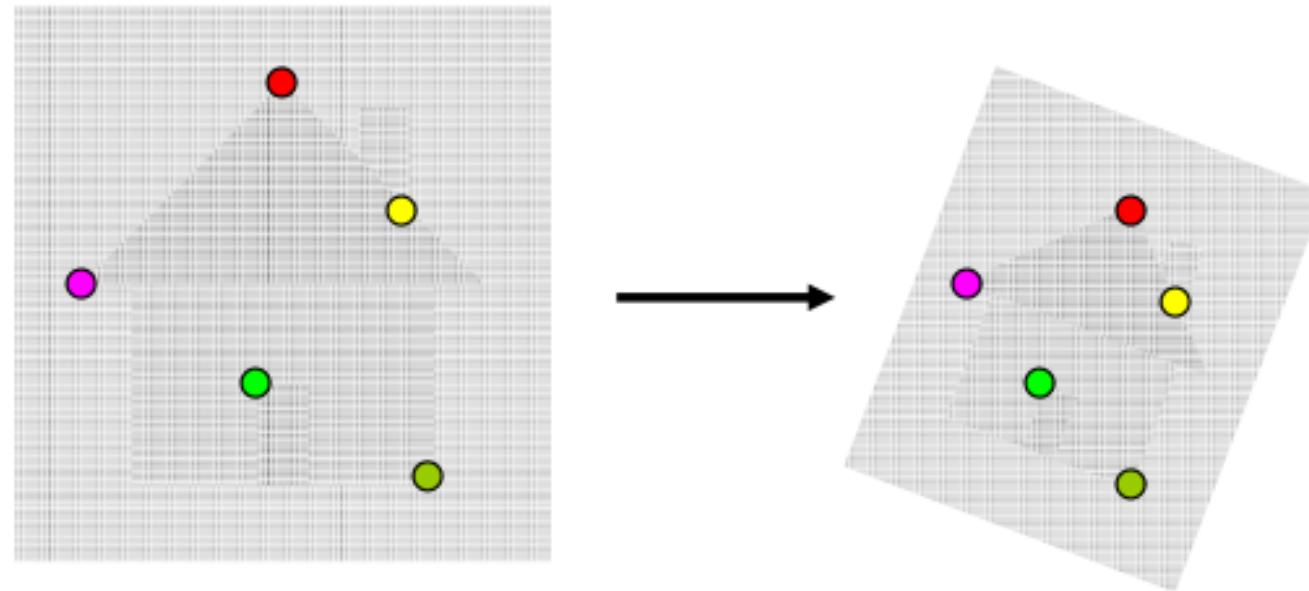


Image alignment



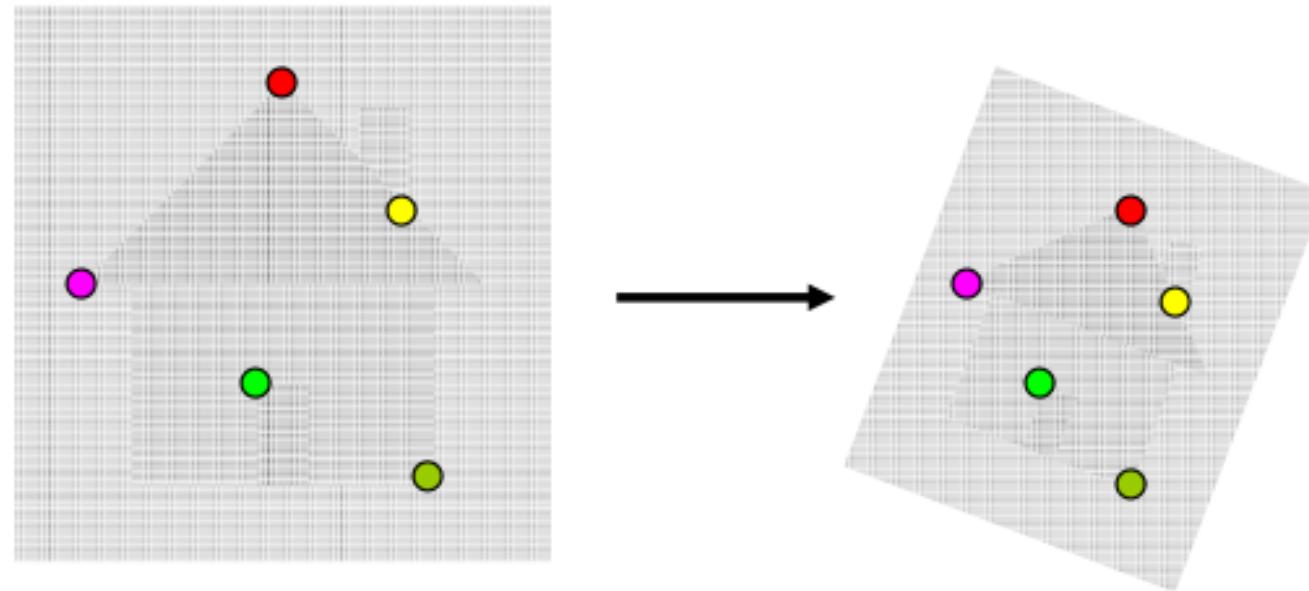
- ▶ Two broad approaches:
 - ▶ Direct (pixel-based) alignment:
 - ▶ Search for alignment where the most pixels agree

Image alignment



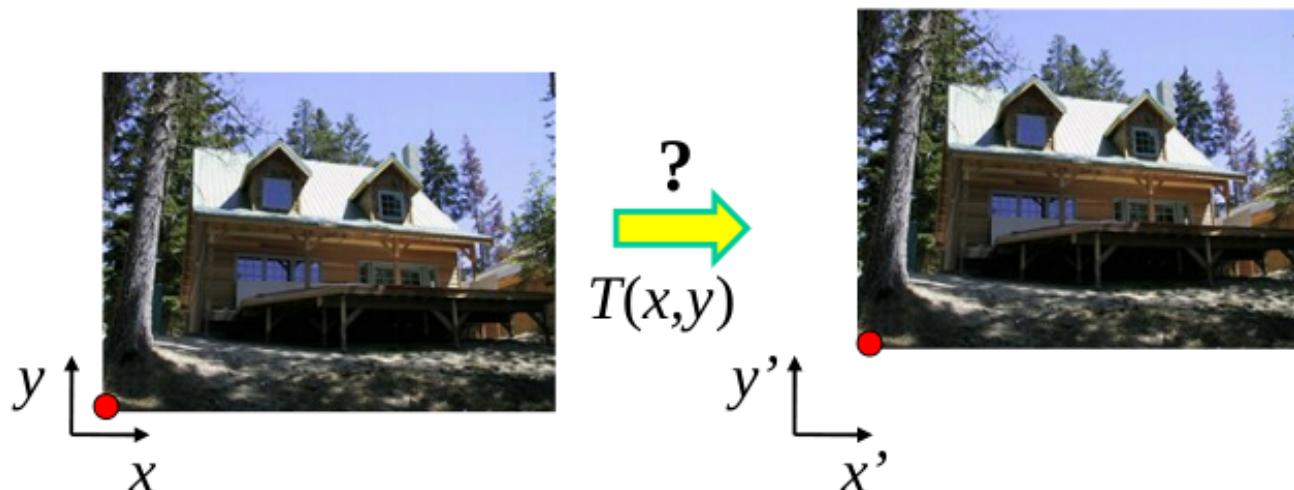
- ▶ Two broad approaches:
 - ▶ Direct (pixel-based) alignment:
 - ▶ Search for alignment where the most pixels agree
 - ▶ Feature-based alignment:
 - ▶ Search for alignment where *extracted features* agree
 - ▶ Can be verified using pixel-based alignment

Image alignment



- ▶ Two broad approaches:
 - ▶ Direct (pixel-based) alignment:
 - ▶ Search for alignment where the most pixels agree
 - ▶ Feature-based alignment:
 - ▶ Search for alignment where *extracted features* agree
 - ▶ Can be verified using pixel-based alignment
 - ▶ **How many features do we need?**

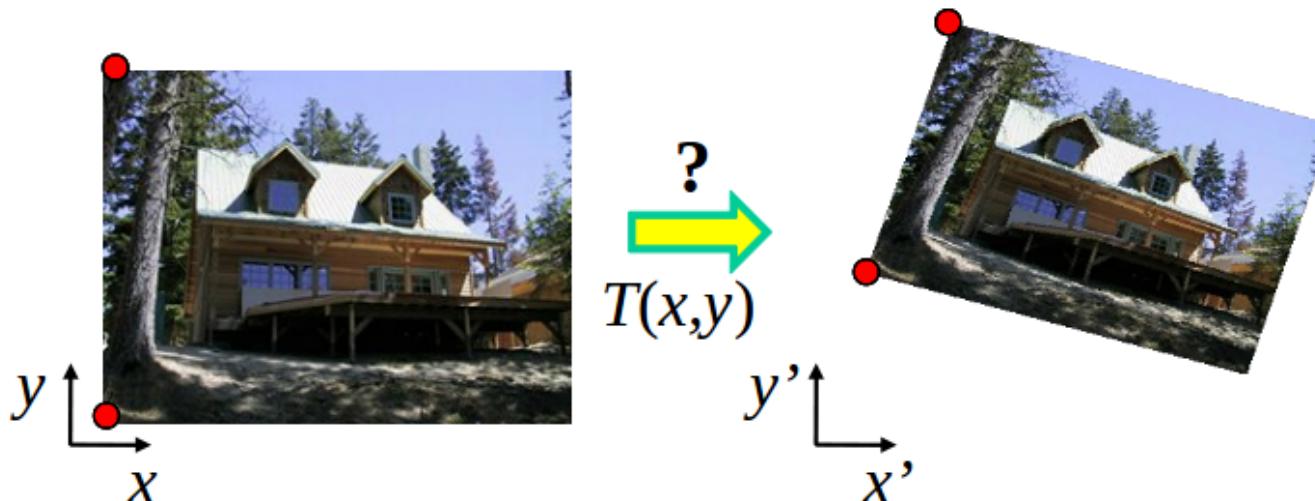
Translation: # correspondences?



- ▶ How many correspondences needed for translation?
- ▶ How many degrees of freedom?
- ▶ What is the transformation matrix?

$$T = \begin{bmatrix} 1 & 0 & x' - x \\ 0 & 1 & y' - y \\ 0 & 0 & 1 \end{bmatrix}$$

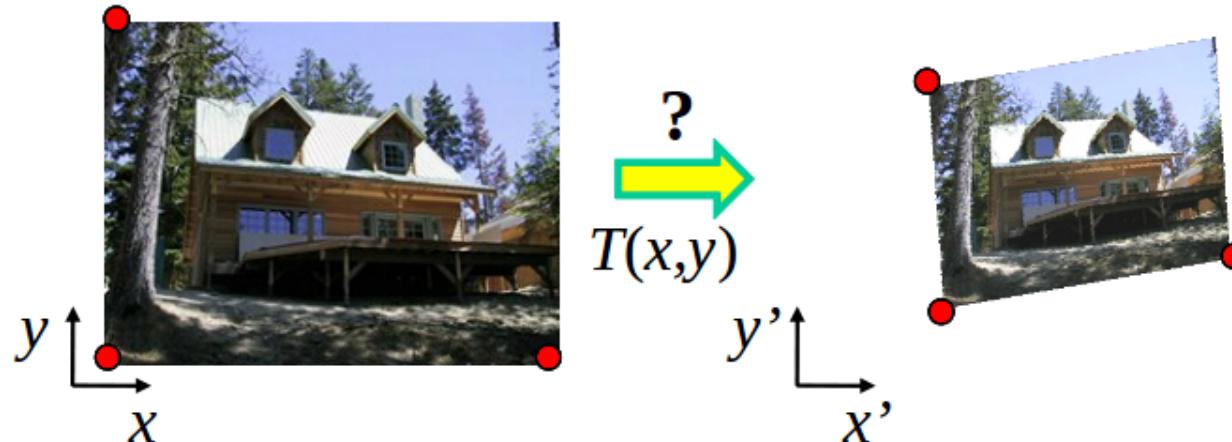
Euclidean: # correspondences?



- ▶ How many correspondences needed for translation + rotation?
- ▶ How many degrees of freedom?
- ▶ What is the transformation matrix?

$$T = \begin{bmatrix} \cos(\theta) & \sin(\theta) & x' - x \\ -\sin(\theta) & \cos(\theta) & y' - y \\ 0 & 0 & 1 \end{bmatrix}$$

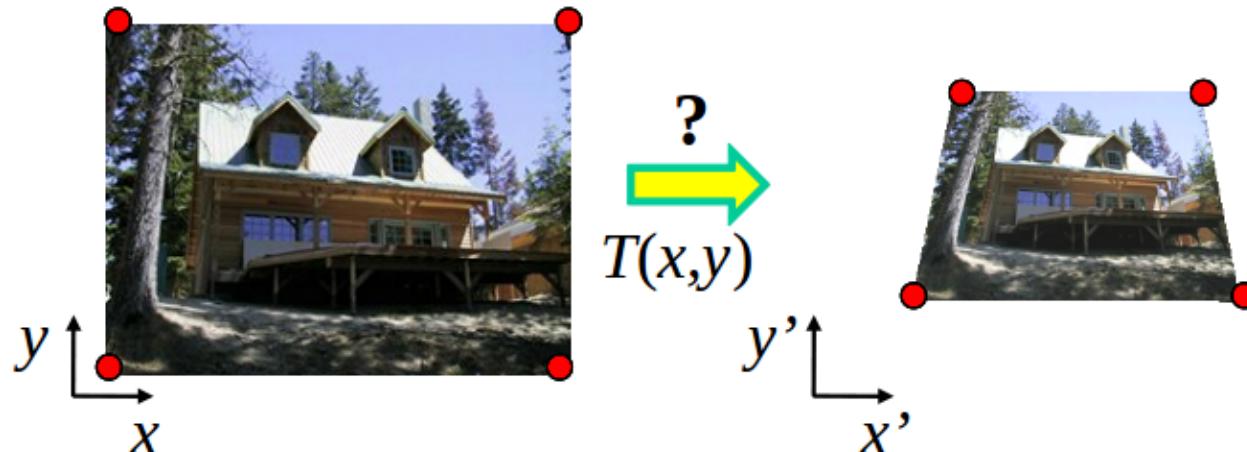
Affine: # correspondences?



- ▶ How many correspondences needed for affine?
- ▶ How many degrees of freedom?
- ▶ What is the transformation matrix?

$$T = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix}$$

Projective: # correspondences?



- ▶ How many correspondences needed for translation?
- ▶ How many degrees of freedom?
- ▶ What is the transformation matrix?

$$T = \begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix}$$

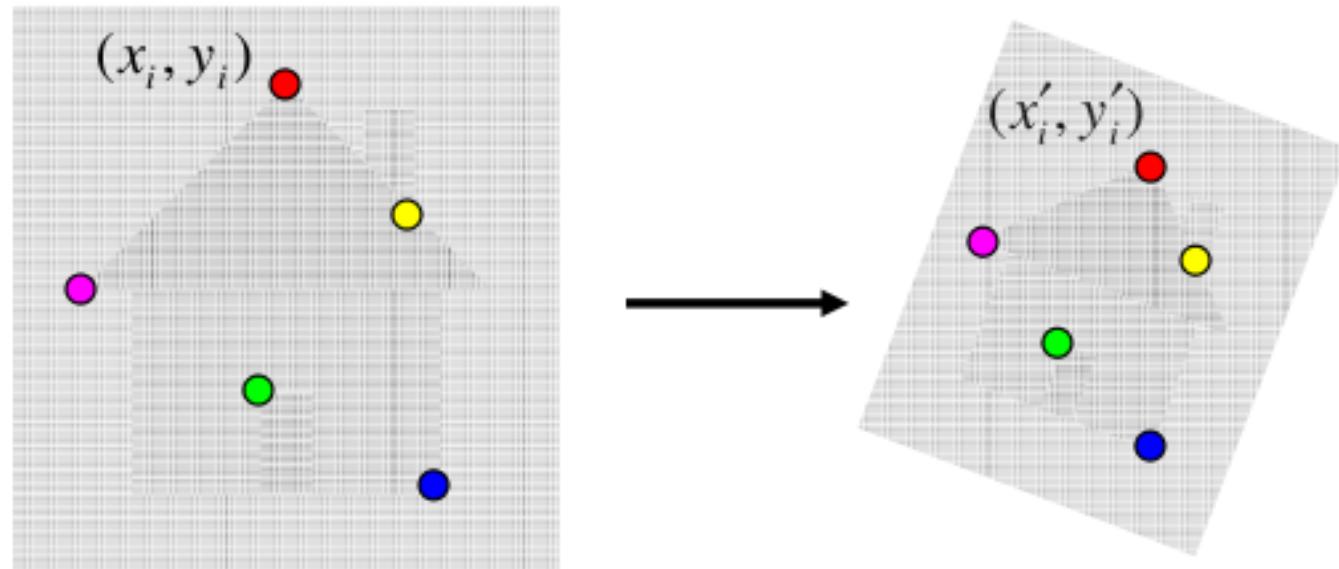
Fitting an affine transformation



- ▶ Translation and Euclidean are trivial to compute (you should start with these for practice!)
- ▶ Let's look at affine models: **approximates** perspective projection of planar objects.

Fitting an affine transformation

- ▶ Assuming we know the correspondences, how do we get the transformation?



$$\begin{bmatrix} x'_i \\ y'_i \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix}$$

Fitting an affine transformation

- ▶ How the heck do we solve for this guy (standard Gaussian elimination)???

Fitting an affine transformation

- ▶ How the heck do we solve for this guy (standard Gaussian elimination)???

$$\begin{bmatrix} x'_i \\ y'_i \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix} \implies$$

$$\begin{bmatrix} x_i & y_i & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & x_i & y_i & 1 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \end{bmatrix} \begin{bmatrix} a \\ b \\ c \\ d \\ e \\ f \end{bmatrix} = \begin{bmatrix} x'_i \\ y'_i \\ \vdots \\ \vdots \\ \vdots \\ \vdots \end{bmatrix}$$

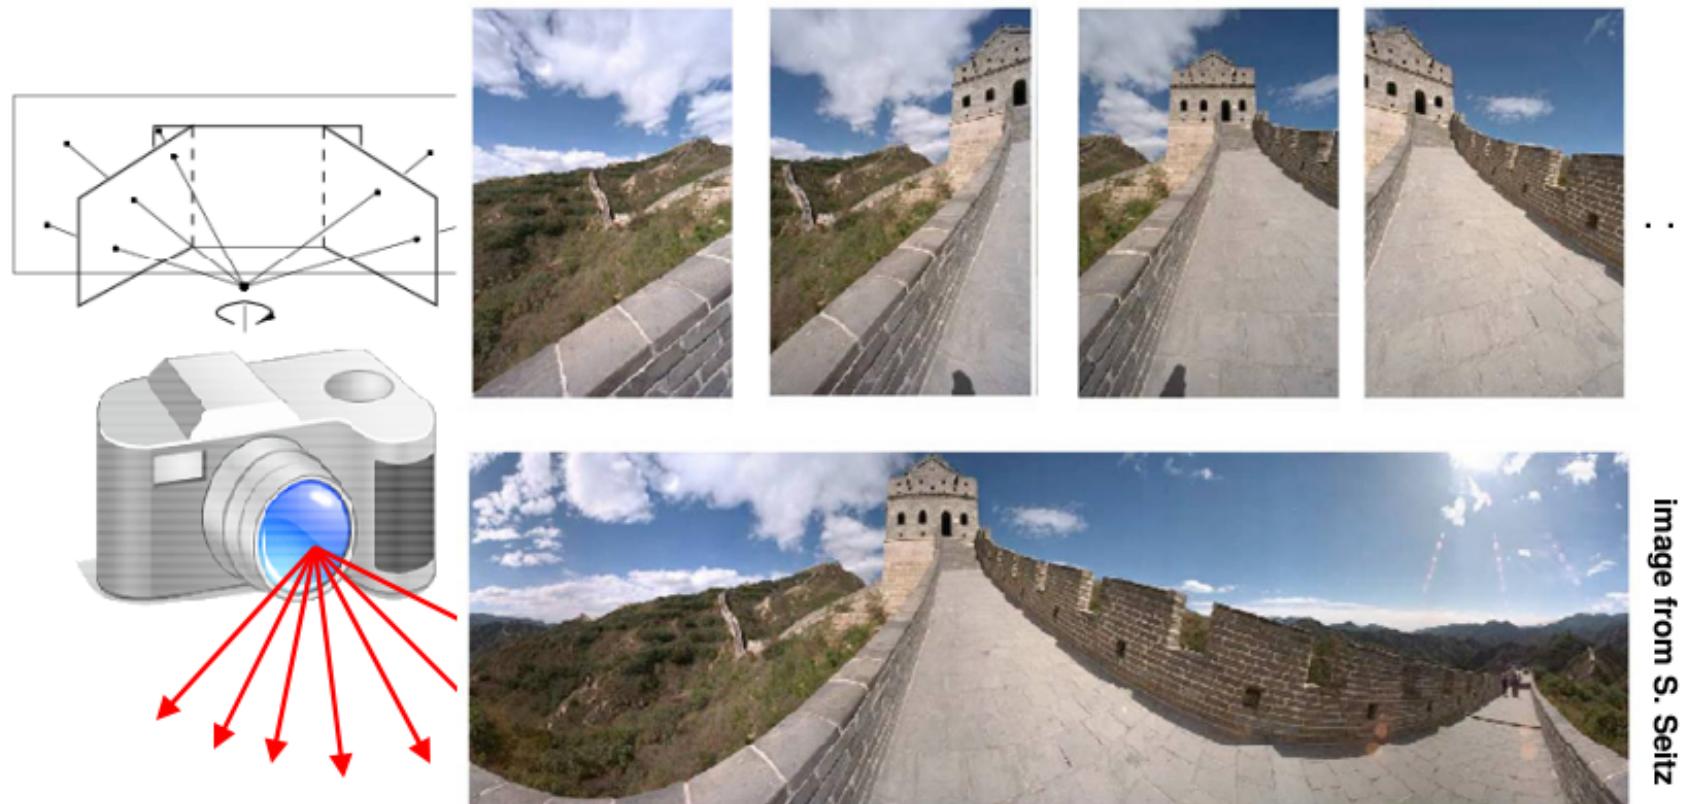
Fitting an affine transformation

- ▶ How the heck do we solve for this guy (standard Gaussian elimination)???
- ▶ How many corresponding points do we need?

$$\begin{bmatrix} x'_i \\ y'_i \\ 1 \end{bmatrix} = \begin{bmatrix} a & b & c \\ d & e & f \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_i \\ y_i \\ 1 \end{bmatrix} \implies$$

$$\begin{bmatrix} x_i & y_i & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & x_i & y_i & 1 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \vdots \end{bmatrix} \begin{bmatrix} a \\ b \\ c \\ d \\ e \\ f \end{bmatrix} = \begin{bmatrix} x'_i \\ y'_i \\ \vdots \\ \vdots \\ \vdots \\ \vdots \end{bmatrix}$$

Creating panoramas



- ▶ General idea: Given a few images of the same scene, create a wider view by combining.

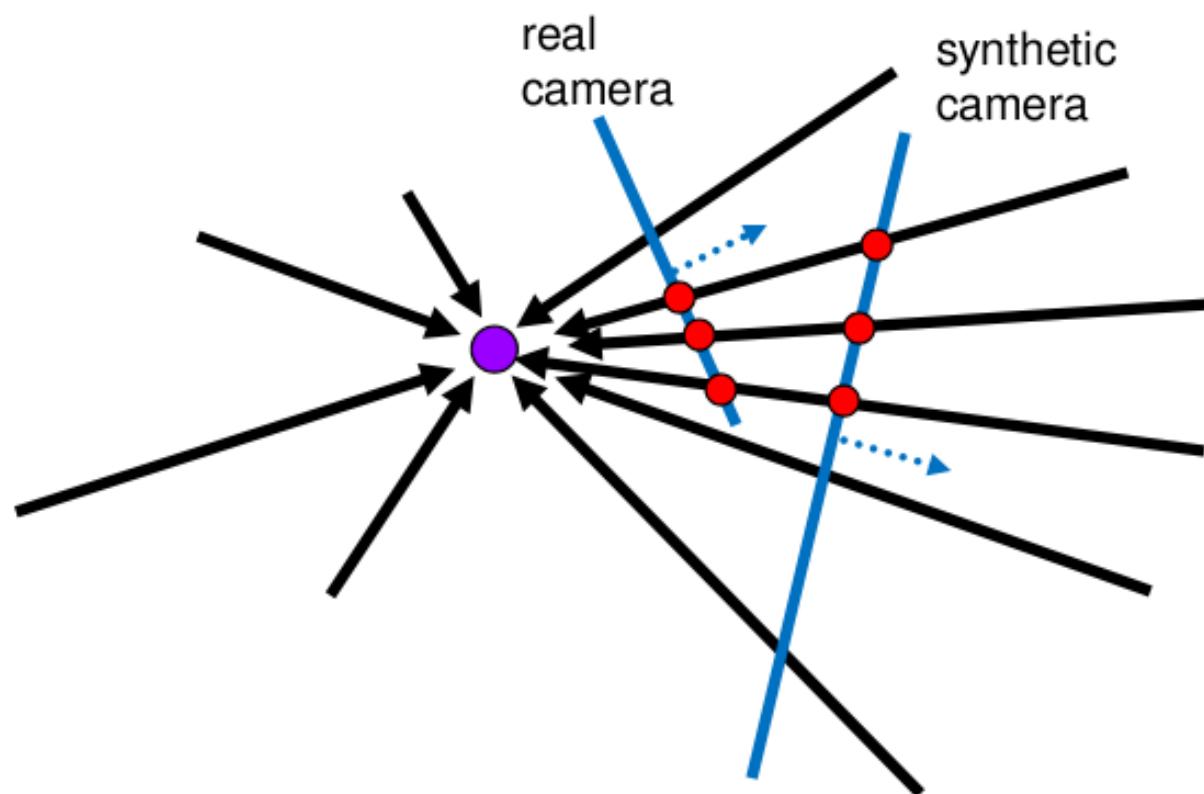
How?

- ▶ Basic procedure:
 - ▶ Take a sequence of images from the same position
 - ▶ Rotate the camera about its optical center
 - ▶ Compute the transformation between the second image and the first
 - ▶ Transform (warp) the second image to overlap with the first
 - ▶ Blend the two together
 - ▶ If there are more than a single image, repeat!

How?

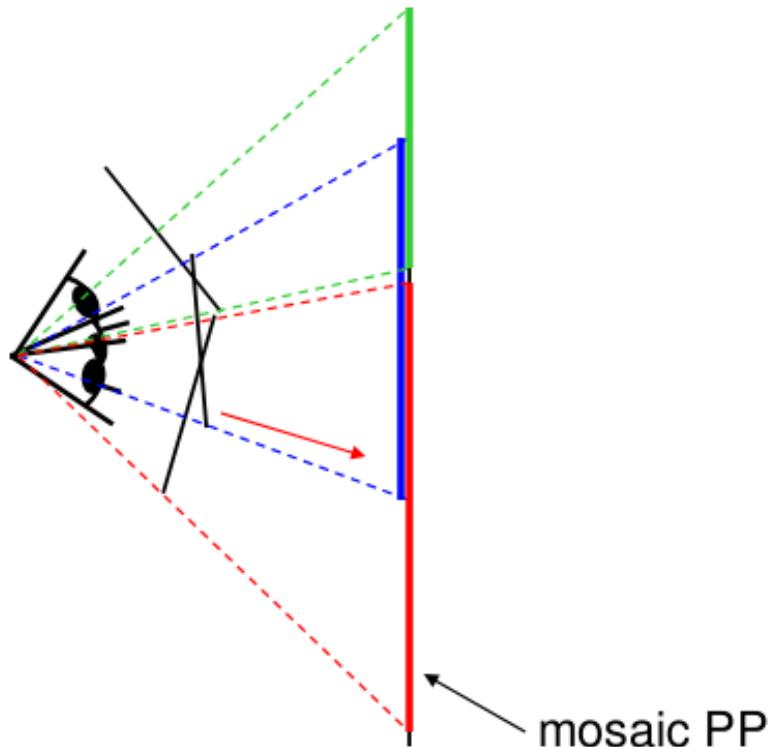
- ▶ Basic procedure:
 - ▶ Take a sequence of images from the same position
 - ▶ Rotate the camera about its optical center
 - ▶ Compute the transformation between the second image and the first
 - ▶ Transform (warp) the second image to overlap with the first
 - ▶ Blend the two together
 - ▶ If there are more than a single image, repeat!
- ▶ ...**hold on a minute!** why does this work?
 - ▶ What about the 3-D geometry of the scene?
 - ▶ Why aren't we using it?

Generating synthetic views



- ▶ We can generate *any* synthetic camera view as long as it has the **same center of projection!**

Image reprojection

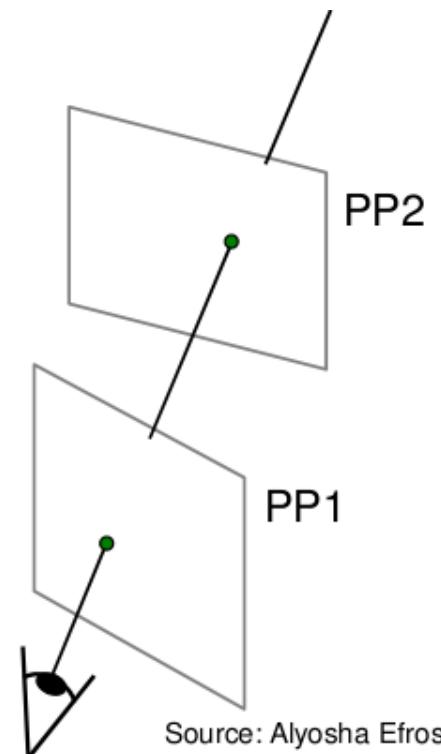


- ▶ The mosaic has a natural interpretation in 3-D
 - ▶ The images are reprojected onto a common plane
 - ▶ Form the mosaic on this new plane
 - ▶ This implies that the mosaic is really just a *synthetic wide-angle camera*

Homography

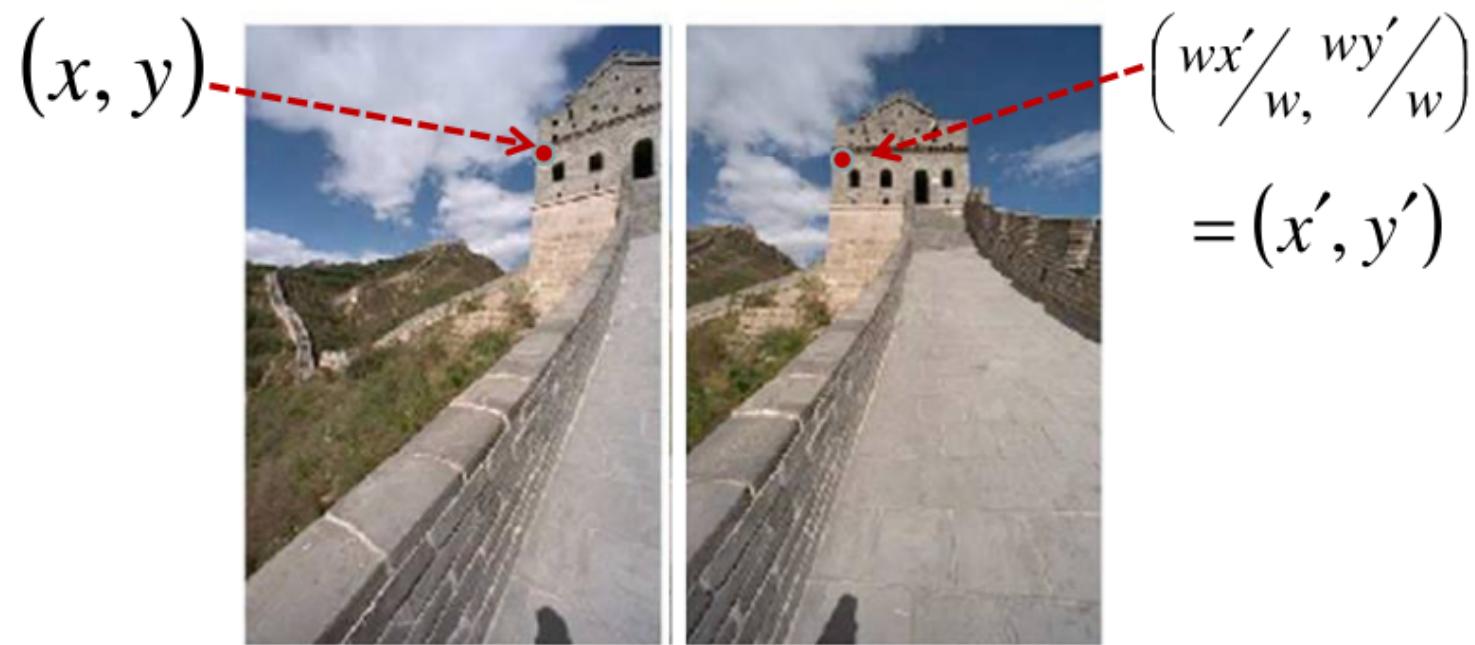
- ▶ How do we relate two images from the same camera center?
 - ▶ i.e., how do we map a pixel from PP1 to PP2?
- ▶ We need to think of this as a 2-D image warp from one image to another
- ▶ A projective transformation is a mapping between any two PPs with the same center of projection!
 - ▶ A rectangle maps to an arbitrary quadrilateral
 - ▶ Parallel lines do not remain parallel
 - ▶ However straight lines are preserved!
- ▶ We refer to this transformation as a **homography**

$$\underbrace{\begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix}}_{\mathbf{p}'} = \underbrace{\begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix}}_H \underbrace{\begin{bmatrix} x \\ y \\ 1 \end{bmatrix}}_{\mathbf{p}}$$



Source: Alyosha Efros

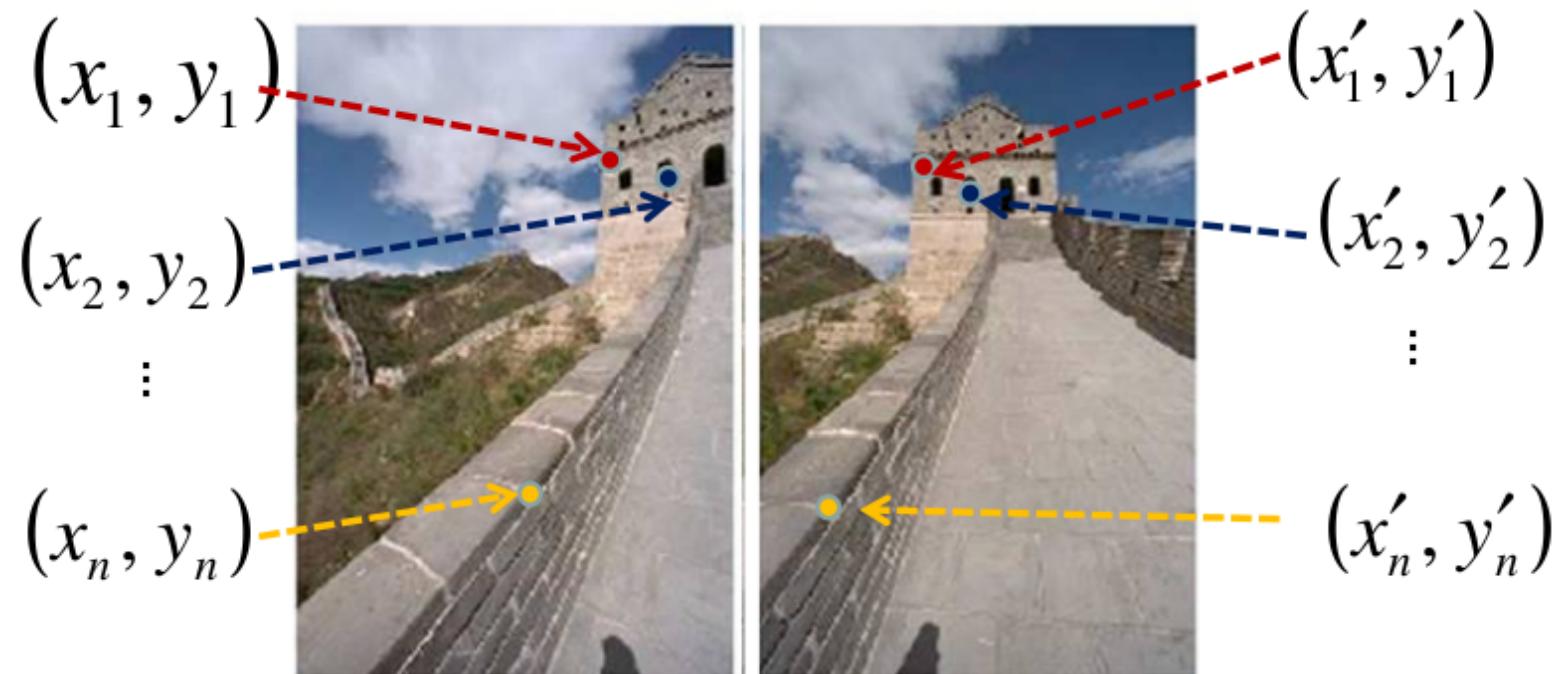
Homography



- ▶ To apply our homography H
 - ▶ Compute $\mathbf{p}' = H\mathbf{p}$
 - ▶ Convert \mathbf{p}' from homogeneous coords to image coords (normalize)

$$\underbrace{\begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix}}_{\mathbf{p}'} = \underbrace{\begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix}}_H \underbrace{\begin{bmatrix} x \\ y \\ 1 \end{bmatrix}}_{\mathbf{p}}$$

Homography



- ▶ To **compute** our homography H given pairs of corresponding feature points n two images, **similar to our affine transformation**, we need to set up an equation where the parameters of H are the unknowns

How??? - Direct Linear Transform (DLT)

- Recall, we're trying to solve: $\mathbf{p}' = H\mathbf{p}$ where

$$\underbrace{\begin{bmatrix} wx' \\ wy' \\ w \end{bmatrix}}_{\mathbf{p}'} = \underbrace{\begin{bmatrix} a & b & c \\ d & e & f \\ g & h & i \end{bmatrix}}_H \underbrace{\begin{bmatrix} x \\ y \\ 1 \end{bmatrix}}_{\mathbf{p}}$$

- Important note:** These are homogeneous vectors! $\therefore \mathbf{p}' \neq H\mathbf{p}$
 - They have the same direction but differ in magnitude by a non-zero scale factor!
- We also note that $\mathbf{p}' = H\mathbf{p}$ can be written as $\mathbf{p}' \times H\mathbf{p} = 0$ (**Proof? Trivial and left to you!**)

How??? - Direct Linear Transform (DLT)

- ▶ Denote the j^{th} -row of H by \mathbf{h}^j , then for the i^{th} point correspondence we get

$$H\mathbf{p}_i = \begin{bmatrix} \mathbf{h}^1\mathbf{p}_i \\ \mathbf{h}^2\mathbf{p}_i \\ \mathbf{h}^3\mathbf{p}_i \end{bmatrix}$$

Letting $\mathbf{p}' = [x'_i, y'_i, w'_i]^T$, the cross product becomes

$$\mathbf{p}'_i \times H\mathbf{p}_i = \begin{bmatrix} y'_i \mathbf{h}^3\mathbf{p}_i - w'_i \mathbf{h}^2\mathbf{p}_i \\ w'_i \mathbf{h}^1\mathbf{p}_i - x'_i \mathbf{h}^3\mathbf{p}_i \\ x'_i \mathbf{h}^2\mathbf{p}_i - y'_i \mathbf{h}^1\mathbf{p}_i \end{bmatrix}$$

Note: $\mathbf{h}^j\mathbf{p}_i = [\mathbf{p}_i]^T[\mathbf{h}^j]^T$ (i.e. this is a dot-product!)

How??? - Direct Linear Transform (DLT)

- ▶ Therefore, for the i^{th} point correspondence, we can re-write this as:

$$\begin{bmatrix} \mathbf{0}^T & -w'_i \mathbf{p}_i^T & y'_i \mathbf{p}_i^T \\ w'_i \mathbf{p}_i^T & \mathbf{0}^T & -x'_i \mathbf{p}_i^T \\ -y'_i \mathbf{p}_i^T & x'_i \mathbf{p}_i^T & \mathbf{0}^T \end{bmatrix} \begin{bmatrix} (\mathbf{h}^1)^T \\ (\mathbf{h}^2)^T \\ (\mathbf{h}^3)^T \end{bmatrix} = \mathbf{0}$$

Notice however that row three is a linear combination of rows one and two! (multiply row 1 by $-\frac{x'_i}{w'_i}$ and row two by $-\frac{y'_i}{w'_i}$ and add them to get row three)

- ▶ Therefore, w.l.o.g., we can take the first two rows to get:

$$\underbrace{\begin{bmatrix} \mathbf{0}^T & -w'_i \mathbf{p}_i^T & y'_i \mathbf{p}_i^T \\ w'_i \mathbf{p}_i^T & \mathbf{0}^T & -x'_i \mathbf{p}_i^T \end{bmatrix}}_{A_i} \underbrace{\begin{bmatrix} (\mathbf{h}^1)^T \\ (\mathbf{h}^2)^T \\ (\mathbf{h}^3)^T \end{bmatrix}}_{\mathbf{h}} = \mathbf{0}$$

which is of the form $A_i \mathbf{h} = \mathbf{0}$ where $A_i \in \mathbb{R}^{2 \times 9}$ and $\mathbf{h} \in \mathbb{R}^{9 \times 1}$ contains the entries of H

How??? - Direct Linear Transform (DLT)

- ▶ Because we are looking for a projective transformation, we need a minimum of 4-corresponding points, i.e., $i = 1, 2, 3, 4$
- ▶ This will yield some $A\mathbf{h} = 0$ where $A \in \mathbb{R}^{8 \times 9}$ and again $\mathbf{h} \in \mathbb{R}^{9 \times 1}$
- ▶ Question: What is the rank of A ?

How??? - Direct Linear Transform (DLT)

- ▶ Because we are looking for a projective transformation, we need a minimum of 4-corresponding points, i.e., $i = 1, 2, 3, 4$
- ▶ This will yield some $A\mathbf{h} = 0$ where $A \in \mathbb{R}^{8 \times 9}$ and again $\mathbf{h} \in \mathbb{R}^{9 \times 1}$
- ▶ Question: What is the rank of A ?
- ▶ Answer: At most A has rank 8!

How??? - Direct Linear Transform (DLT)

- ▶ Because we are looking for a projective transformation, we need a minimum of 4-corresponding points, i.e., $i = 1, 2, 3, 4$
- ▶ This will yield some $A\mathbf{h} = 0$ where $A \in \mathbb{R}^{8 \times 9}$ and again $\mathbf{h} \in \mathbb{R}^{9 \times 1}$
- ▶ Question: What is the rank of A ?
- ▶ Answer: At most A has rank 8!
- ▶ Question: What does such an \mathbf{h} represent?

How??? - Direct Linear Transform (DLT)

- ▶ Because we are looking for a projective transformation, we need a minimum of 4-corresponding points, i.e., $i = 1, 2, 3, 4$
- ▶ This will yield some $A\mathbf{h} = 0$ where $A \in \mathbb{R}^{8 \times 9}$ and again $\mathbf{h} \in \mathbb{R}^{9 \times 1}$
- ▶ Question: What is the rank of A ?
- ▶ Answer: At most A has rank 8!
- ▶ Question: What does such an \mathbf{h} represent?
- ▶ Answer: A must have a 1-dimensional NULL space that is spanned by \mathbf{h}

How??? - Direct Linear Transform (DLT)

- ▶ Because we are looking for a projective transformation, we need a minimum of 4-corresponding points, i.e., $i = 1, 2, 3, 4$
- ▶ This will yield some $A\mathbf{h} = 0$ where $A \in \mathbb{R}^{8 \times 9}$ and again $\mathbf{h} \in \mathbb{R}^{9 \times 1}$
- ▶ Question: What is the rank of A ?
- ▶ Answer: At most A has rank 8!
- ▶ Question: What does such an \mathbf{h} represent?
- ▶ Answer: A must have a 1-dimensional NULL space that is spanned by \mathbf{h}
- ▶ How on earth do we find a non-trivial vector such that $A\mathbf{h} = 0$?

How??? - Direct Linear Transform (DLT)

- ▶ Because we are looking for a projective transformation, we need a minimum of 4-corresponding points, i.e., $i = 1, 2, 3, 4$
- ▶ This will yield some $A\mathbf{h} = 0$ where $A \in \mathbb{R}^{8 \times 9}$ and again $\mathbf{h} \in \mathbb{R}^{9 \times 1}$
- ▶ Question: What is the rank of A ?
- ▶ Answer: At most A has rank 8!
- ▶ Question: What does such an \mathbf{h} represent?
- ▶ Answer: A must have a 1-dimensional NULL space that is spanned by \mathbf{h}
- ▶ How on earth do we find a non-trivial vector such that $A\mathbf{h} = 0$?
- ▶ Use the SVD!

Recall: singular value decomposition

- Given a matrix $A \in \mathbb{R}^{8 \times 9}$, the SVD is

$$A = U\Sigma V^T$$

where $UU^T = I \in \mathbb{R}^{8 \times 8}$, $VV^T = I \in \mathbb{R}^{9 \times 9}$, and $\Sigma \in \mathbb{R}^{8 \times 9} = \text{diag}(\sigma_1, \sigma_2, \dots, \sigma_8)$ with $\sigma_9 = 0$

- Because U and V are unitary, we can write

$$AV = U\Sigma$$

and letting $U = [\mathbf{u}_1, \mathbf{u}_2, \dots]$, $V = [\mathbf{v}_1, \mathbf{v}_2, \dots]$ we get

$$A\mathbf{v}_i = \mathbf{u}_i\sigma_i = \sigma_i\mathbf{u}_i$$

Recall: singular value decomposition

- ▶ So what happens for $i = 9$?

$$A\mathbf{v}_9 = \sigma_9 \mathbf{u}_9$$

- but there is no \mathbf{u}_9 and $\sigma_9 = 0$, so we must get that $A\mathbf{v}_9 = 0$ (try it!)
- ▶ What does this tell us?

Recall: singular value decomposition

- ▶ So what happens for $i = 9$?

$$A\mathbf{v}_9 = \sigma_9 \mathbf{u}_9$$

- but there is no \mathbf{u}_9 and $\sigma_9 = 0$, so we must get that $A\mathbf{v}_9 = 0$ (try it!)
- ▶ What does this tell us?
- ▶ We have found our NULL space vector!

Recall: singular value decomposition

- ▶ So what happens for $i = 9$?

$$A\mathbf{v}_9 = \sigma_9 \mathbf{u}_9$$

- but there is no \mathbf{u}_9 and $\sigma_9 = 0$, so we must get that $A\mathbf{v}_9 = 0$ (try it!)
- ▶ What does this tell us?
- ▶ We have found our NULL space vector!
- ▶ If $\mathbf{v}_9 = \mathbf{h} = [h_1, h_2, h_3, \dots, h_9]^T$, then

$$H = \begin{bmatrix} h_1 & h_2 & h_3 \\ h_4 & h_5 & h_6 \\ h_7 & h_8 & h_9 \end{bmatrix}$$

Algorithmically - numerically unstable!

- ▶ Assume we have at least four matched points: How do we compute the homography H ?
- ▶ Using the above formulation can behave poorly due to scaling (homogeneous coords. $\Rightarrow \mathbf{p} = ([750, 250, 1]^T)^2$)
- ▶ Instead, we should normalize the coordinates (we can always use a similarity transformation to find H !)
- ▶ Compute the Normalized DLT:
 - ① Translate each coordinate to have zero mean
 - ② Scale each point so the average distance to the origin is $\sqrt{2}$
 - ③ Both of these are done using a similarity transformation:

$$\tilde{\mathbf{p}} = T\mathbf{p} \quad \tilde{\mathbf{p}}' = T'\mathbf{p}'$$

- ④ Compute \tilde{H} using the standard DLT in normalized coordinates
- ⑤ Unnormalize: $H = T'^{-1}\tilde{H}T \rightarrow \mathbf{p}' = T'^{-1}\tilde{H}T\mathbf{p} = H\mathbf{p}$

Similarity transformation

- Great, so how do we find such a similarity transformation?

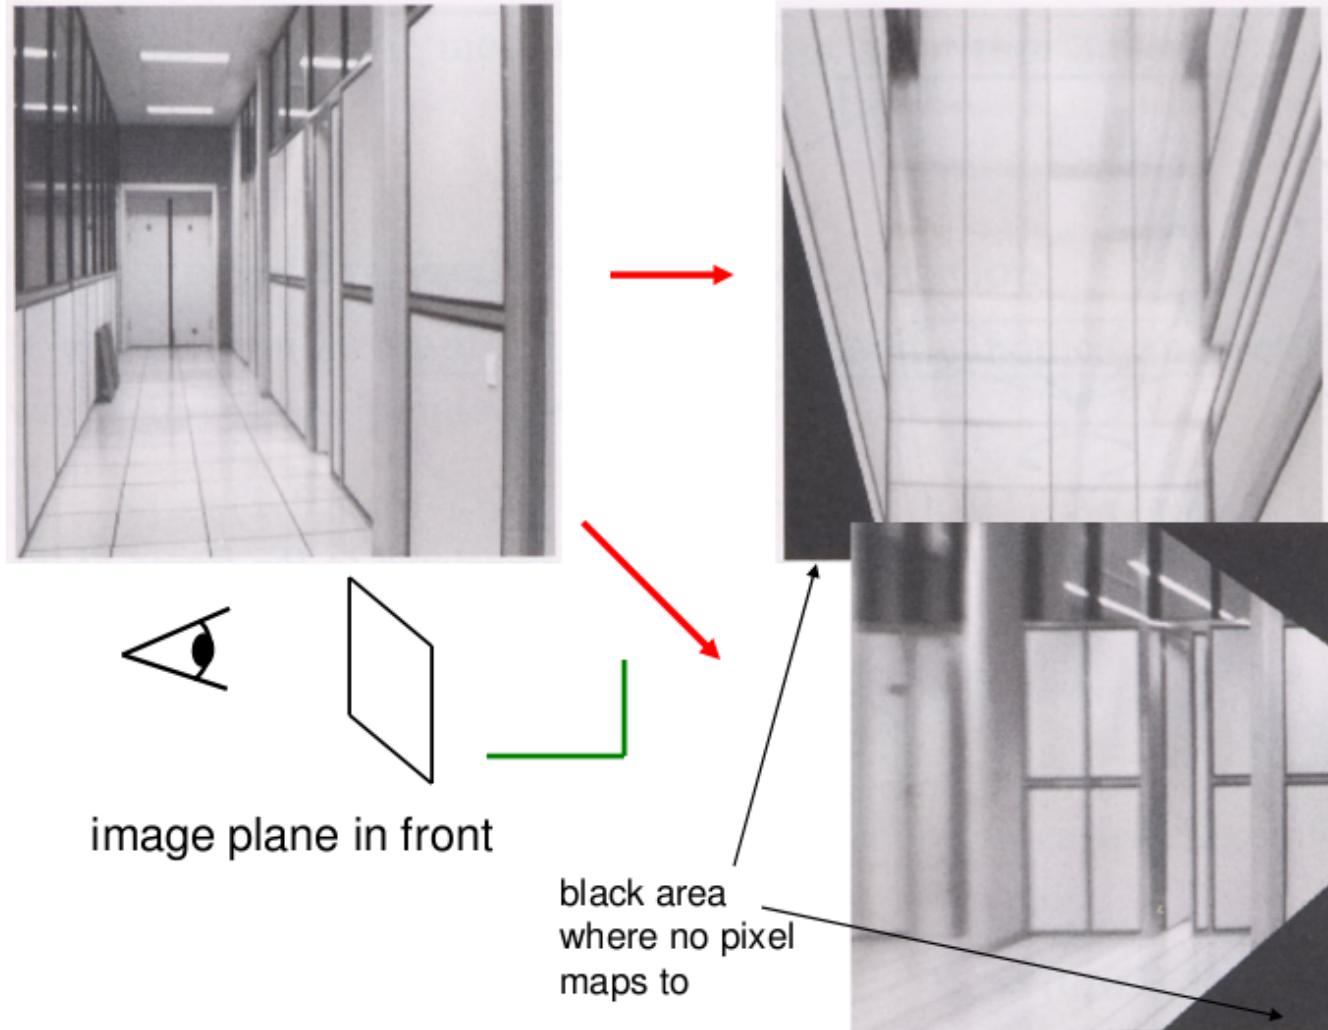
$$T = \begin{bmatrix} \alpha & 0 & -\alpha\mu_x \\ 0 & \alpha & -\alpha\mu_y \\ 0 & 0 & 1 \end{bmatrix}$$

- ① Compute the mean of your points \mathbf{p} and denote them by μ_x, μ_y
- ② Compute the scale factor such that the average distance of each point is $\sqrt{2}$ from the origin

$$\frac{1}{n} \sum_{i=1}^n \|\mathbf{p}_i - \mu\| = \sqrt{2} \implies \alpha = \frac{n\sqrt{2}}{\sum_{i=1}^n \|\mathbf{p}_i - \mu\|}$$

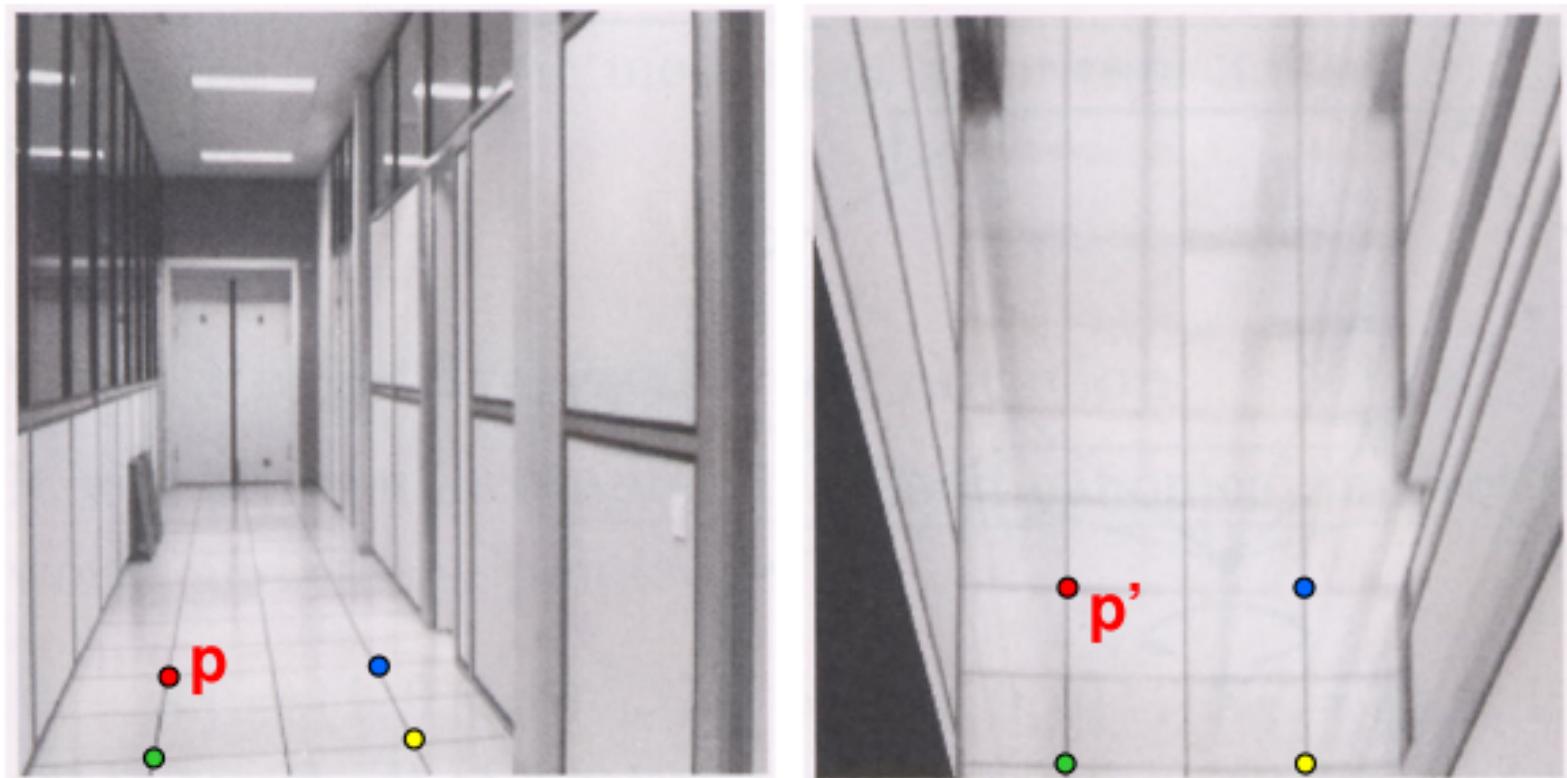
- ③ Only on the 2-D points - inhomogeneous!

Examples - Image warping with homographies



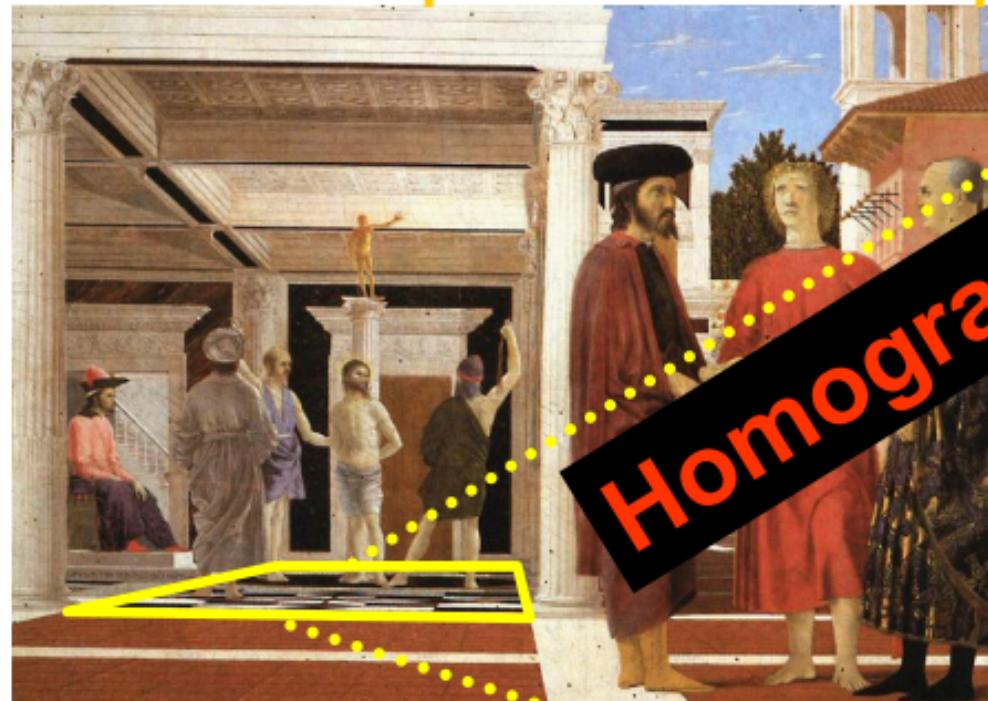
Source: Steve Seitz

Examples - Image warping with homographies



Examples - Image warping with homographies

What is the shape of the b/w floor pattern?



Homography



The floor (enlarged)



Automatically
rectified floor

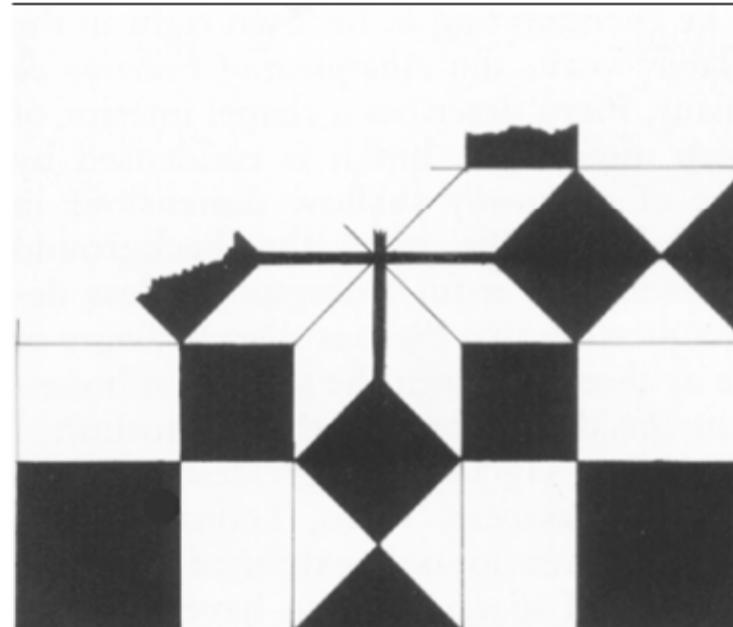
Slide from Criminisi

Examples - Image warping with homographies

Automatic rectification



Slide from Criminisi



From Martin Kemp *The Science of Art*
(manual reconstruction)

Examples - Image warping with homographies



St. Lucy Altarpiece, D. Veneziano

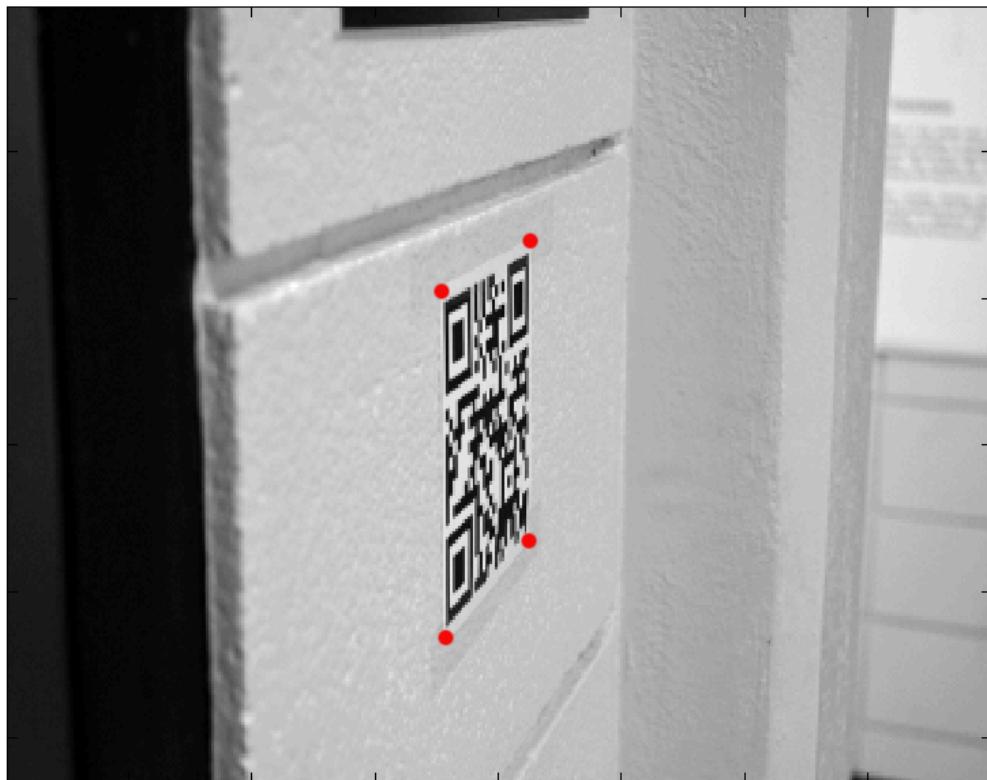
Slide from Criminisi

What is the (complicated) shape of the floor pattern?



Automatically rectified floor

Examples - Image warping with homographies

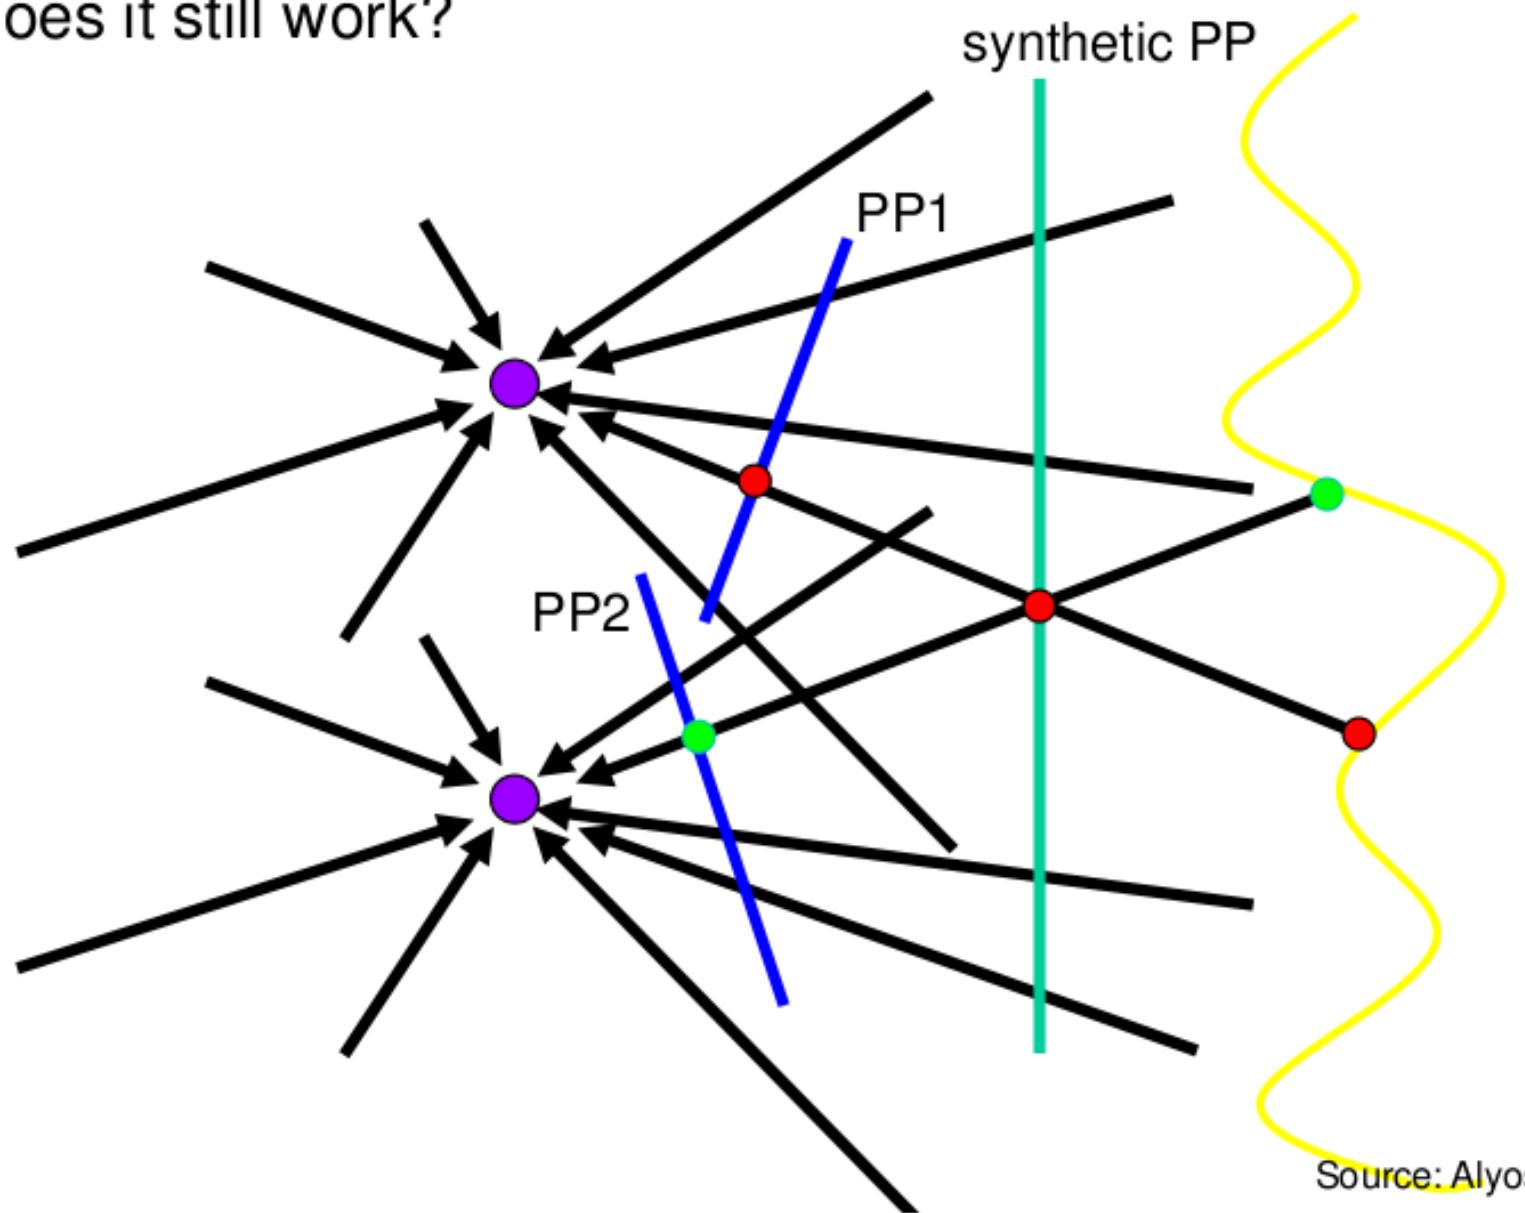


Examples - Image warping with homographies



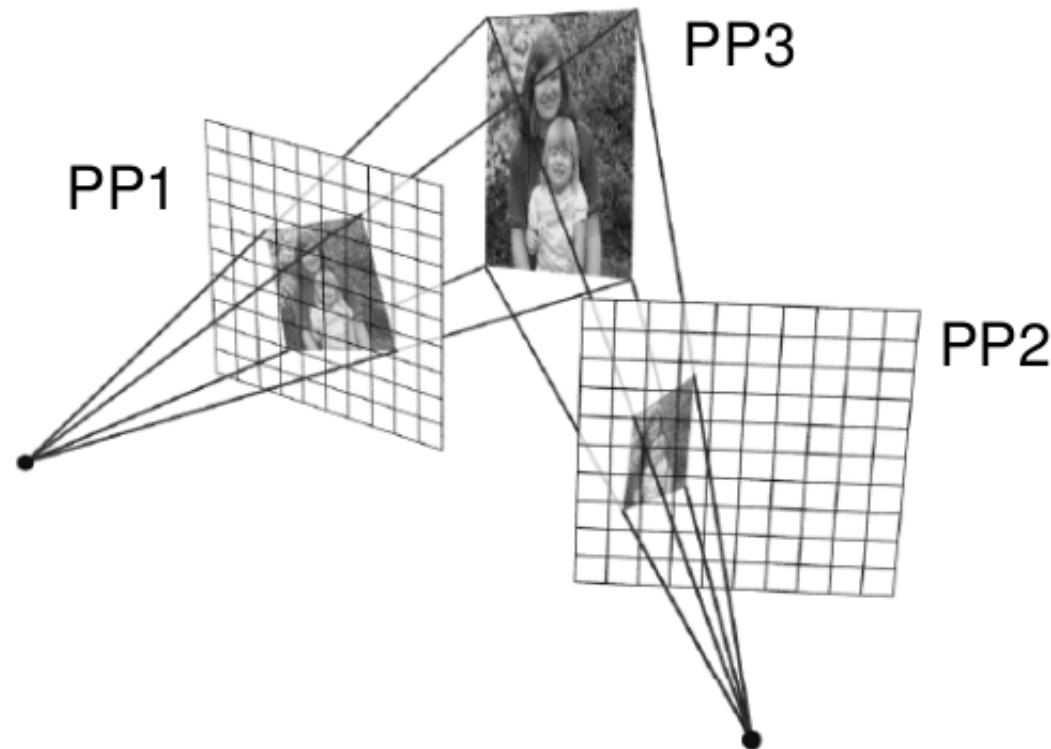
What happens if we change the camera center?

Does it still work?



Source: Alyosha Efros

What happens if we change the camera center?



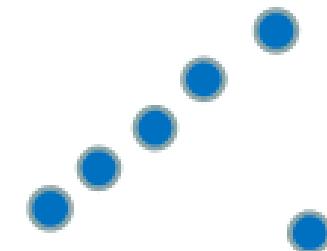
- ▶ PP3 is the projection plane for both centers of projection, so we are okay as long as the scene is relatively far from the camera's local frame
 - ▶ This is how big aerial photographs are stitched!

Recap: How to stitch together a panorama?

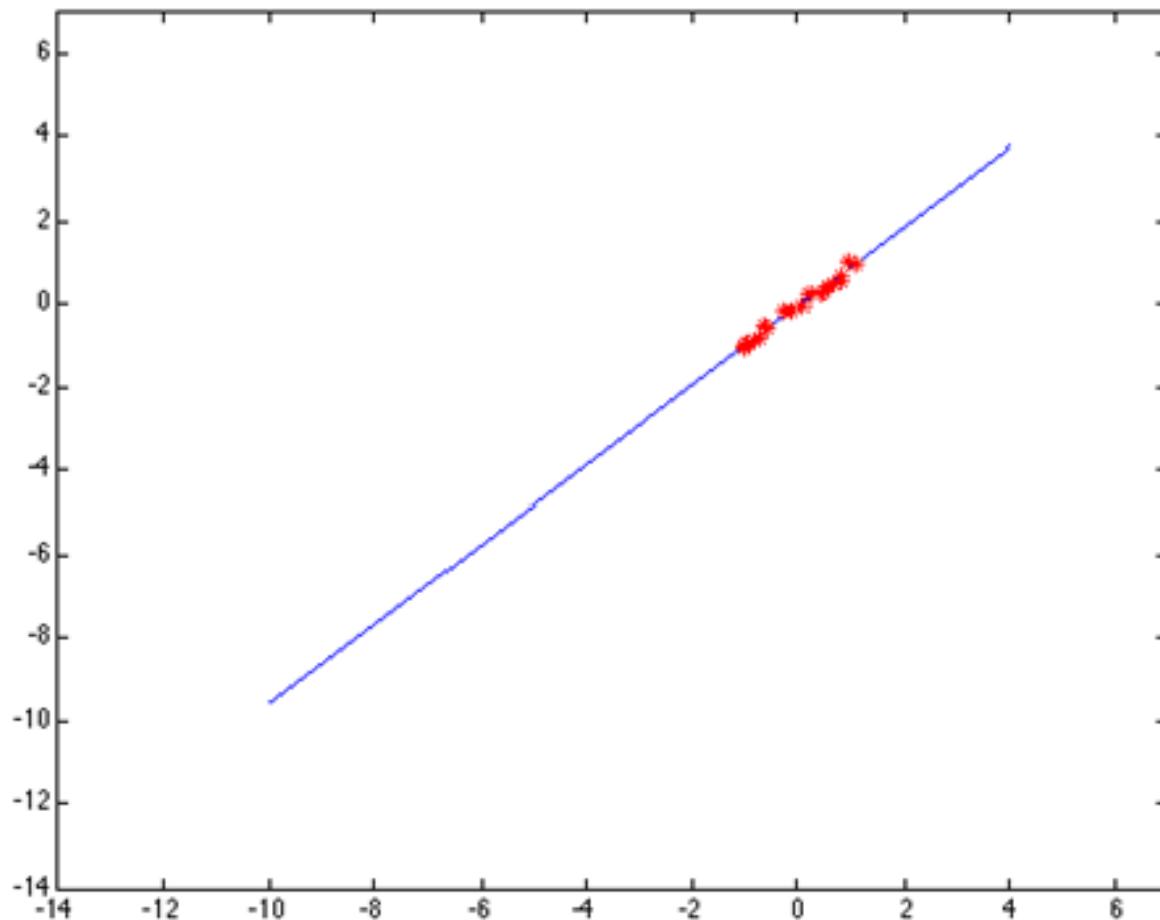
- ▶ Take a sequence of images from the same position
 - ▶ Rotate the camera about its optical center
- ▶ Compute the transformation between the second image and the first by finding common features (correspondences)
- ▶ Transform the second image to overlap the first
- ▶ Blend the two images together (later)
- ▶ If there are more than two images, repeat!

What if some your correspondences are off? - Outliers

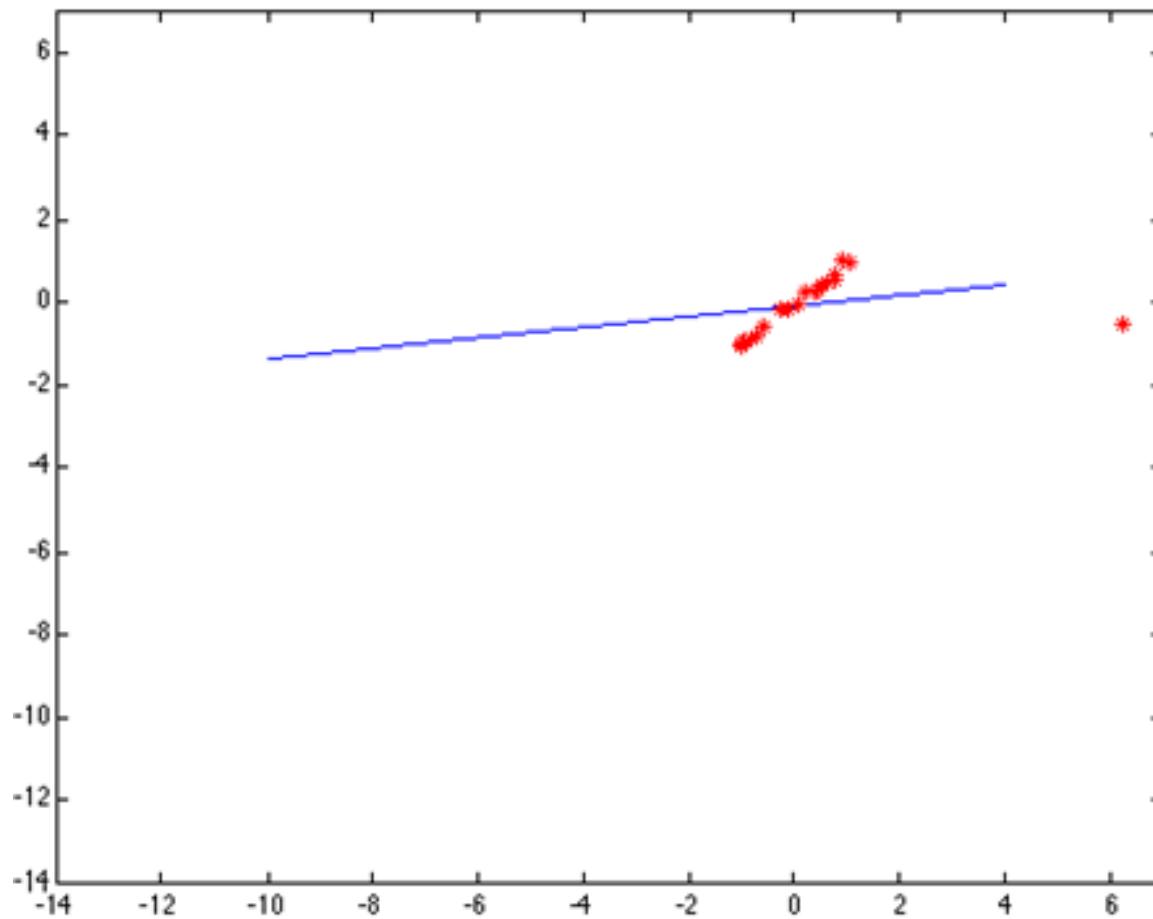
- ▶ **Outliers** will ruin the quality of the homography estimate
 - ▶ Points that you think correspond but don't
 - ▶ Edge point that is really noise or doesn't belong to the line we're fitting



Example: outliers effect our least squares fit



Example: outliers effect our least squares fit



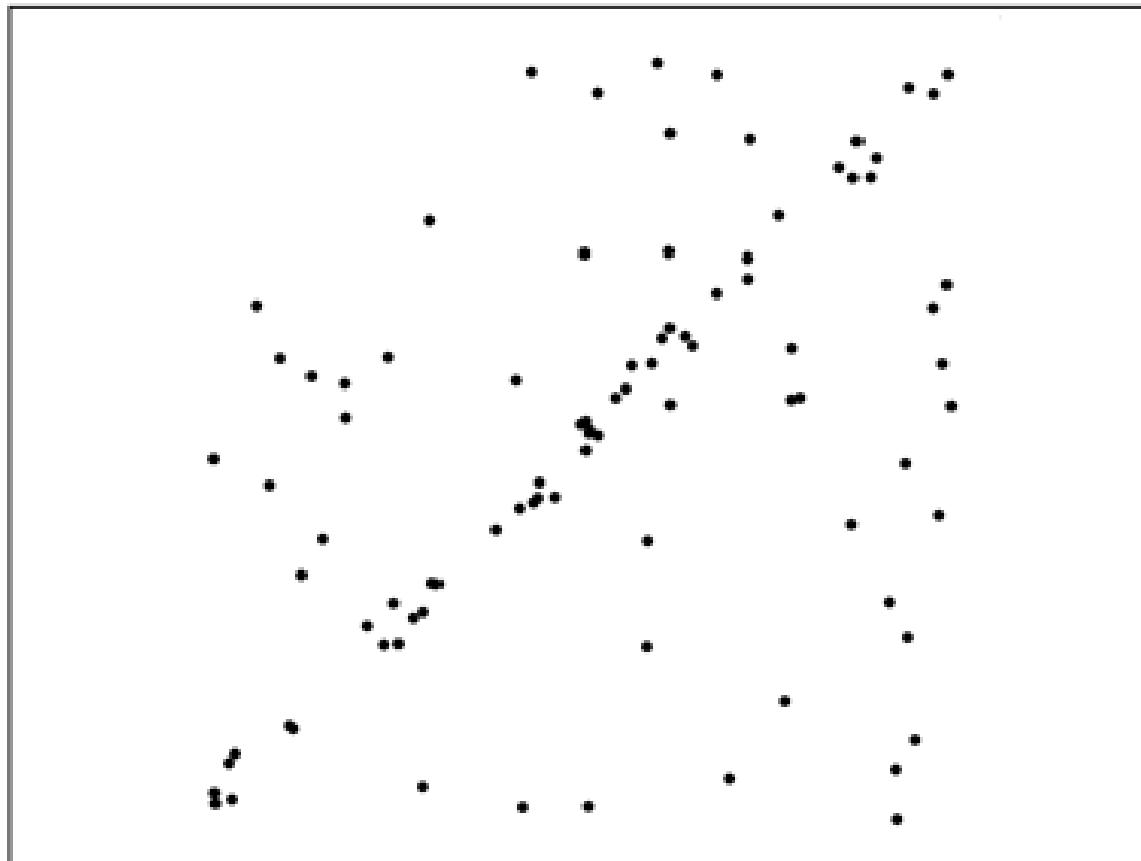
How can we cope with this? - RANSAC

- ▶ RANdom Sample and Consensus
- ▶ **Approach:** we want to avoid the impact of outliers, so let's look for *inliers* and only use those for our computation!
- ▶ **intuition:** if an outlier is chosen to compute the current fit, then the resulting line won't have much support from the rest of the points.

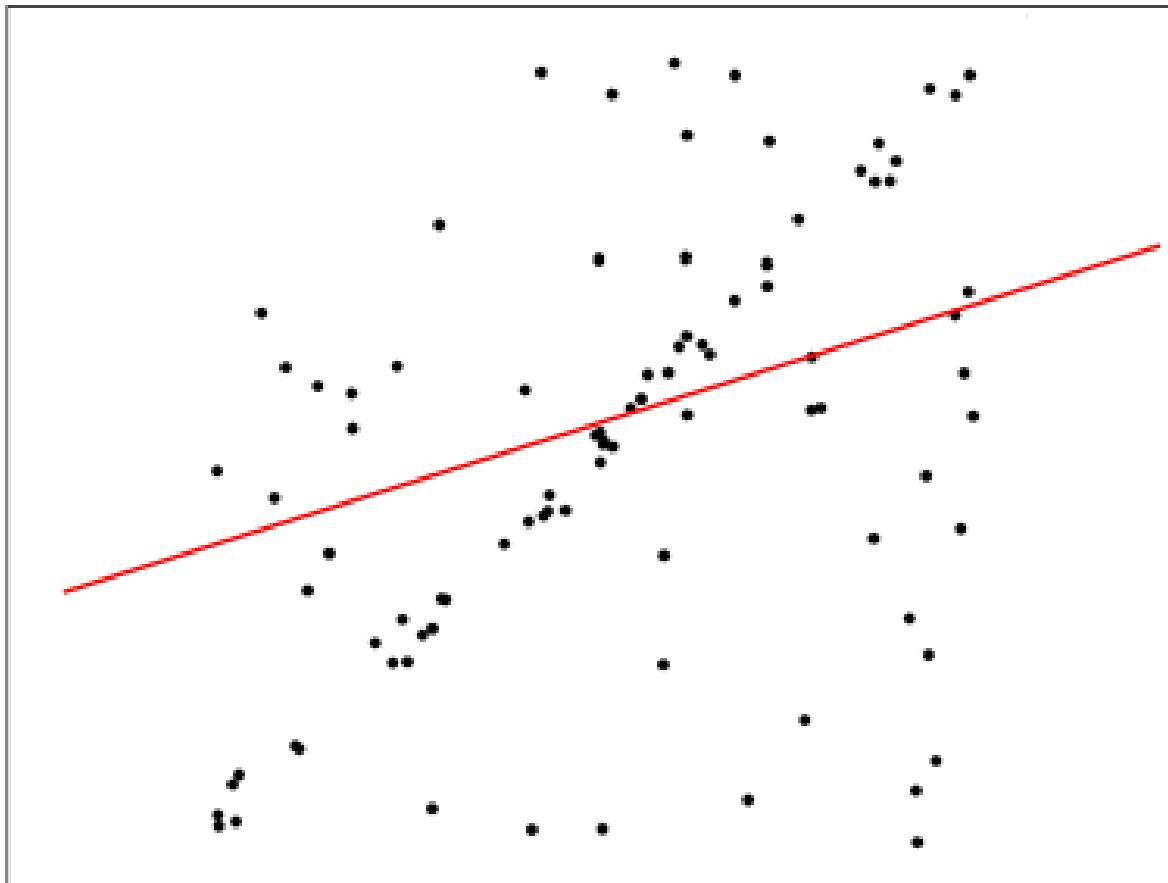
RANSAC: General form

- ▶ RANSAC loop:
 - ① Randomly select a *seed group* of points on which to compute our estimation (e.g. a group of putative matches)
 - ② Compute the transform (homography in our case) from this *seed* group
 - ③ Find the *inliers* resulting from this transform
 - ④ If the number of inliers is sufficiently large, re-compute an estimate of the transform using all the inliers
 - ⑤ If NOT, repeat the process!
- ▶ Goal: keep the transform with the largest number of inliers!

RANSAC Example

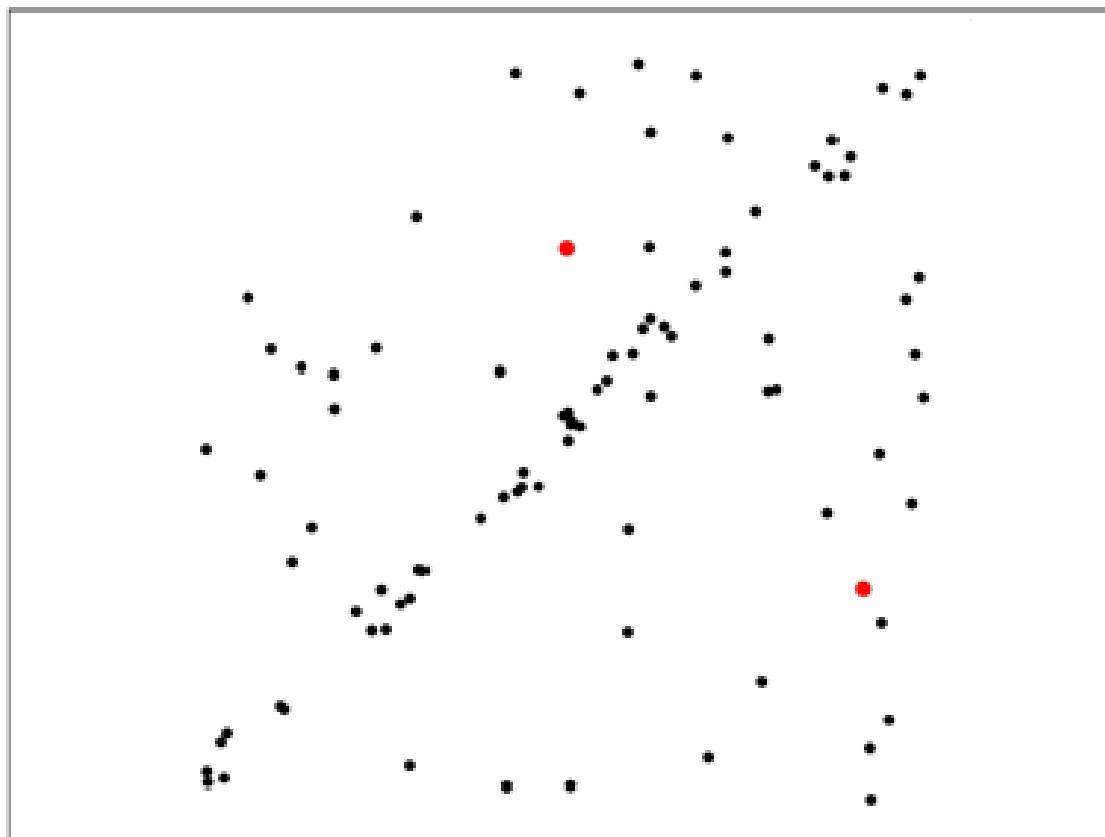


RANSAC Example



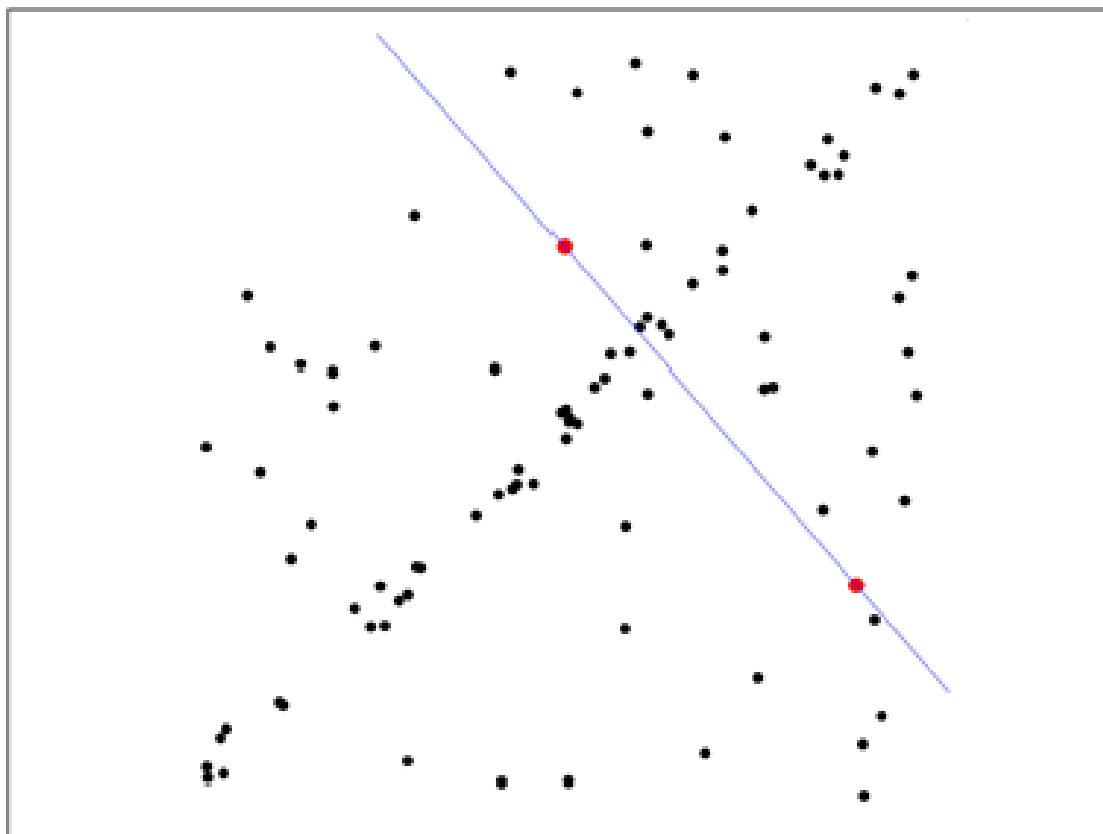
Least-squares fit

RANSAC Example



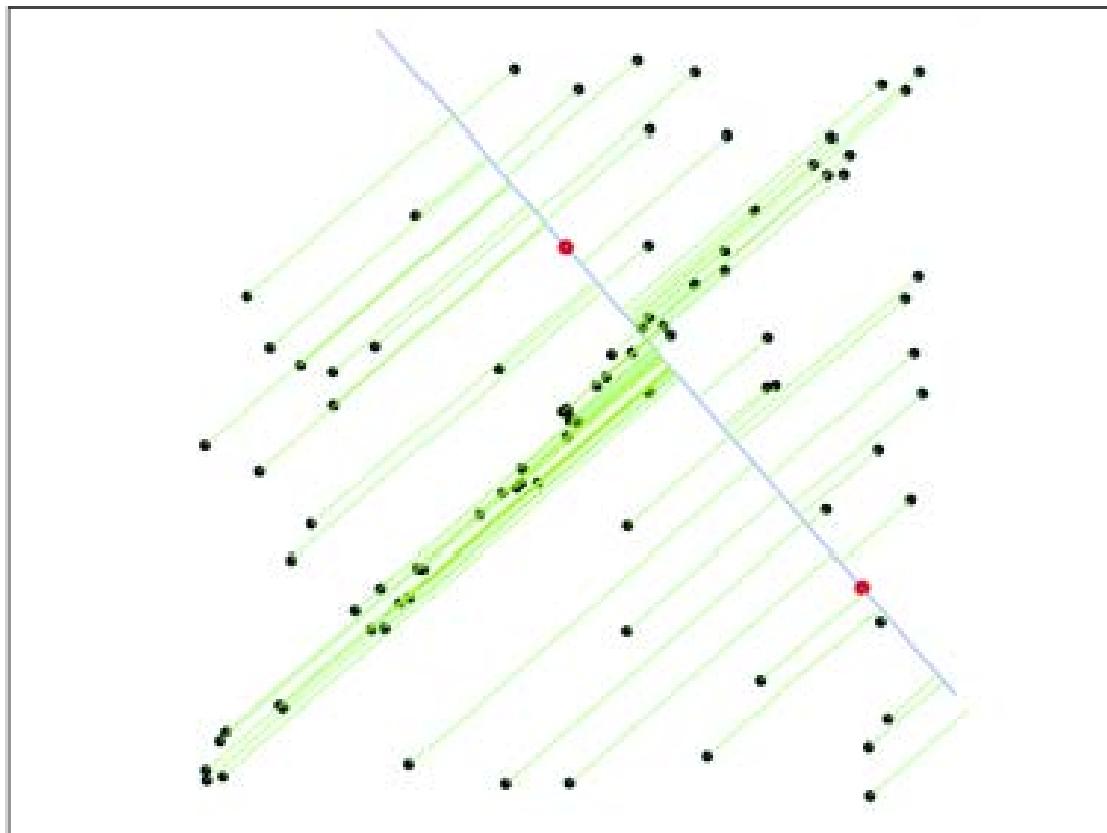
1. Randomly select minimal subset of points

RANSAC Example



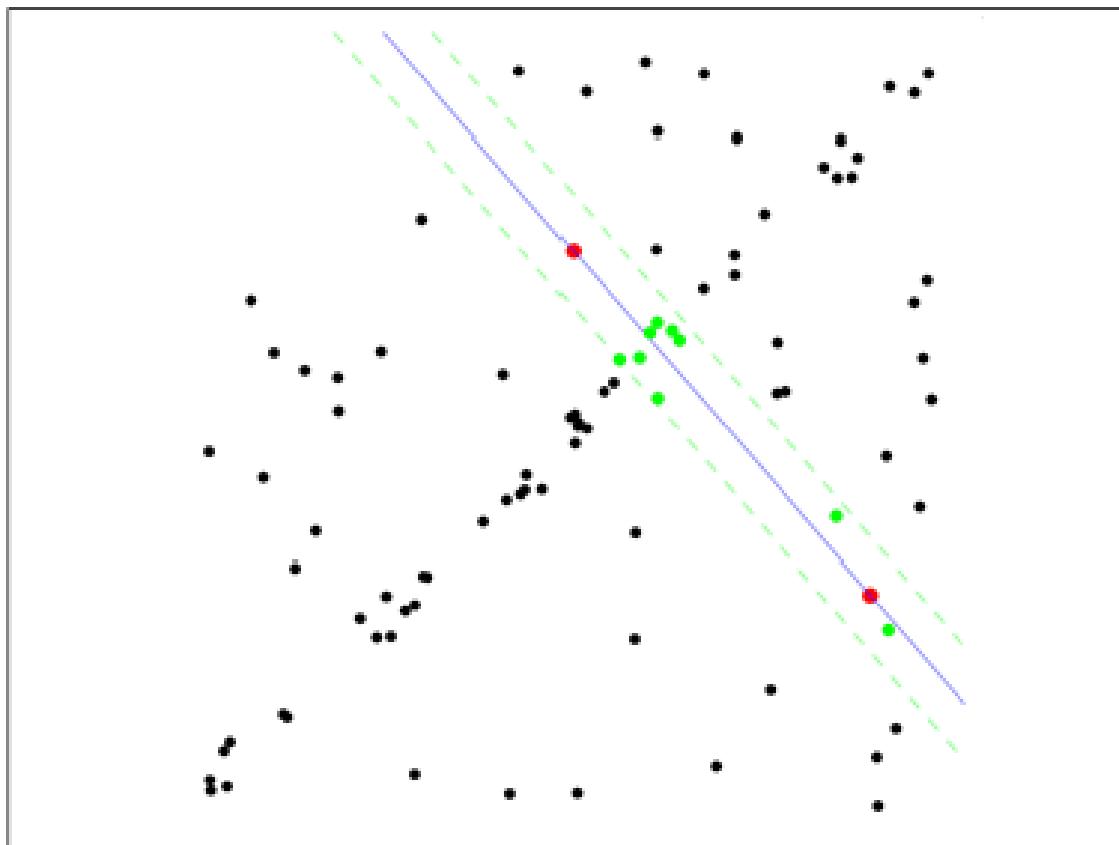
1. Randomly select minimal subset of points
2. Hypothesize a model

RANSAC Example



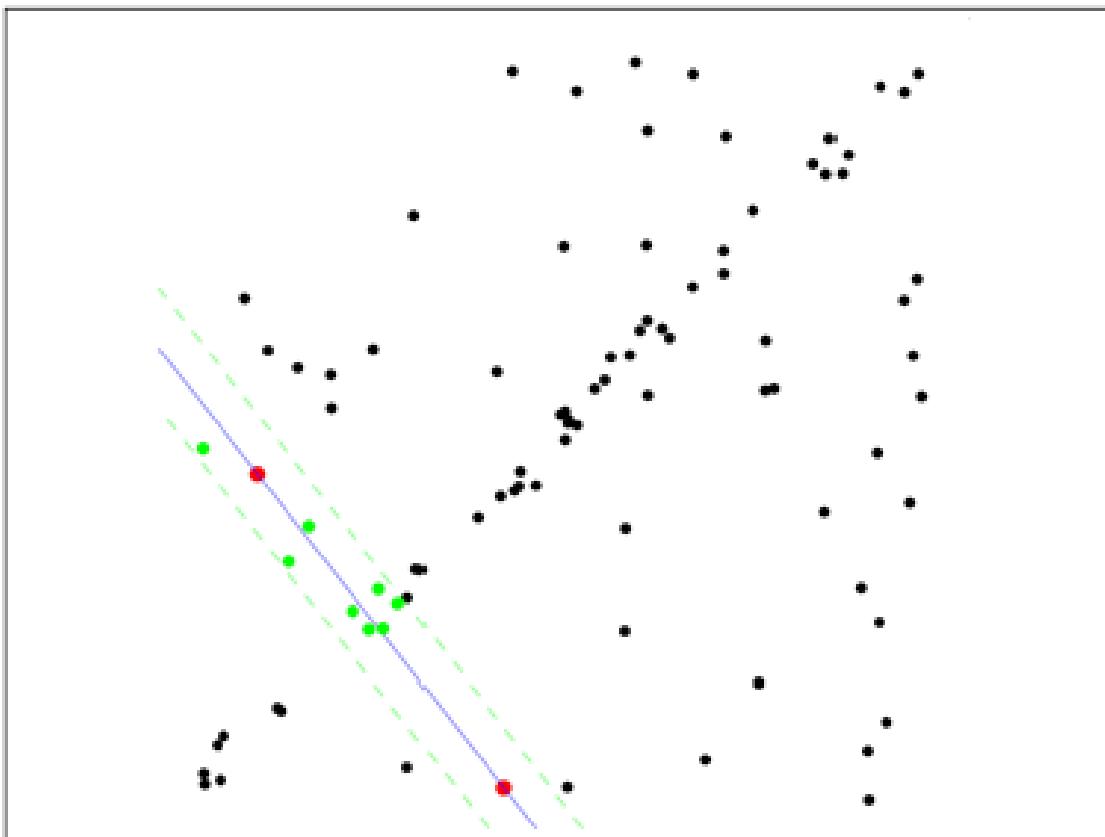
1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function

RANSAC Example



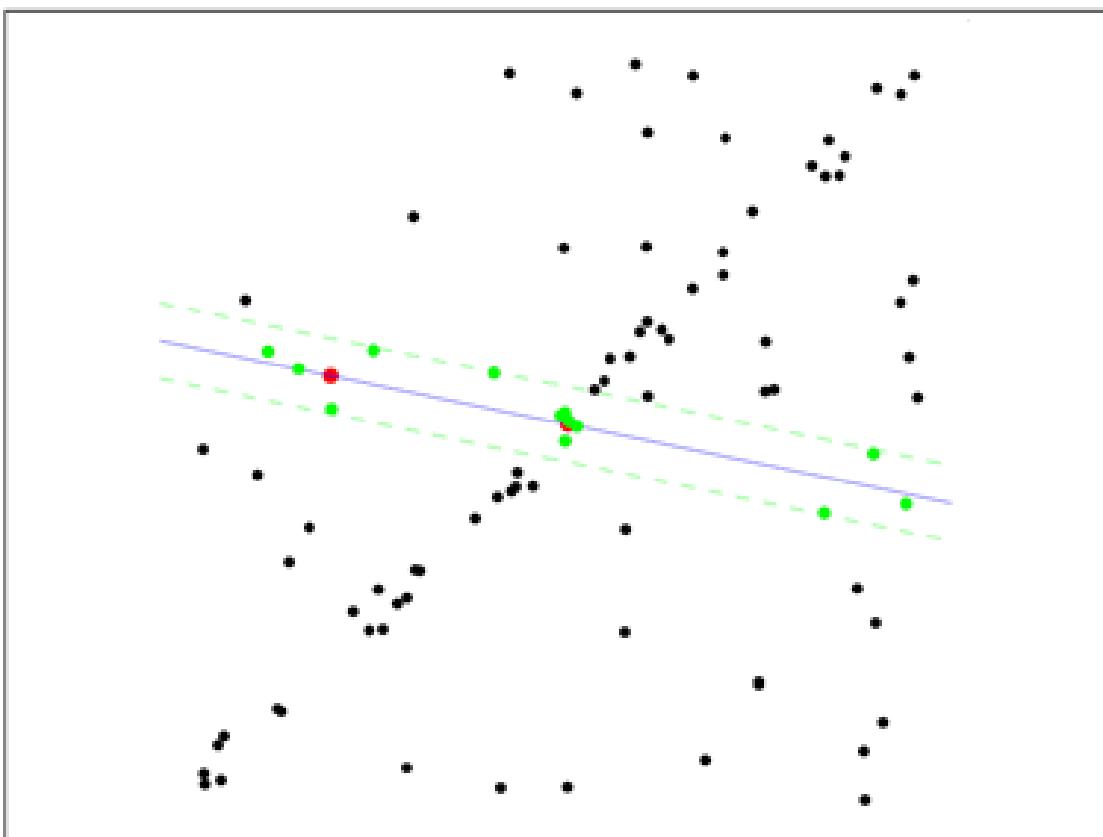
1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model

RANSAC Example



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop

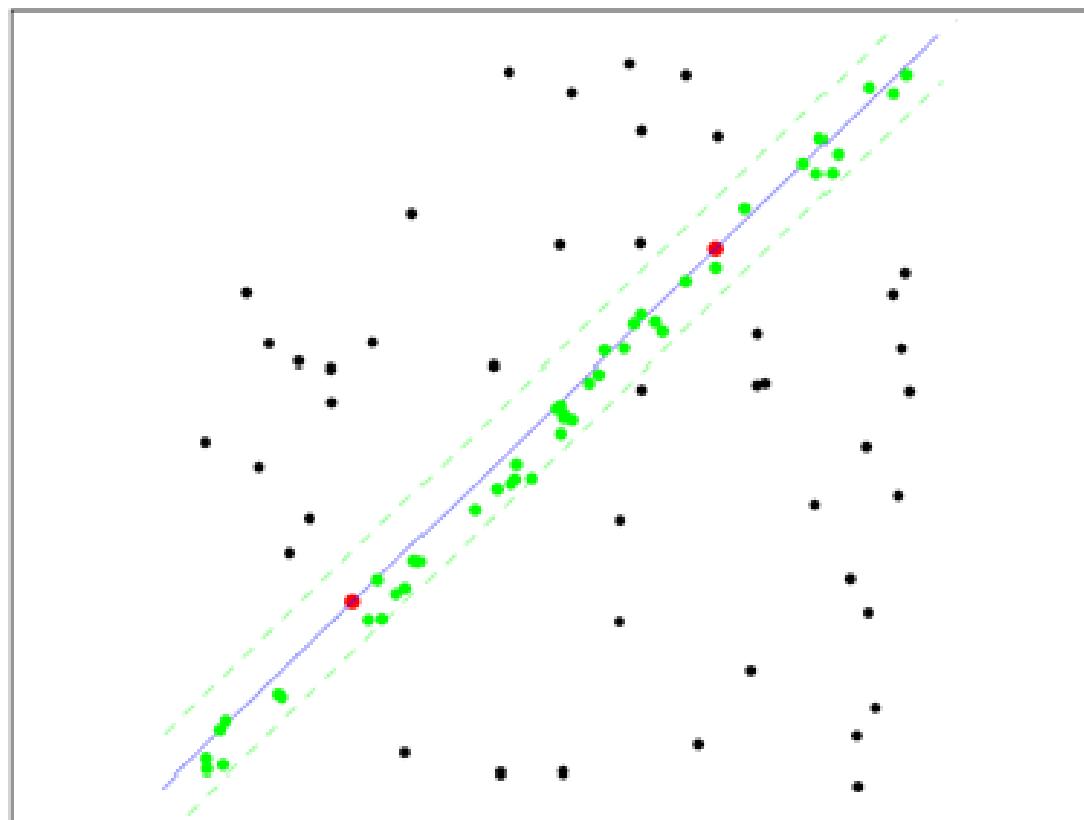
RANSAC Example



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop

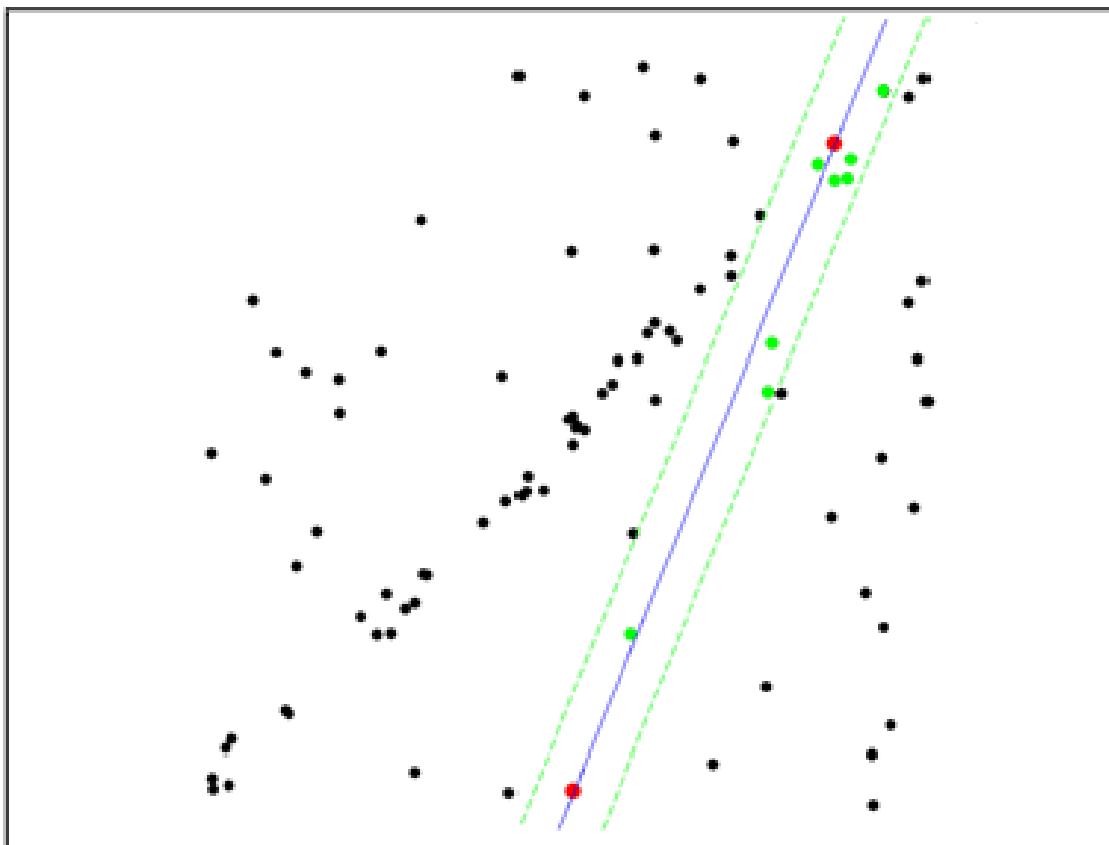
RANSAC Example

Uncontaminated sample



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop

RANSAC Example



1. Randomly select minimal subset of points
2. Hypothesize a model
3. Compute error function
4. Select points consistent with model
5. Repeat *hypothesize-and-verify* loop

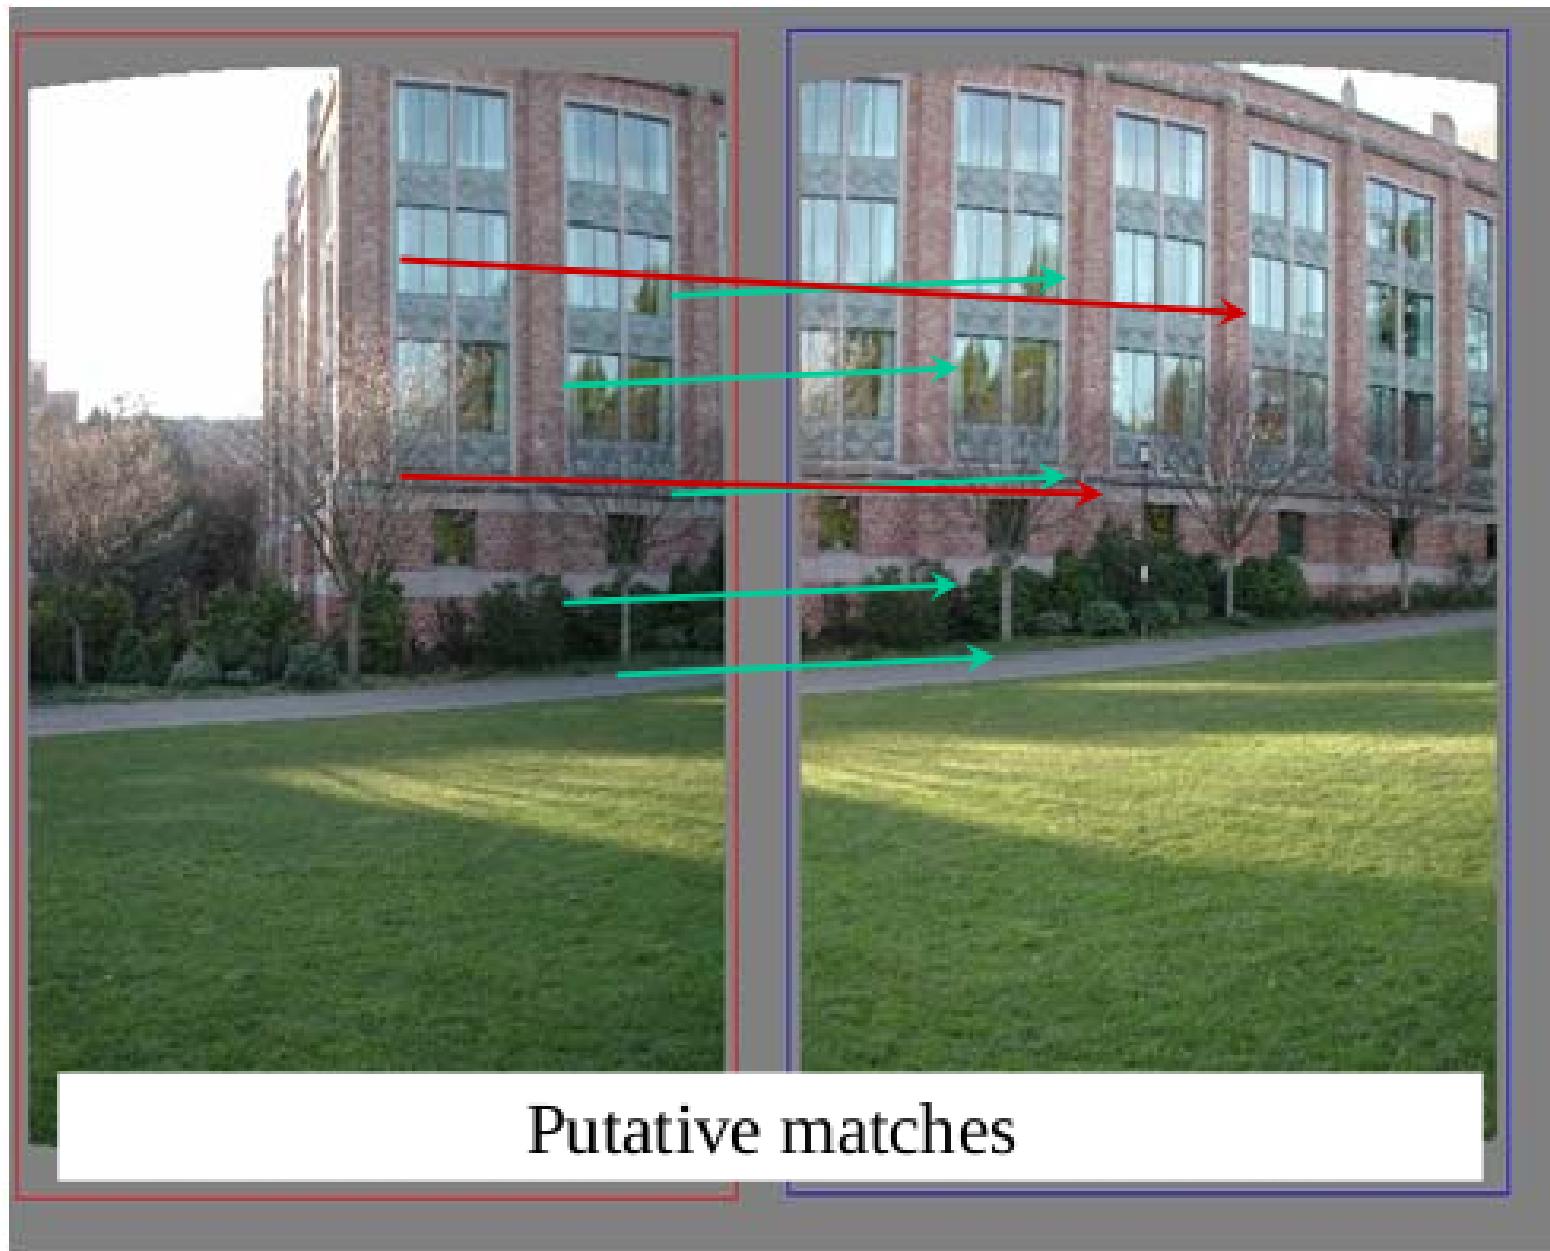
RANSAC: for line fitting

- ▶ Repeat **N** times
 - ① Randomly select a **s** points at random
 - ② Fit a line to these points
 - ③ Find the inliers to this line among the remaining points (i.e., points whose distance from the line is less than **t**)
 - ④ If there are **d** or more inliers, accept the line and refit using all the inliers

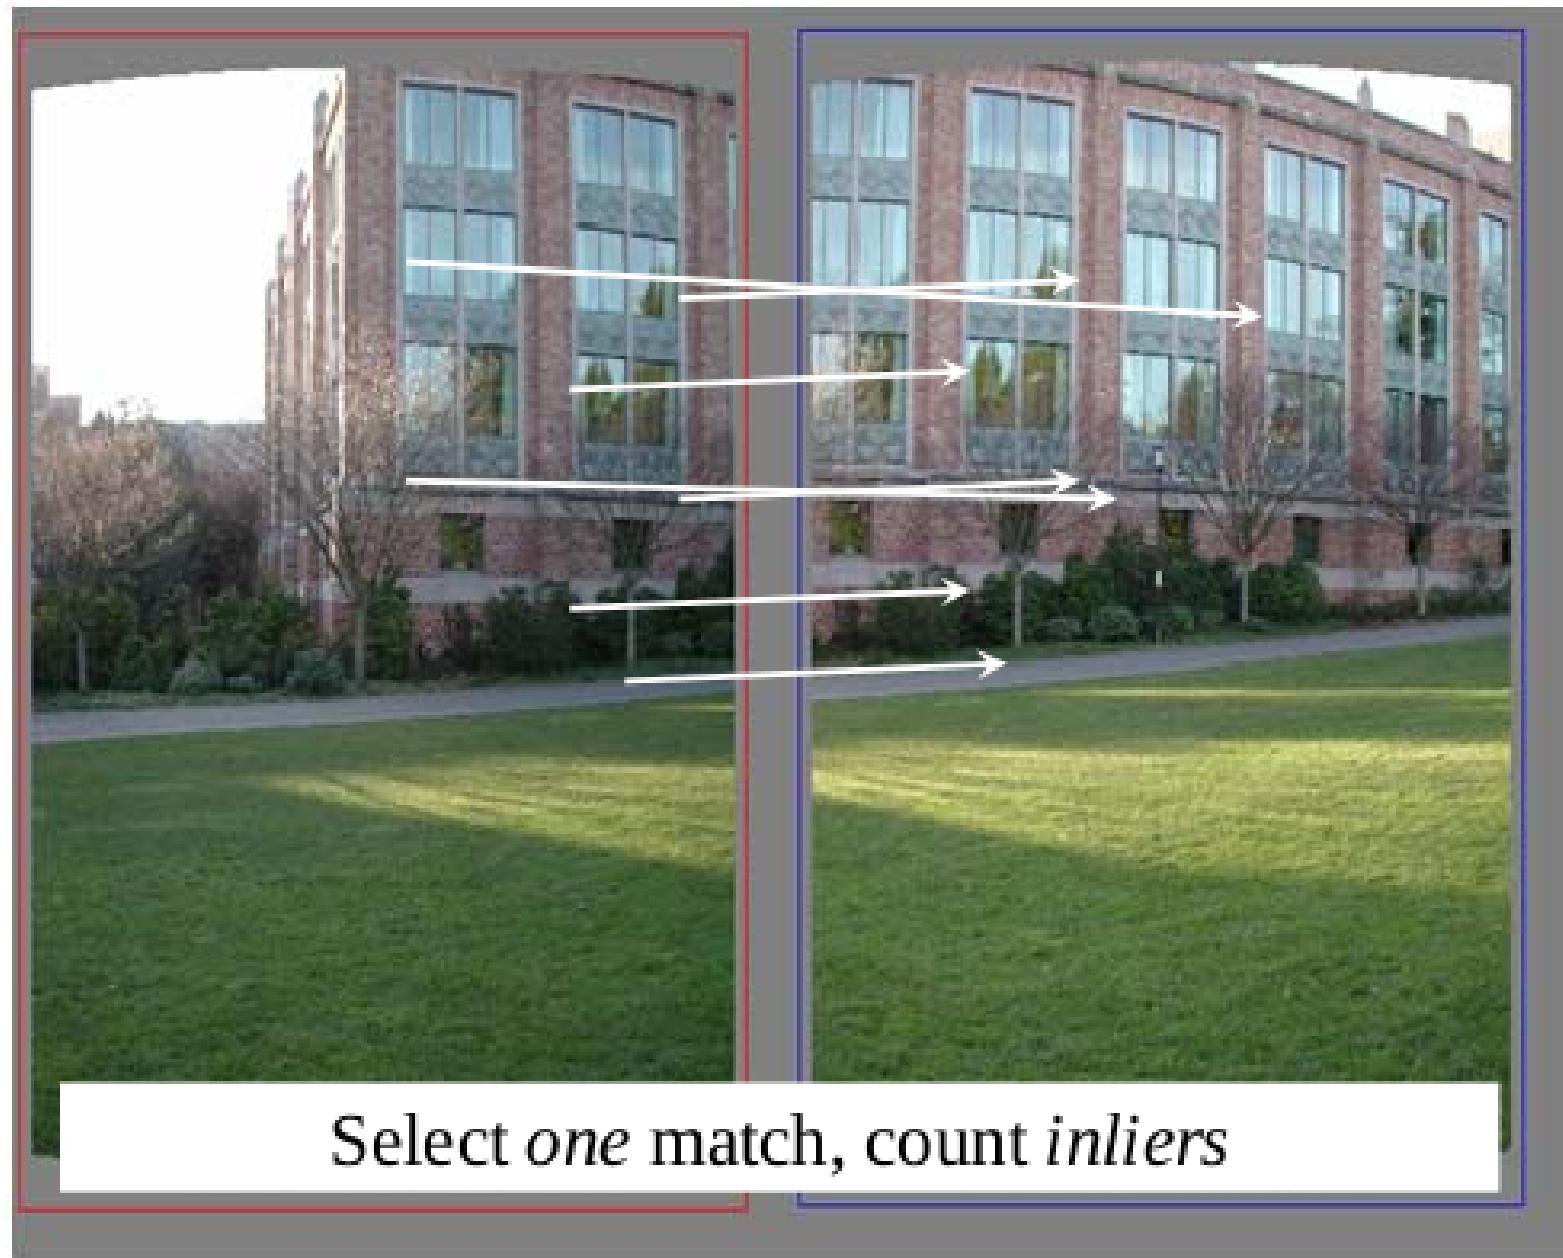
RANSAC: for line fitting

- ▶ This is an example of how to fit a model (line)!
- ▶ How would we do this for a general transformation?

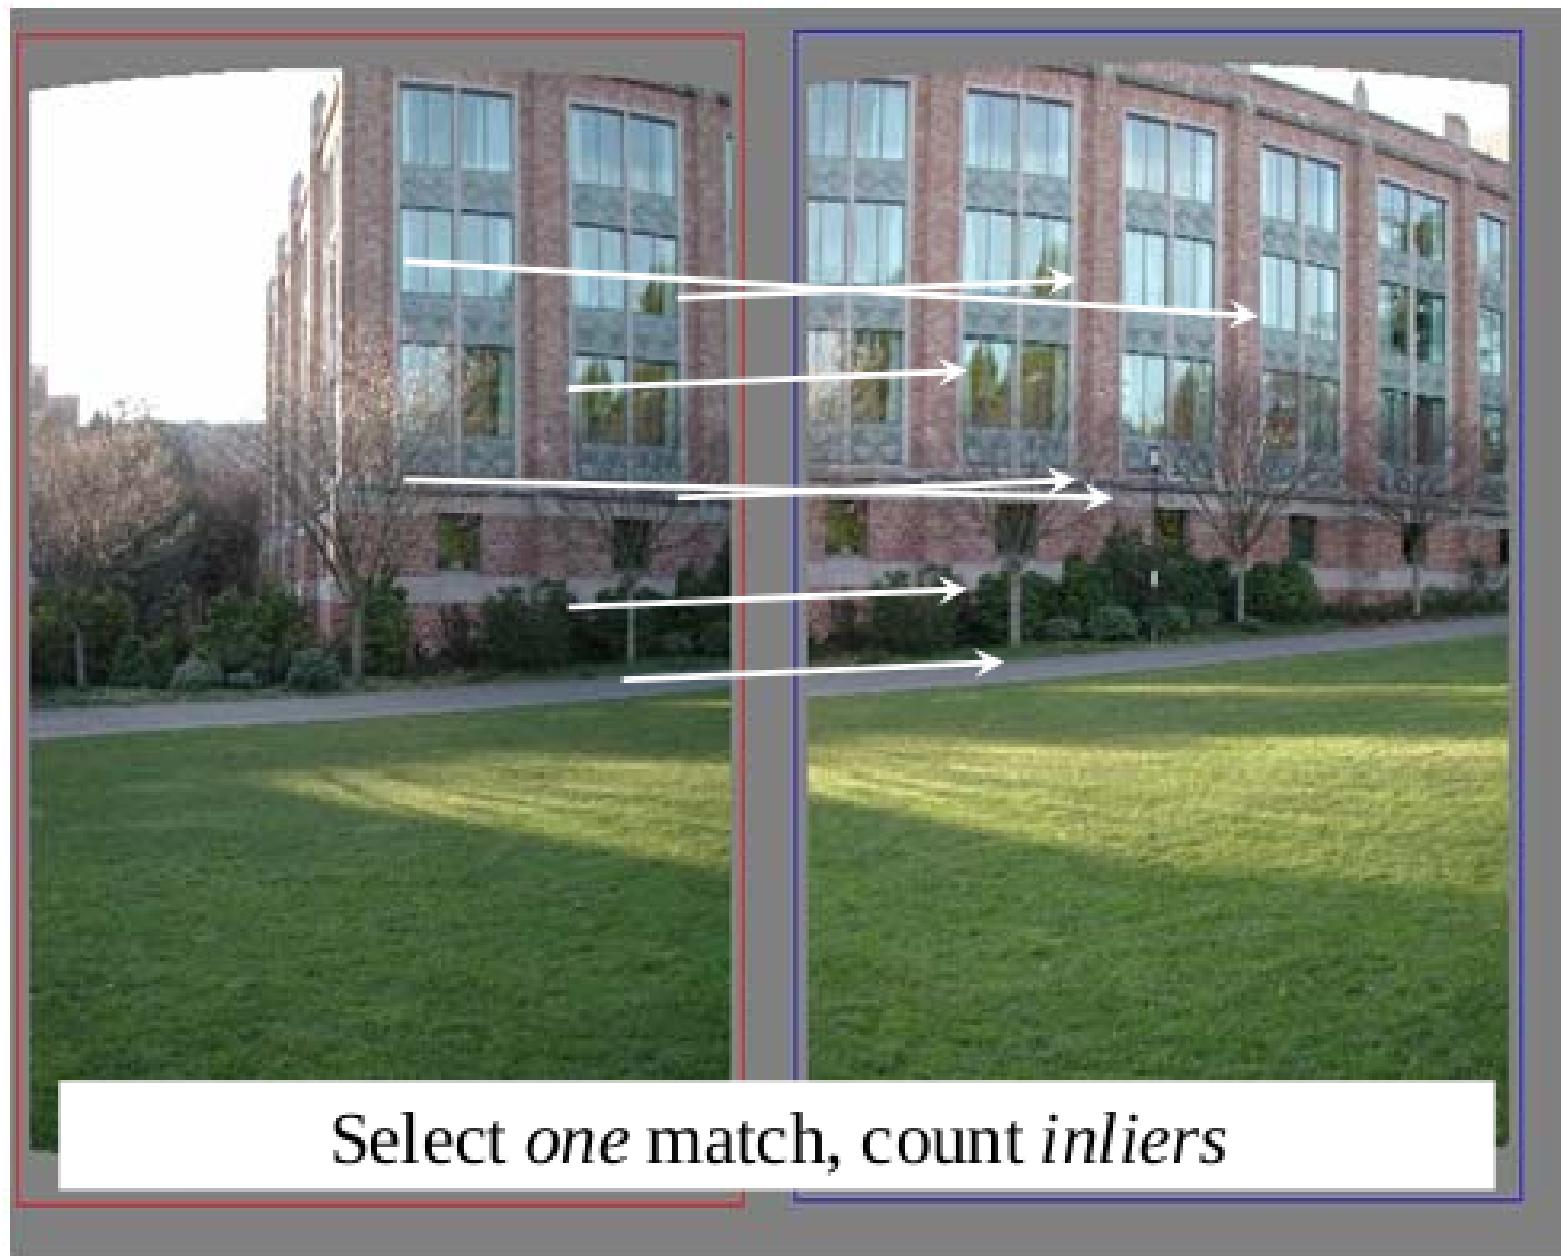
RANSAC Example: Translation



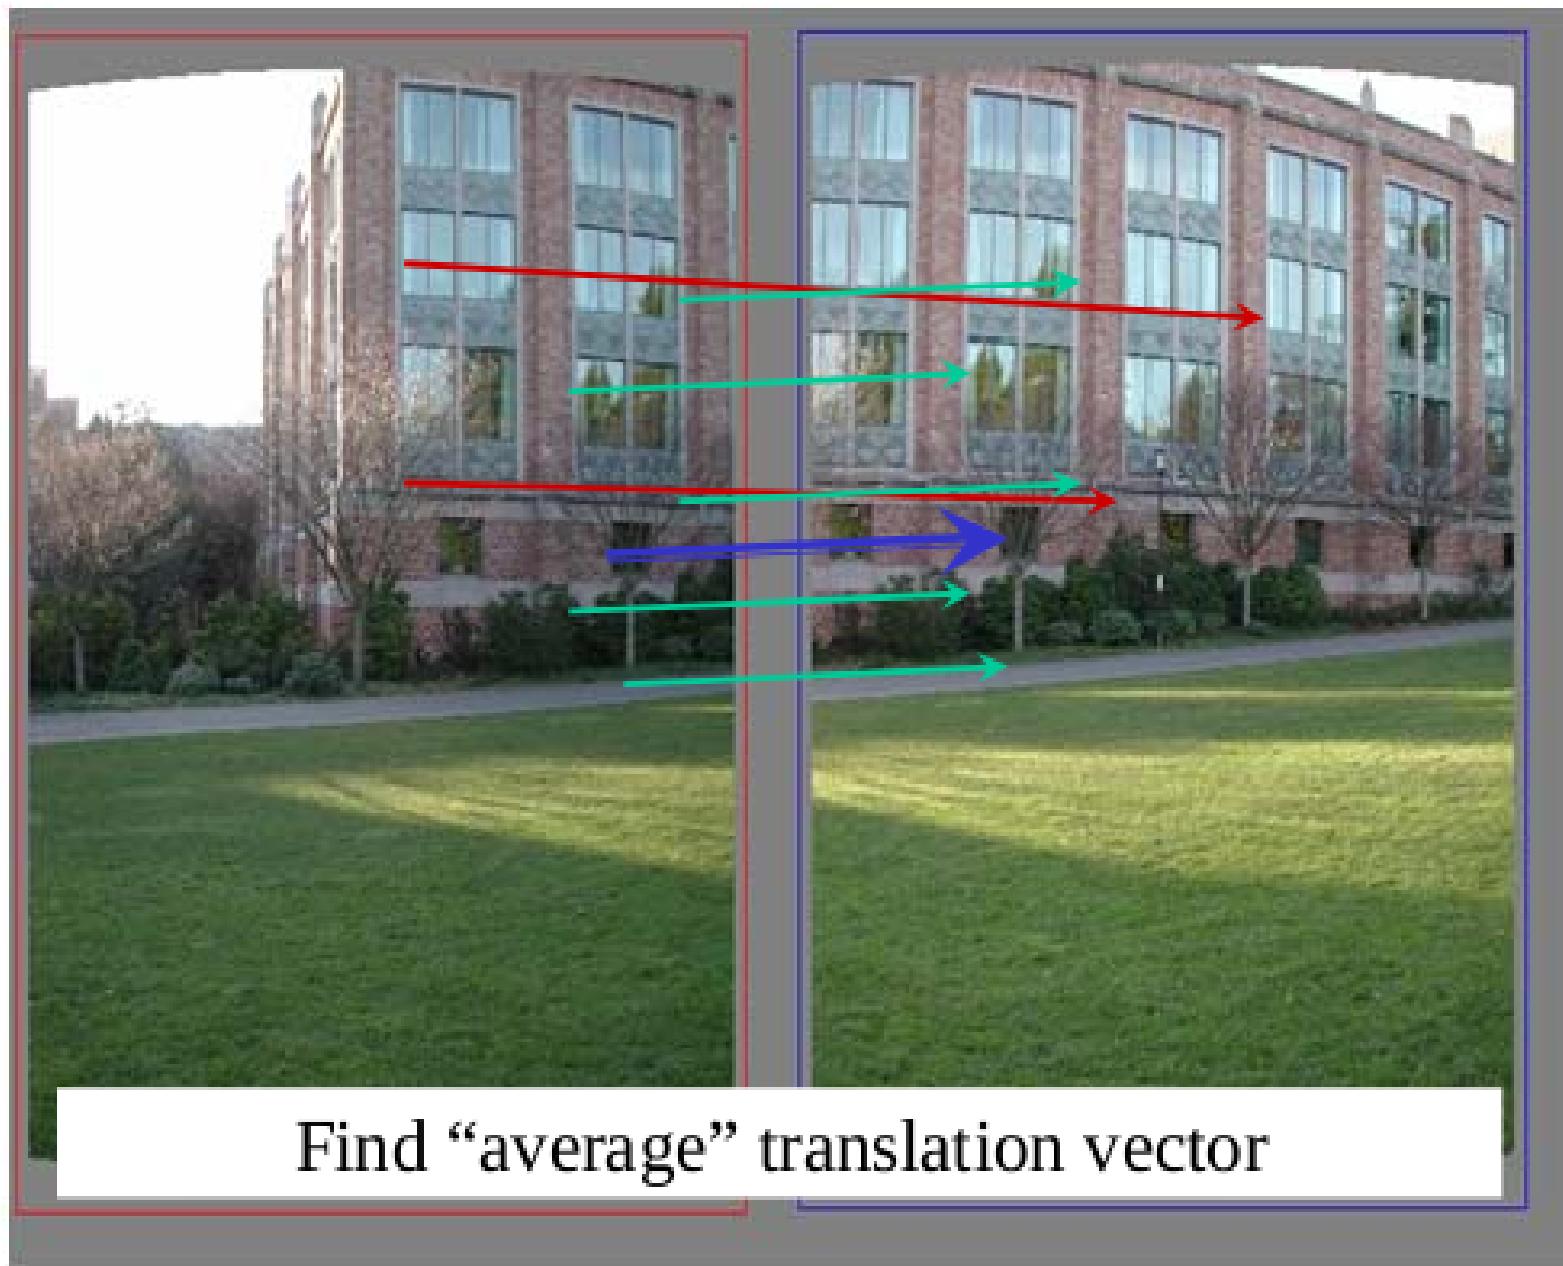
RANSAC Example: Translation



RANSAC Example: Translation



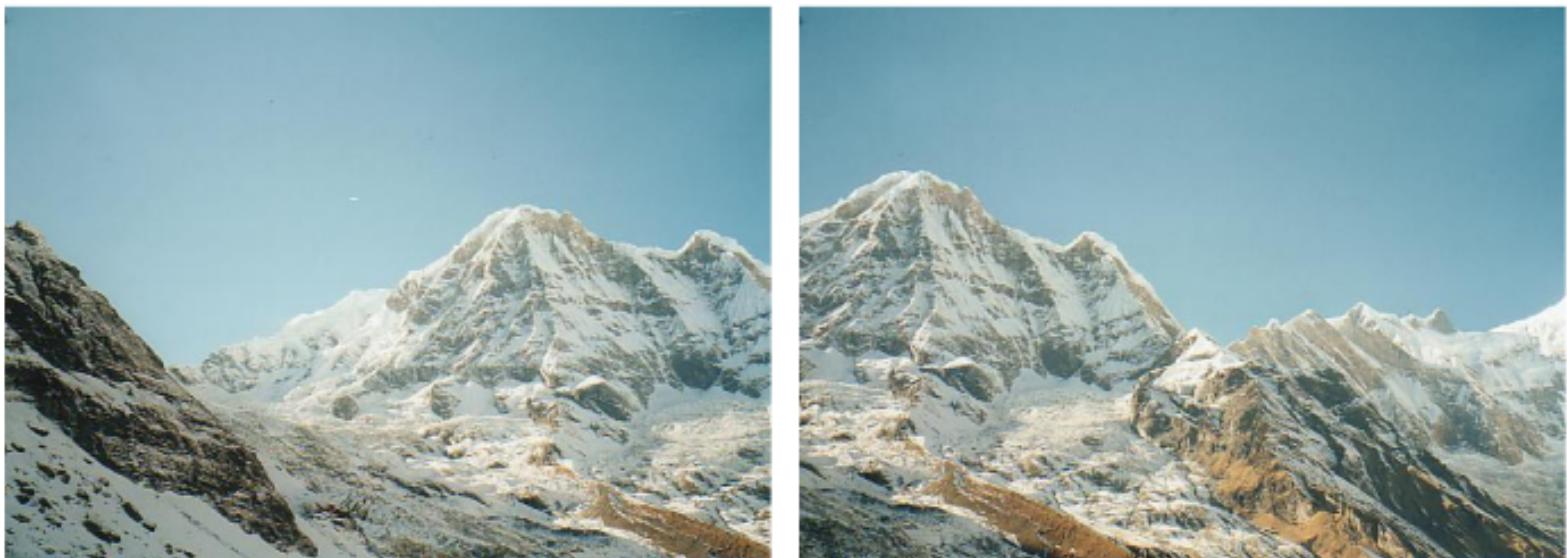
RANSAC Example: Translation



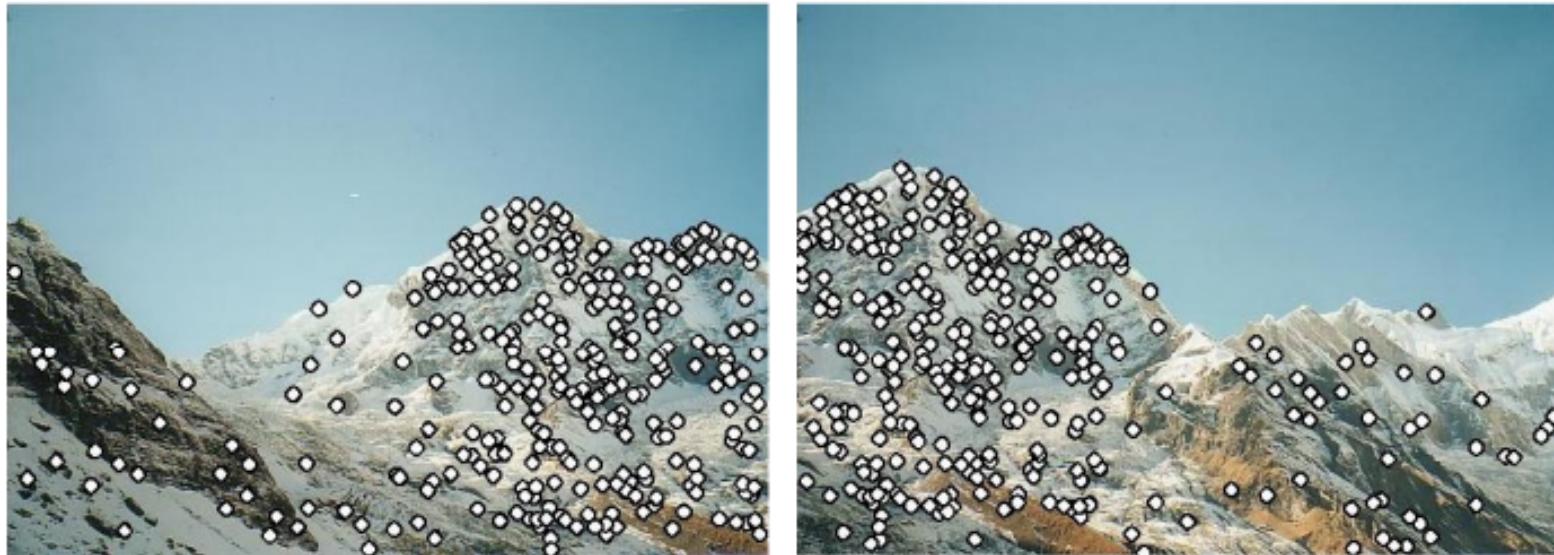
RANSAC pros and cons

- ▶ Pros
 - ▶ Simple and general
 - ▶ Applicable to many different problems
 - ▶ Often works well in practice
- ▶ Cons
 - ▶ Lots of parameters to tune
 - ▶ Doesn't work well for low inlier ratios (too many iterations or can fail completely)
 - ▶ Can't always get a good initialization of the model based on the minimum number of samples

Recap: Feature-based alignment



Recap: Feature-based alignment



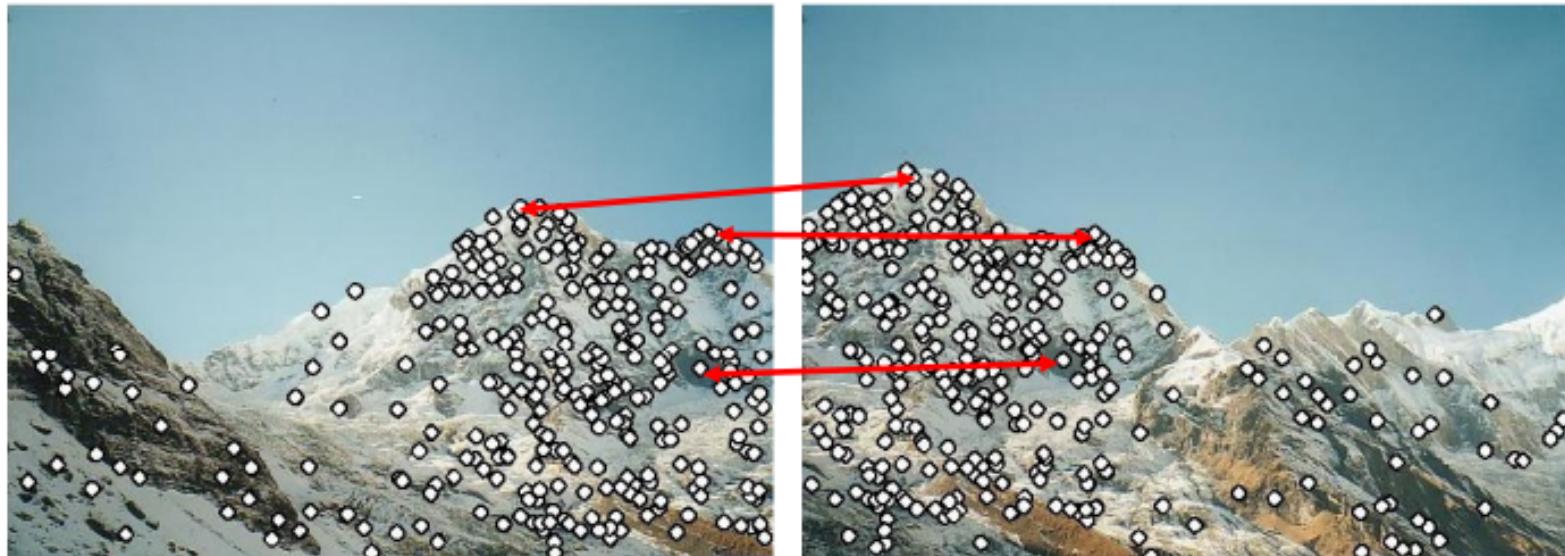
- ▶ Extract features

Recap: Feature-based alignment



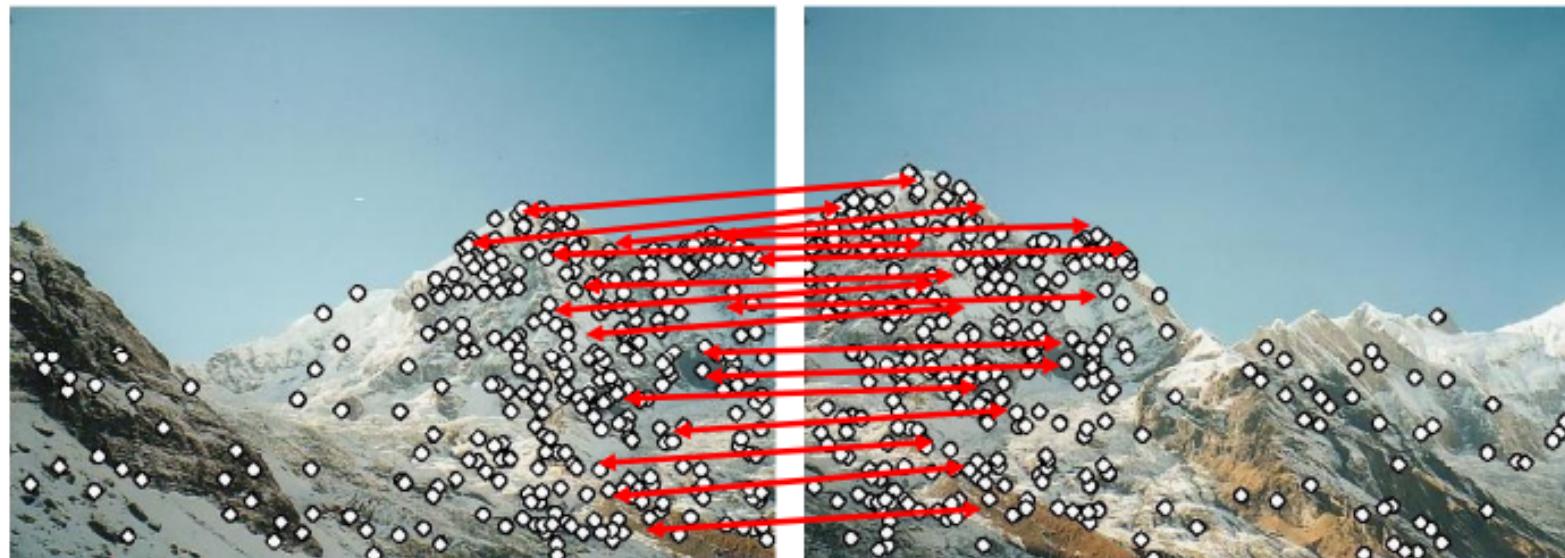
- ▶ Extract features
- ▶ Compute *punitive* matches

Recap: Feature-based alignment



- ▶ Extract features
- ▶ Compute *punitive* matches
- ▶ Loop:
 - ▶ Hypothesize transformation H (small group of punitive matches that are related by H)

Recap: Feature-based alignment



- ▶ Extract features
- ▶ Compute *punitive* matches
- ▶ Loop:
 - ▶ Hypothesize transformation H (small group of punitive matches that are related by H)
 - ▶ Verify transformation (search for other matches consistent with H)

Recap: Feature-based alignment



- ▶ Extract features
- ▶ Compute *punitive* matches
- ▶ Loop:
 - ▶ Hypothesize transformation H (small group of punitive matches that are related by H)
 - ▶ Verify transformation (search for other matches consistent with H)

Image blending - feathering

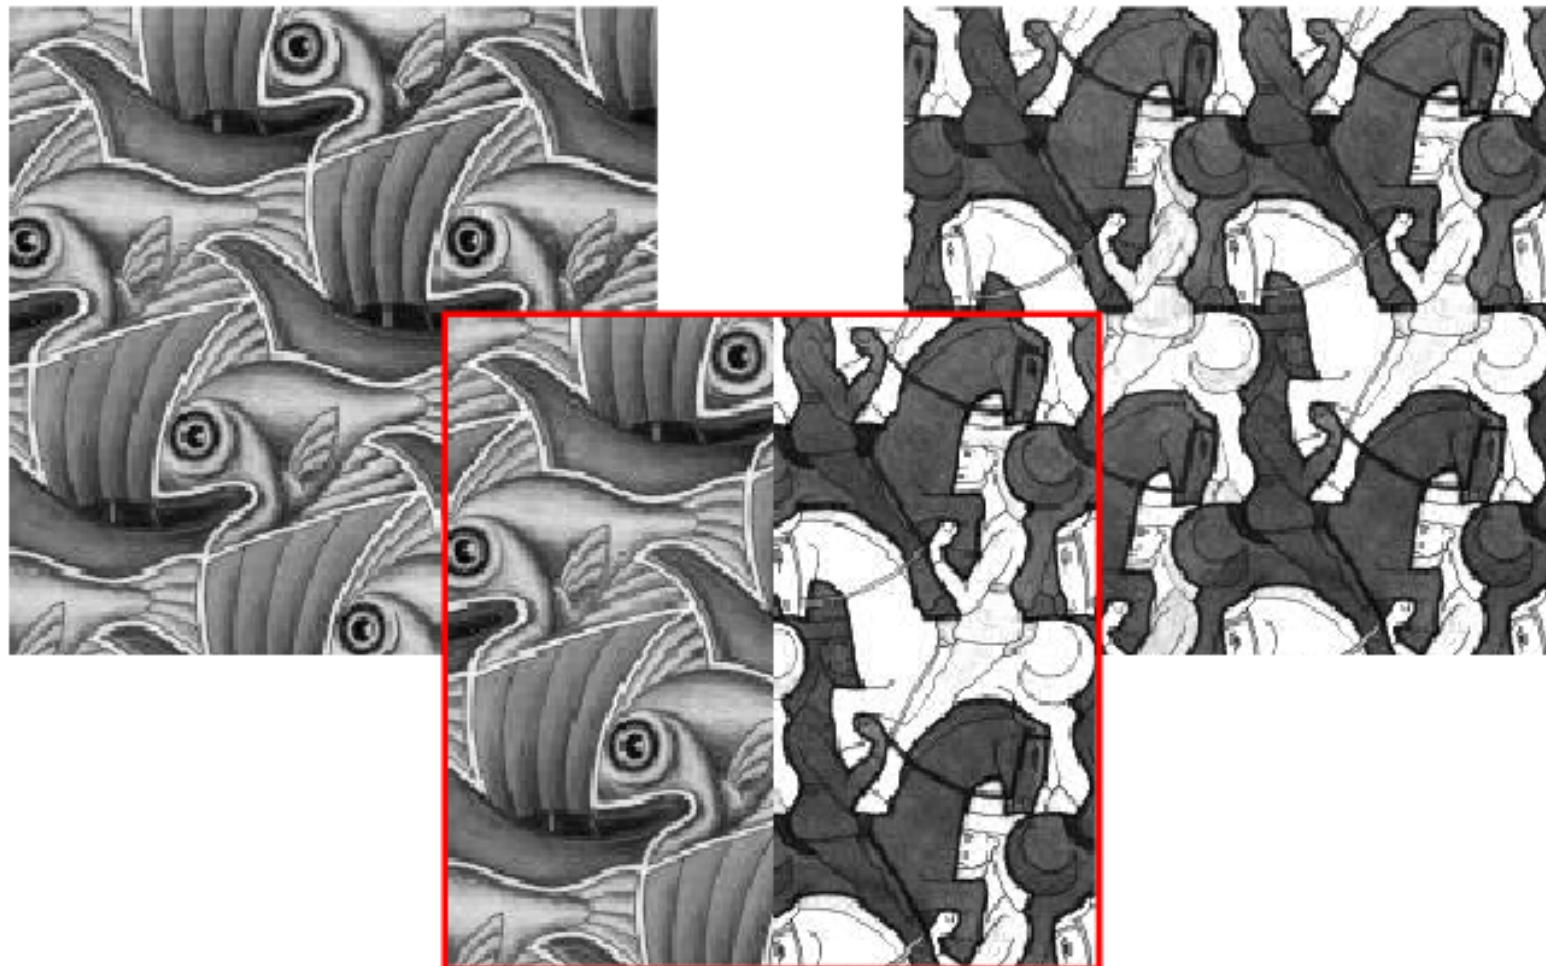


Image blending - feathering

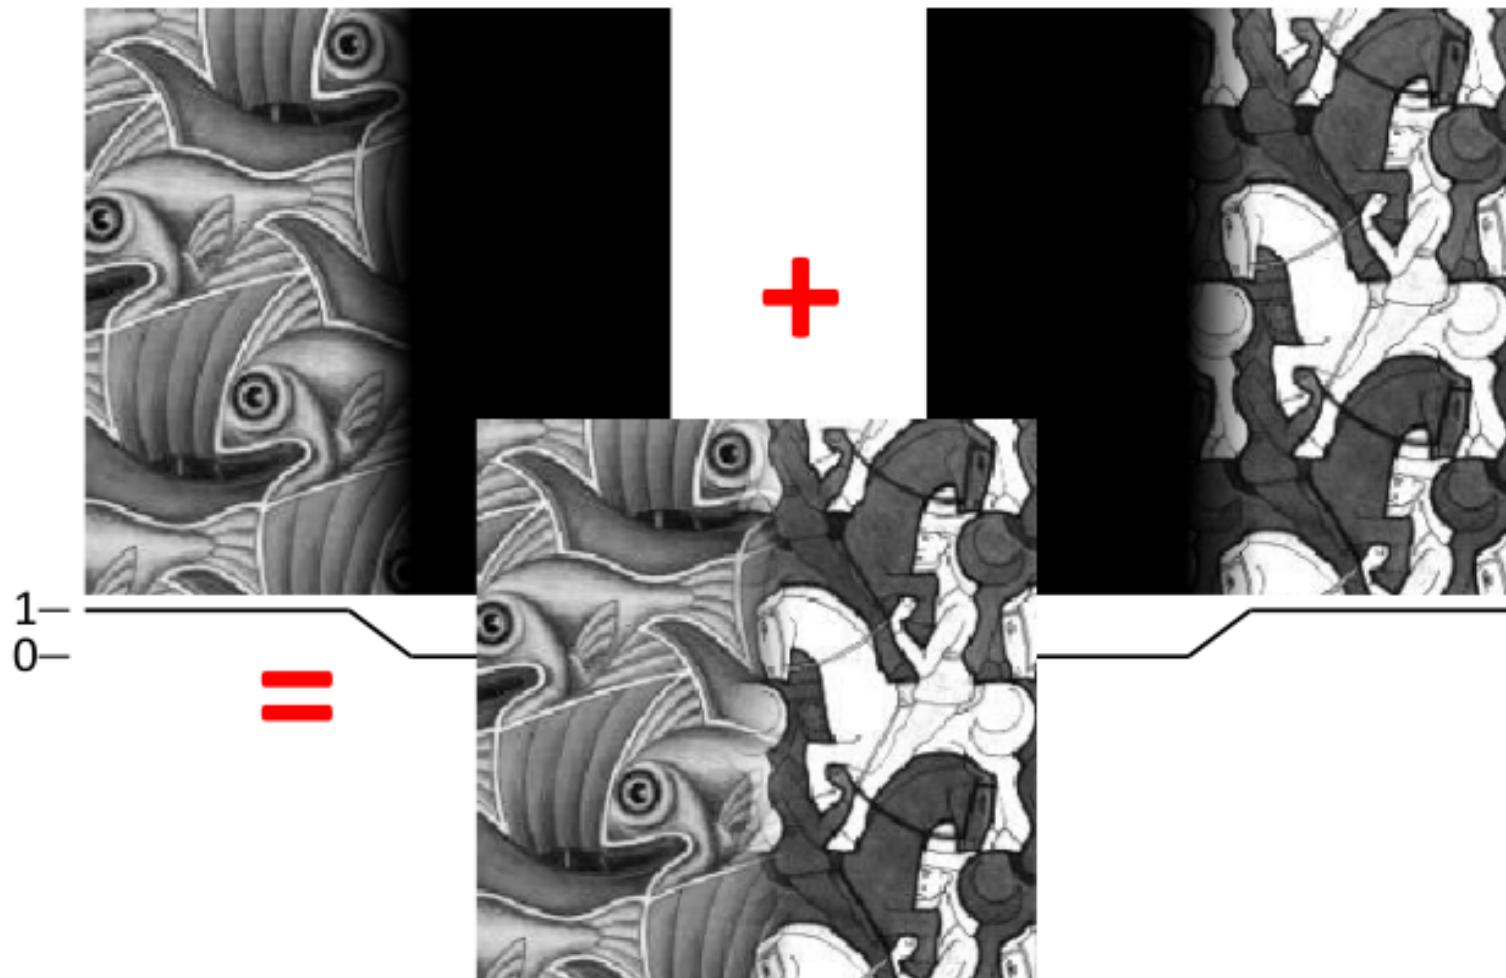


Image blending - feathering

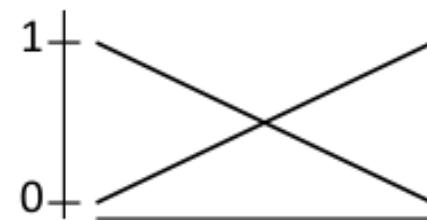
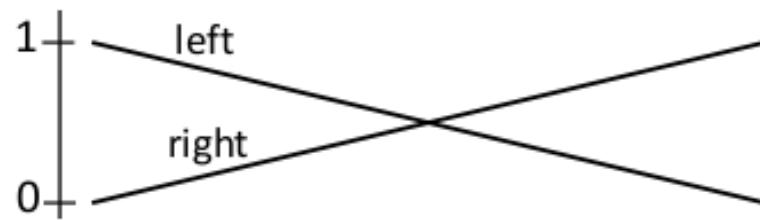
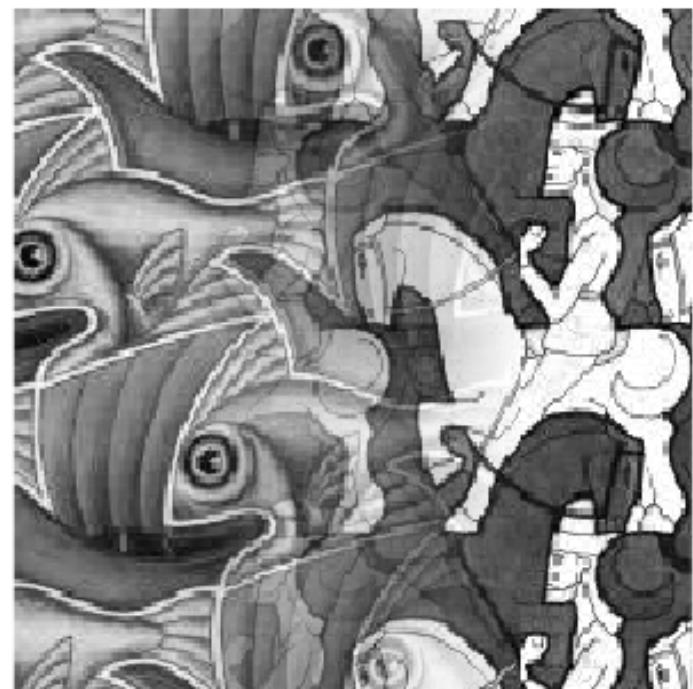
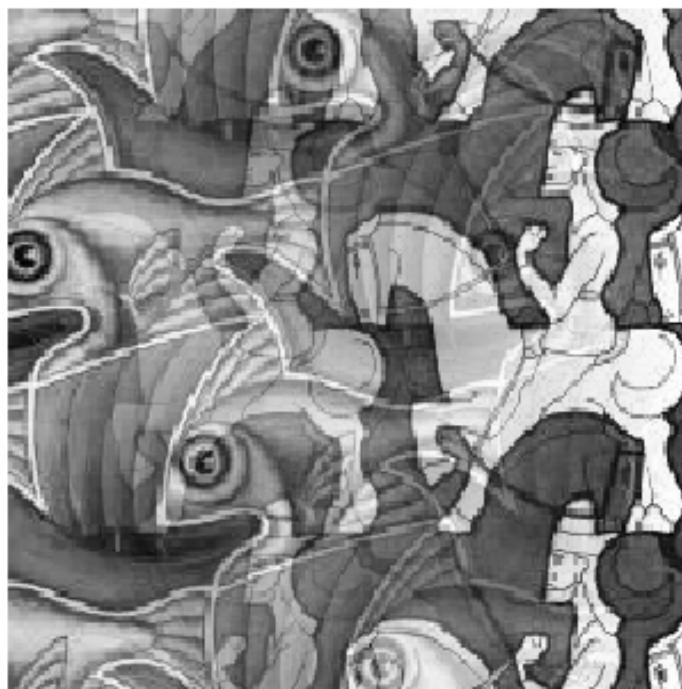


Image blending - feathering

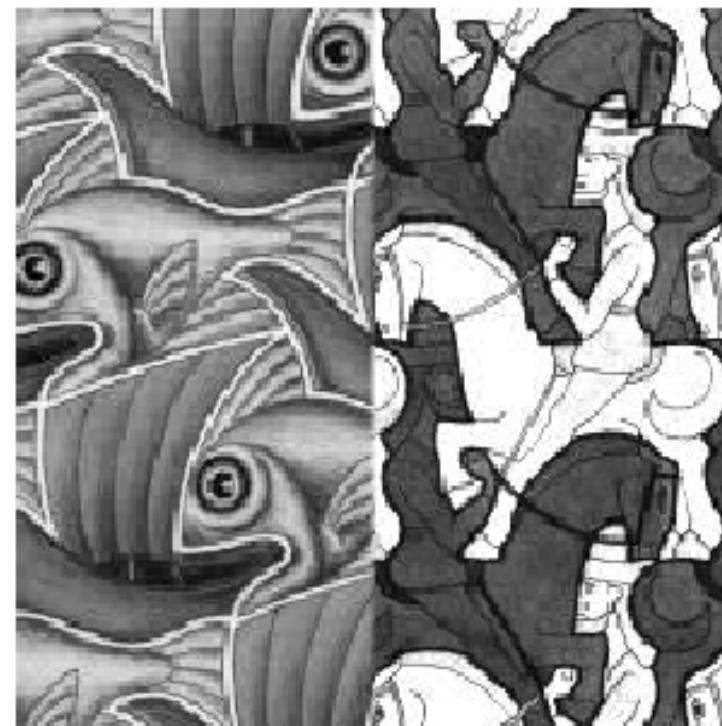
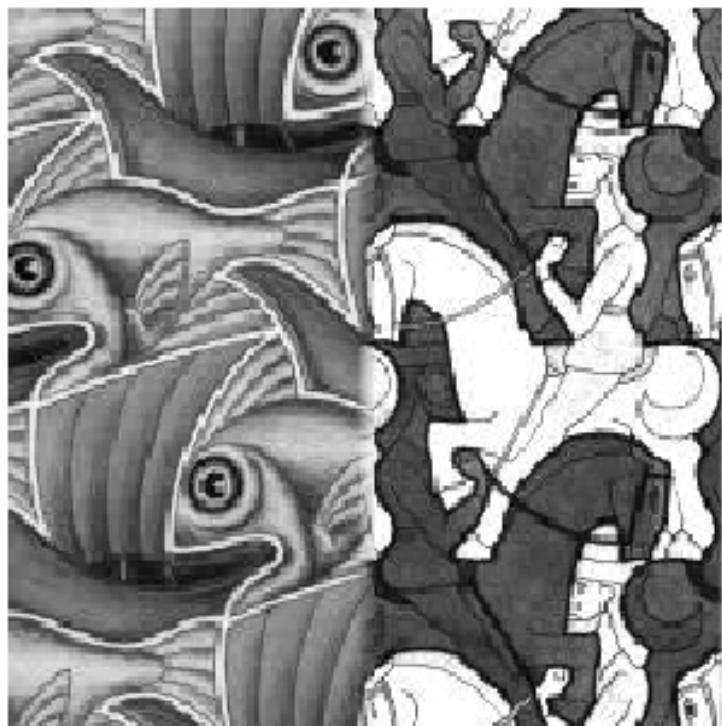


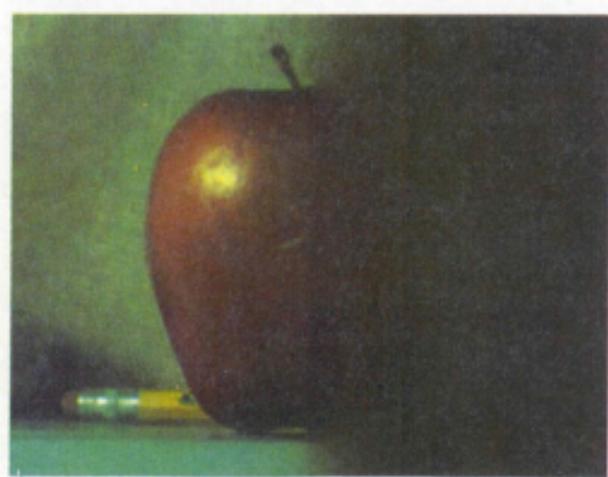
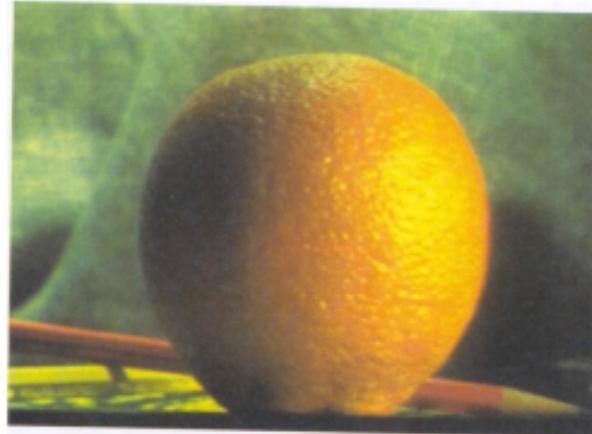
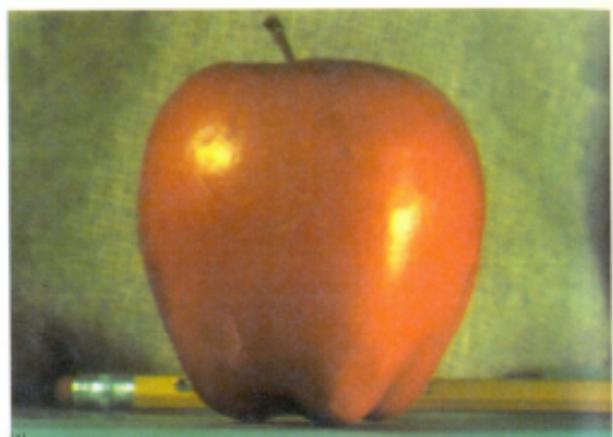
Image blending - feathering



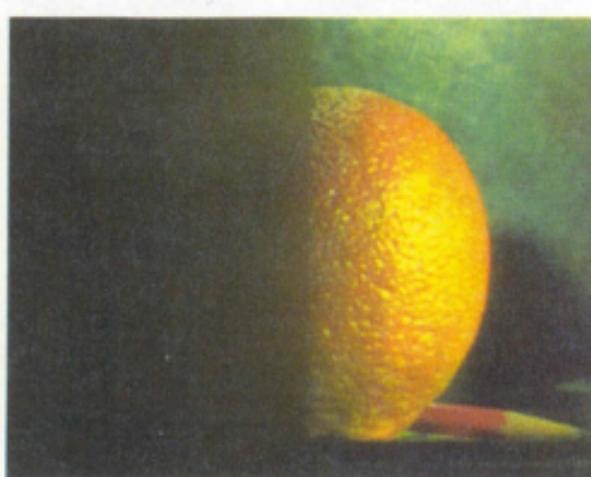
“Optimal” window: smooth but not ghosted

- Doesn't always work...

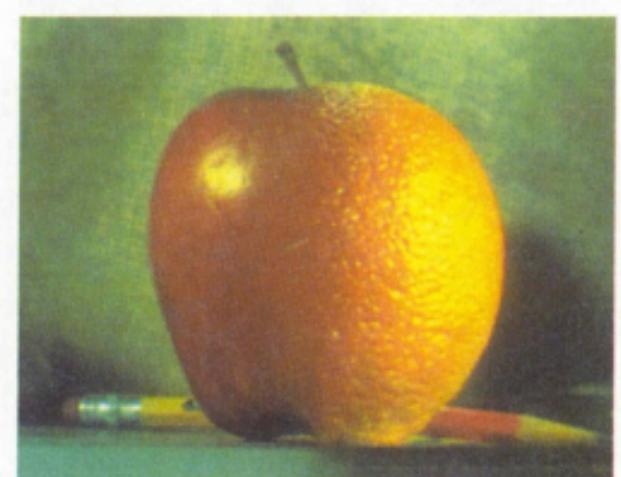
Image blending - pyramid blending



(d)



(h)



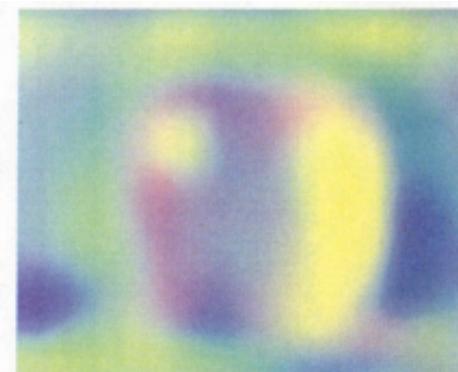
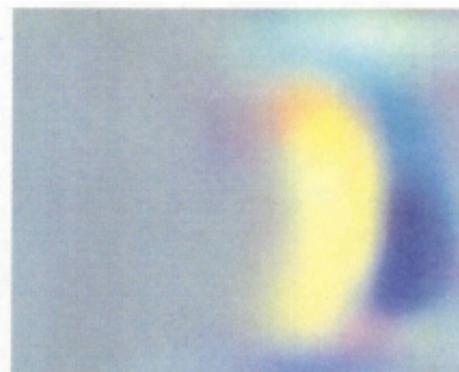
(l)

Burt, P. J. and Adelson, E. H., [A multiresolution spline with applications to image mosaics](#), ACM Transactions on Graphics, 42(4), October 1983, 217-236.

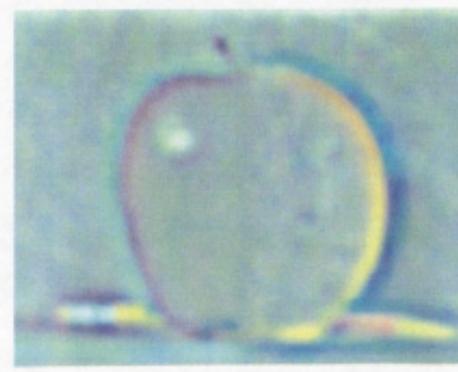
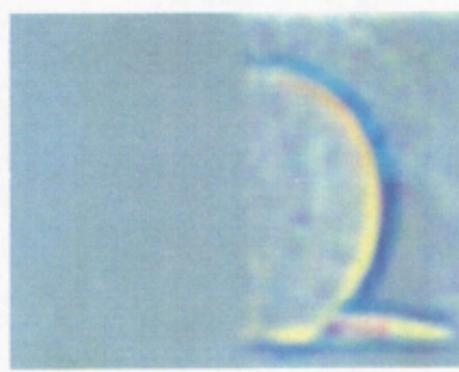
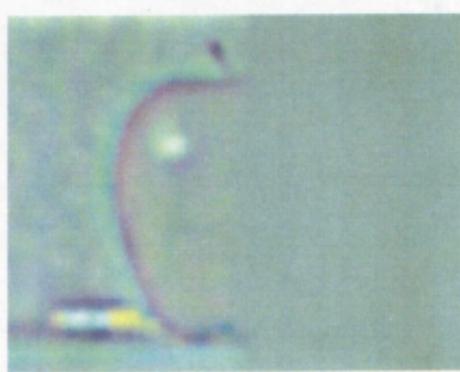
Szeliski

Image blending - pyramid blending

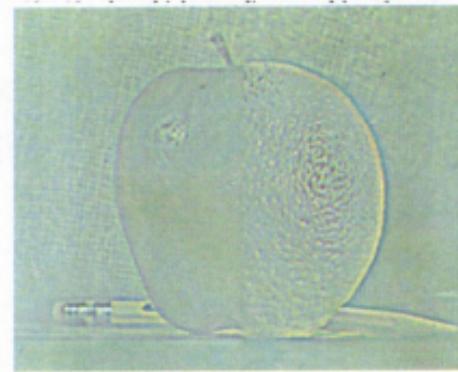
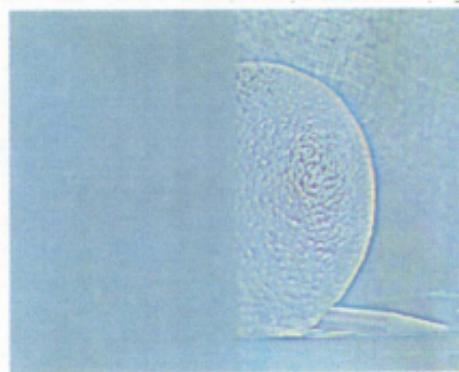
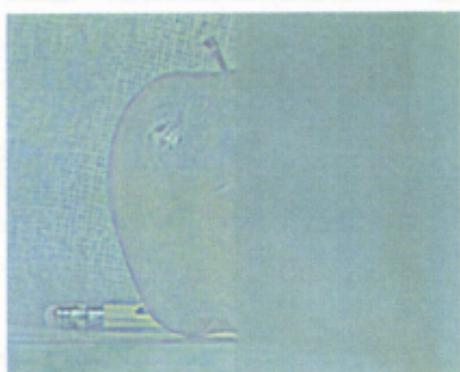
Laplacian
level
4



Laplacian
level
2



Laplacian
level
0



left pyramid

right pyramid

blended pyramid

Image blending - pyramid blending

1. Compute Laplacian pyramid
2. Compute Gaussian pyramid on *weight* image (can put this in A channel)
3. Blend Laplacians using Gaussian blurred weights
4. Reconstruct the final image
 - Q: How do we compute the original weights?
 - A: For horizontal panorama, use *mid-lines*
 - Q: How about for a general “3D” panorama?

Image blending - Poisson Image Editing



- Blend the gradients of the two images, then integrate
- For more info: Perez et al, SIGGRAPH 2003

Slide credits

- ▶ Trevor Darrell
- ▶ Kristen Grauman
- ▶ Forsyth and Ponce
- ▶ Rich Szeliski
- ▶ David Lowe
- ▶ Mathew Brown