

Topic 1 - Features

Jesse Hoey
School of Computing
University of Dundee

January 22, 2009

Scale Space

- at what scale are objects to be analysed?
- images contain information at many scales



Scale Space

- at what scale are objects to be analysed?
- images contain information at many scales
- best strategy: represent information at many scales



Gaussian Pyramid

- Pyramid of successively smaller images
- Smooth first,
- Interpolate and Downsample.
- Each Image gives image information at a particular scale

2D Convolution - Continuous Math

$$\begin{aligned} f(x, y) &= h(x, y) * g(x, y) \\ &= (f * g)(x, y) \\ &= \int_{-\infty}^{\infty} \int_{-\infty}^{\infty} h(x', y') g(x - x', y - y') dx' dy' \end{aligned}$$

Low Pass Filtering - Gaussian

$$h(x, y) = e^{-\frac{1}{2} \frac{(x^2 + y^2)}{\sigma^2}}$$

Is Separable:

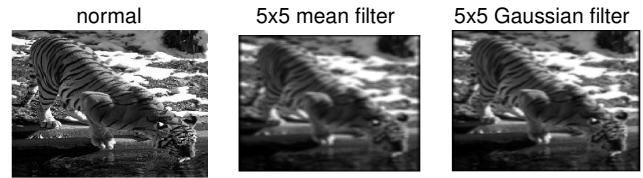
$$\begin{aligned} h(x, y) &= e^{-\frac{1}{2} \frac{x^2}{\sigma^2} + \frac{y^2}{\sigma^2}} \\ &= e^{-\frac{1}{2} \frac{x^2}{\sigma^2}} e^{-\frac{1}{2} \frac{y^2}{\sigma^2}} \end{aligned}$$

Low Pass Filtering - Gaussian

Advantages of Gaussian Kernel

- ▶ Separable
- ▶ Rotationally symmetric
- ▶ Smoothly falls to zero at edges
- ▶ Fourier Transform is a Gaussian

Low Pass Filtering - Gaussian



Transformation Equations

Affine transformation:

$$\begin{aligned} x' &= a_0x + a_1y + a_2 \\ y' &= b_0x + b_1y + b_2 \end{aligned}$$

As an array equation

$$\begin{bmatrix} x' \\ y' \end{bmatrix} = \begin{bmatrix} a_0 & a_1 \\ b_0 & b_1 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix} + \begin{bmatrix} a_2 \\ b_2 \end{bmatrix}$$

Or, using *homogeneous coordinates*

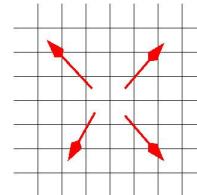
$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a_0 & a_1 & a_2 \\ b_0 & b_1 & b_2 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Scaling

1. Equation: Scaling Matrix

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} s & 0 & 0 \\ 0 & s & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

2.



Scaling Issues

	Small → Large	Large → Small
Technique:	Add Pixels (oversample or copy)	Delete Pixels (subsample)
Problems:		
Solutions:		

Scaling Issues

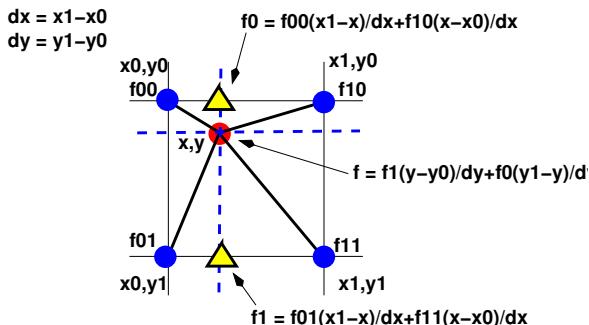
	Small → Large	Large → Small
Technique:	Add Pixels (oversample or copy)	Delete Pixels (subsample)
Problems:	Patchiness or “Blockiness”	Aliasing and missing data
Solutions:	smooth after	smooth before

Aliasing:

need to sample at **twice highest frequency**
if sample rate decreases **then** highest frequency *must decrease*, else aliasing

How to downsample

Bilinear Interpolation



What are interesting points at some scale?

Edges: compute Image Gradients

$$\vec{\nabla}f(x, y) = \frac{\partial f(x, y)}{\partial x} \vec{x} + \frac{\partial f(x, y)}{\partial y} \vec{y}$$

Simple gradient kernels (Prewitt)

$$\begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix} \quad \begin{bmatrix} -1 & -1 & -1 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{bmatrix}$$

Sobel Kernels (give more weight to on-axis pixels)

$$\begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix} \quad \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

Image Gradients

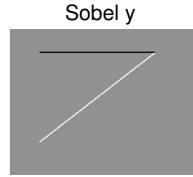
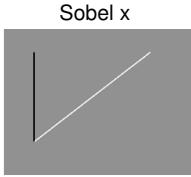
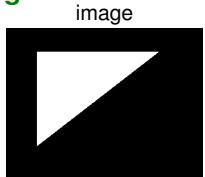
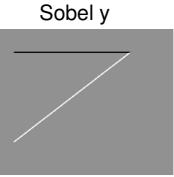
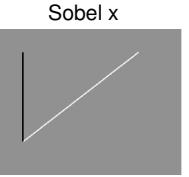
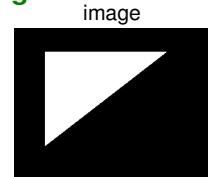
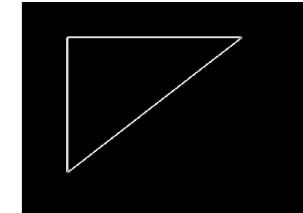


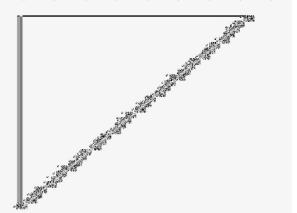
Image Gradients



Magnitude: $g_x^2 + g_y^2$

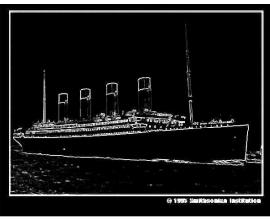


phase $\text{atan}(g_y/g_x)$



From Gradients to Edges

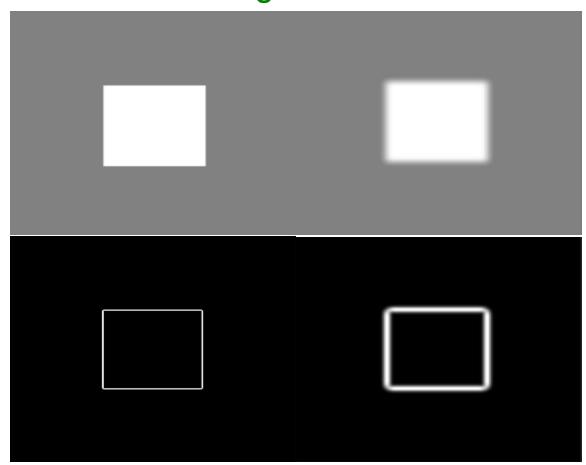
image and gradient magnitude:



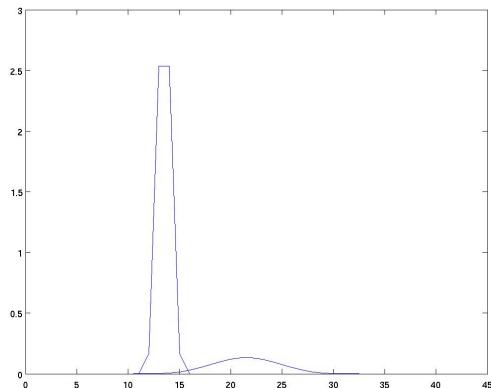
last step: threshold the magnitude.



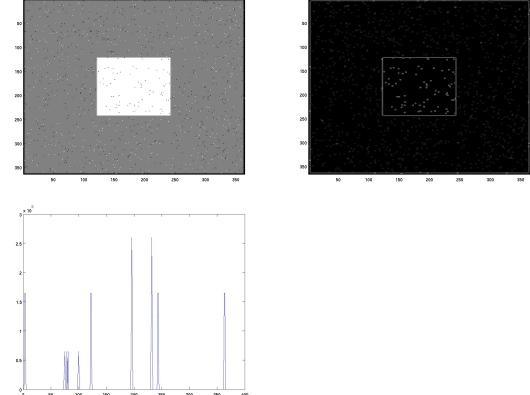
From Gradients to Edges



From Gradients to Edges



From Gradients to Edges



From Gradients to Edges

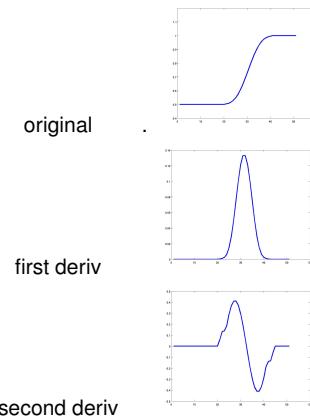
Problems with simple thresholding

- ▶ Sometimes edges are diffuse
- ▶ Noise can produce large gradient magnitudes
- ▶ Entire ridge (not maximum point) is detected

Need: Second order derivative

Second derivative allows us to find *maximum or minimum* values of gradient

From Gradients to Edges



Laplacian

$$\nabla^2 f(x, y) = \frac{\partial^2 f(x, y)}{\partial^2 x} + \frac{\partial^2 f(x, y)}{\partial^2 y}$$

$$\nabla^2 f(x, y) = \frac{\partial^2 f(x, y)}{\partial^2 x} + \frac{\partial^2 f(x, y)}{\partial^2 y}$$

but recall that gradient and convolution are Linear operators, so

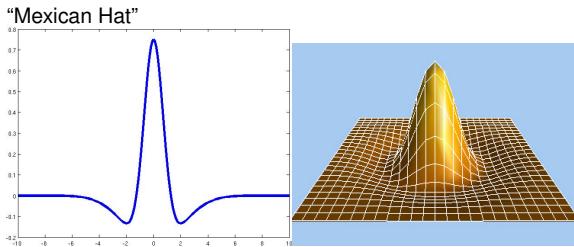
$$\nabla^2(\mathcal{N} * f(x, y)) = (\nabla^2 \mathcal{N}) * f(x, y)$$

taking Laplacian of a smoothed image = smoothing an image with a Laplacian of Gaussian

Laplacian

Laplacian of Gaussian

$$\nabla^2 \mathcal{N}(x; \sigma) = \left(\frac{x^2 - \sigma^2}{\sigma^4} \right) e^{-x^2/2\sigma^2}$$



Scale Space

scale-space representation $L(x; \sigma)$ of $f(x)$ is

$$L(x; \sigma) = f(x) * \mathcal{N}(x; \sigma) = \int_y f(x-y) \mathcal{N}(y; \sigma) dy$$

where $\mathcal{N}(y; \sigma)$ is Gaussian with variance σ : the *scale parameter*.

First derivative of $L(x; \sigma)$ at scale σ is then

$$\frac{\partial}{\partial x} L(x; \sigma) = \left(\frac{\partial}{\partial x} \mathcal{N}(x; \sigma) \right) * f(x)$$

Second derivative of $L(x; \sigma)$ at scale σ is then

$$\frac{\partial^2}{\partial x^2} L(x; \sigma) = \left(\frac{\partial^2}{\partial x^2} \mathcal{N}(x; \sigma) \right) * f(x)$$

Zero crossings of first or second derivative: local minima at each scale

Scale Space: Examples

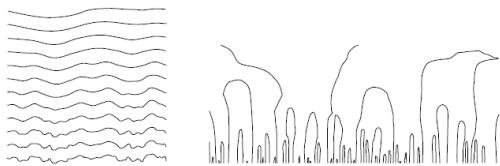


Figure 2: (a) The main idea of a scale-space representation is to generate a one-parameter family of derived signals in which the fine-scale information is successively suppressed. This figure shows a signal which has been successively smoothed by convolution with Gaussian kernels of increasing width. (b) Since new zero-crossings cannot be created by the diffusion equation in the one-dimensional case, the trajectories of zero-crossings in scale-space (here, zero-crossings of the second derivative) form paths across scales that are never closed from below.

Scale Space: Examples

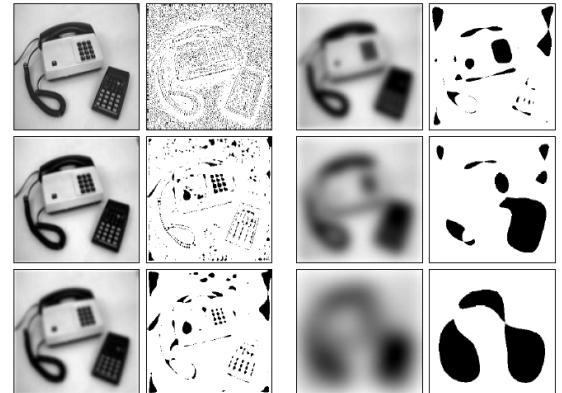


Figure 3: Different levels in the scale-space representation of a two-dimensional image at scale levels $t = 0, 2, 8, 32, 128$ and 512 together with grey-level blobs indicating local minima at each scale.

Laplacian of Gaussian

Approximation to Laplacian of Gaussian is Difference of Gaussians

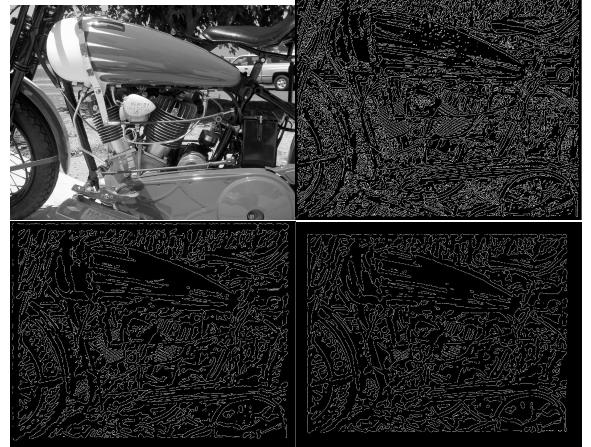
$$e^{-r^2/2\sigma_1^2} - e^{-r^2/2\sigma_2^2}$$

where $\sigma_1 < \sigma_2$

σ_1, σ_2 control the scale

Therefore, can combine Scale-space construction (Gaussian Pyramid) with second-order derivative analysis by subtracting levels of Gaussian Pyramid

Laplacian of Gaussian



Scale Space: biological vision

Receptive field profiles: superpositions of Gaussian derivatives

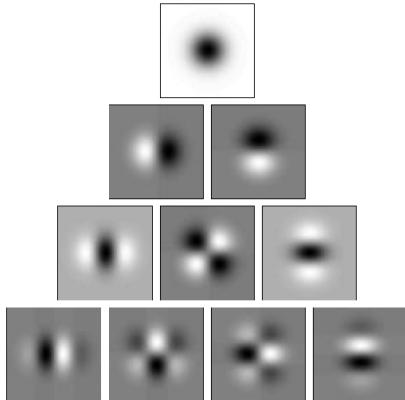


Figure 4: Gaussian derivative kernels up to order four for the two-dimensional case.

Review

1. Interesting stuff at multiple scales in an image
2. Construct **Gaussian Pyramid**
 - 2.1 successive levels are smaller images, represent larger scales
 - 2.2 smooth with Gaussian kernel (low-pass filter) before downsampling to avoid **aliasing**
 - 2.3 downsample using bilinear interpolation
3. Compute Laplacian of Gaussian at each level (or subtract levels if properly done)
 - 3.1 second derivative finds center of edges
 - 3.2 Gaussian filter smooths away noise
4. Find zero-crossings: interesting points
5. But ... edges are difficult to localize

From Edges to Corners

- ▶ Edge cannot be localised exactly
- ▶ Corner can be easily distinguished from its neighbors: good feature?
- ▶ Moravec: Simple Corner detection
- ▶ Harris: based on Moravec, but more complicated

Moravec Corner Detector

- ▶ Look for points with low *self-similarity*
 - ▶ Test each image patch to see if it looks the same as neighboring patches
 - ▶ Use sum of squared differences (SSD) between patches
- $$\sum_{\text{pixels}} (x_i - y_i)^2$$
- ▶ Uniform intensity: patches look similar, SSD is low
 - ▶ Edge pixel: patches look similar in direction parallel to edge, different in direction orthogonal to edge
 - ▶ Corner pixel: patches look different in all directions
 - ▶ Problem?

Moravec Corner Detector

- ▶ Look for points with low *self-similarity*
- ▶ Test each image patch to see if it looks the same as neighboring patches
- ▶ Use sum of squared differences (SSD) between patches

$$\sum_{\text{pixels}} (x_i - y_i)^2$$

- ▶ Uniform intensity: patches look similar, SSD is low
- ▶ Edge pixel: patches look similar in direction parallel to edge, different in direction orthogonal to edge
- ▶ Corner pixel: patches look different in all directions
- ▶ Problem? Non-isotropic: edge not in direction of neighbours .. interest point

Harris Corner Detector

- ▶ deals with isotropy of Moravec
- ▶ Basic idea: orient Moravec detector to the gradient direction

Auto-correlation function

$$c(x, y) = \sum_i [I(x_i, y_i) - I(x_i + \Delta x, y_i + \Delta y)]^2$$

Taylor expand the shifted image

$$I(x_i + \Delta x, y_i + \Delta y) \approx I(x_i, y_i) + \begin{bmatrix} I_x(x_i, y_i) & I_y(x_i, y_i) \end{bmatrix} \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix}$$

substitute into expression for $c(x, y)$

$$c(x, y) \approx \begin{bmatrix} \Delta x \Delta y \end{bmatrix} \begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_y I_x & \sum I_y^2 \end{bmatrix} \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix}$$

Harris Corner Detector

$$c(x, y) \approx [\Delta x \Delta y] \begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_y I_x & \sum I_y^2 \end{bmatrix} \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix}$$

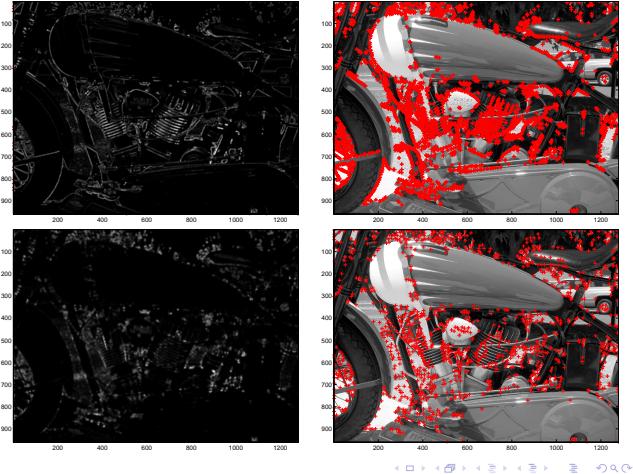
Look at "eigenvalues" of

$$C(x, y) = \begin{bmatrix} \sum I_x^2 & \sum I_x I_y \\ \sum I_y I_x & \sum I_y^2 \end{bmatrix}$$

eigenvalues characterise the change in I with Δx and Δy , both in terms of direction and quantity

1. if both are small: constant intensity patch (change in no direction)
2. if one is large: edge (change in one direction only)
3. if both are large: corner (change in both directions)

Moravec vs. Harris Corner Detectors



Scale Invariant Feature

- ▶ Associate a *characteristic scale* for each feature: the scale at which this feature "lives"
- ▶ The scale at which the operator response is maximal
- ▶ scaled Laplacian of Gaussian best for scale detection:

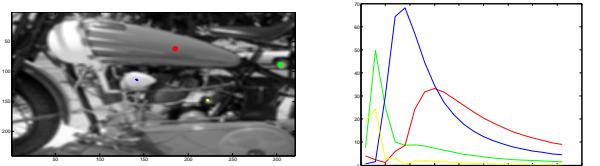
$$\sigma^2 \nabla \mathcal{N}$$

But can show that

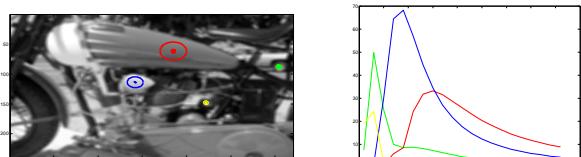
$$\mathcal{N}(x, y, k\sigma) - \mathcal{N}(x, y, \sigma) \approx (k - 1)\sigma^2 \nabla \mathcal{N}$$

- ▶ Therefore, filter with a difference of Gaussians (DoG), and look for maximal response

Characteristic Scale



Characteristic Scale



Scale-Adapted Harris Corner Detector

- ▶ Look for maximal response of Harris detector at each scale
- ▶ For each maximal response point, find the characteristic scale

Affine Invariant Feature

- ▶ Pick a *characteristic orientation (rotation)* for each feature
- ▶ The orientation of the gradient
- ▶ Pick a *characteristic affine scale ratio* for each feature
- ▶ Based on second order gradient

In general, pick a characteristic *affine transformation* for each feature:

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} a_0 & a_1 & a_2 \\ b_0 & b_1 & b_2 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ 1 \end{bmatrix}$$

Object Recognition

Given: an image of an object and a new image

We want to know: is the object in the new image?

How to do it:

1. Find features on the object.
 - 1.1 Each feature has a characteristic *affine transformation*
 - 1.2 Compute local *affine invariant* descriptor at each feature
2. Find features in new image (in same way)
3. **Match** new image features with object features by comparing *affine invariant* descriptors
4. These matches should succeed even if object has undergone an affine transformation

